



# Warp: Advancing Simulation AI with Differentiable GPU computing in Python

Miles Macklin

Director, Simulation Technology, NVIDIA

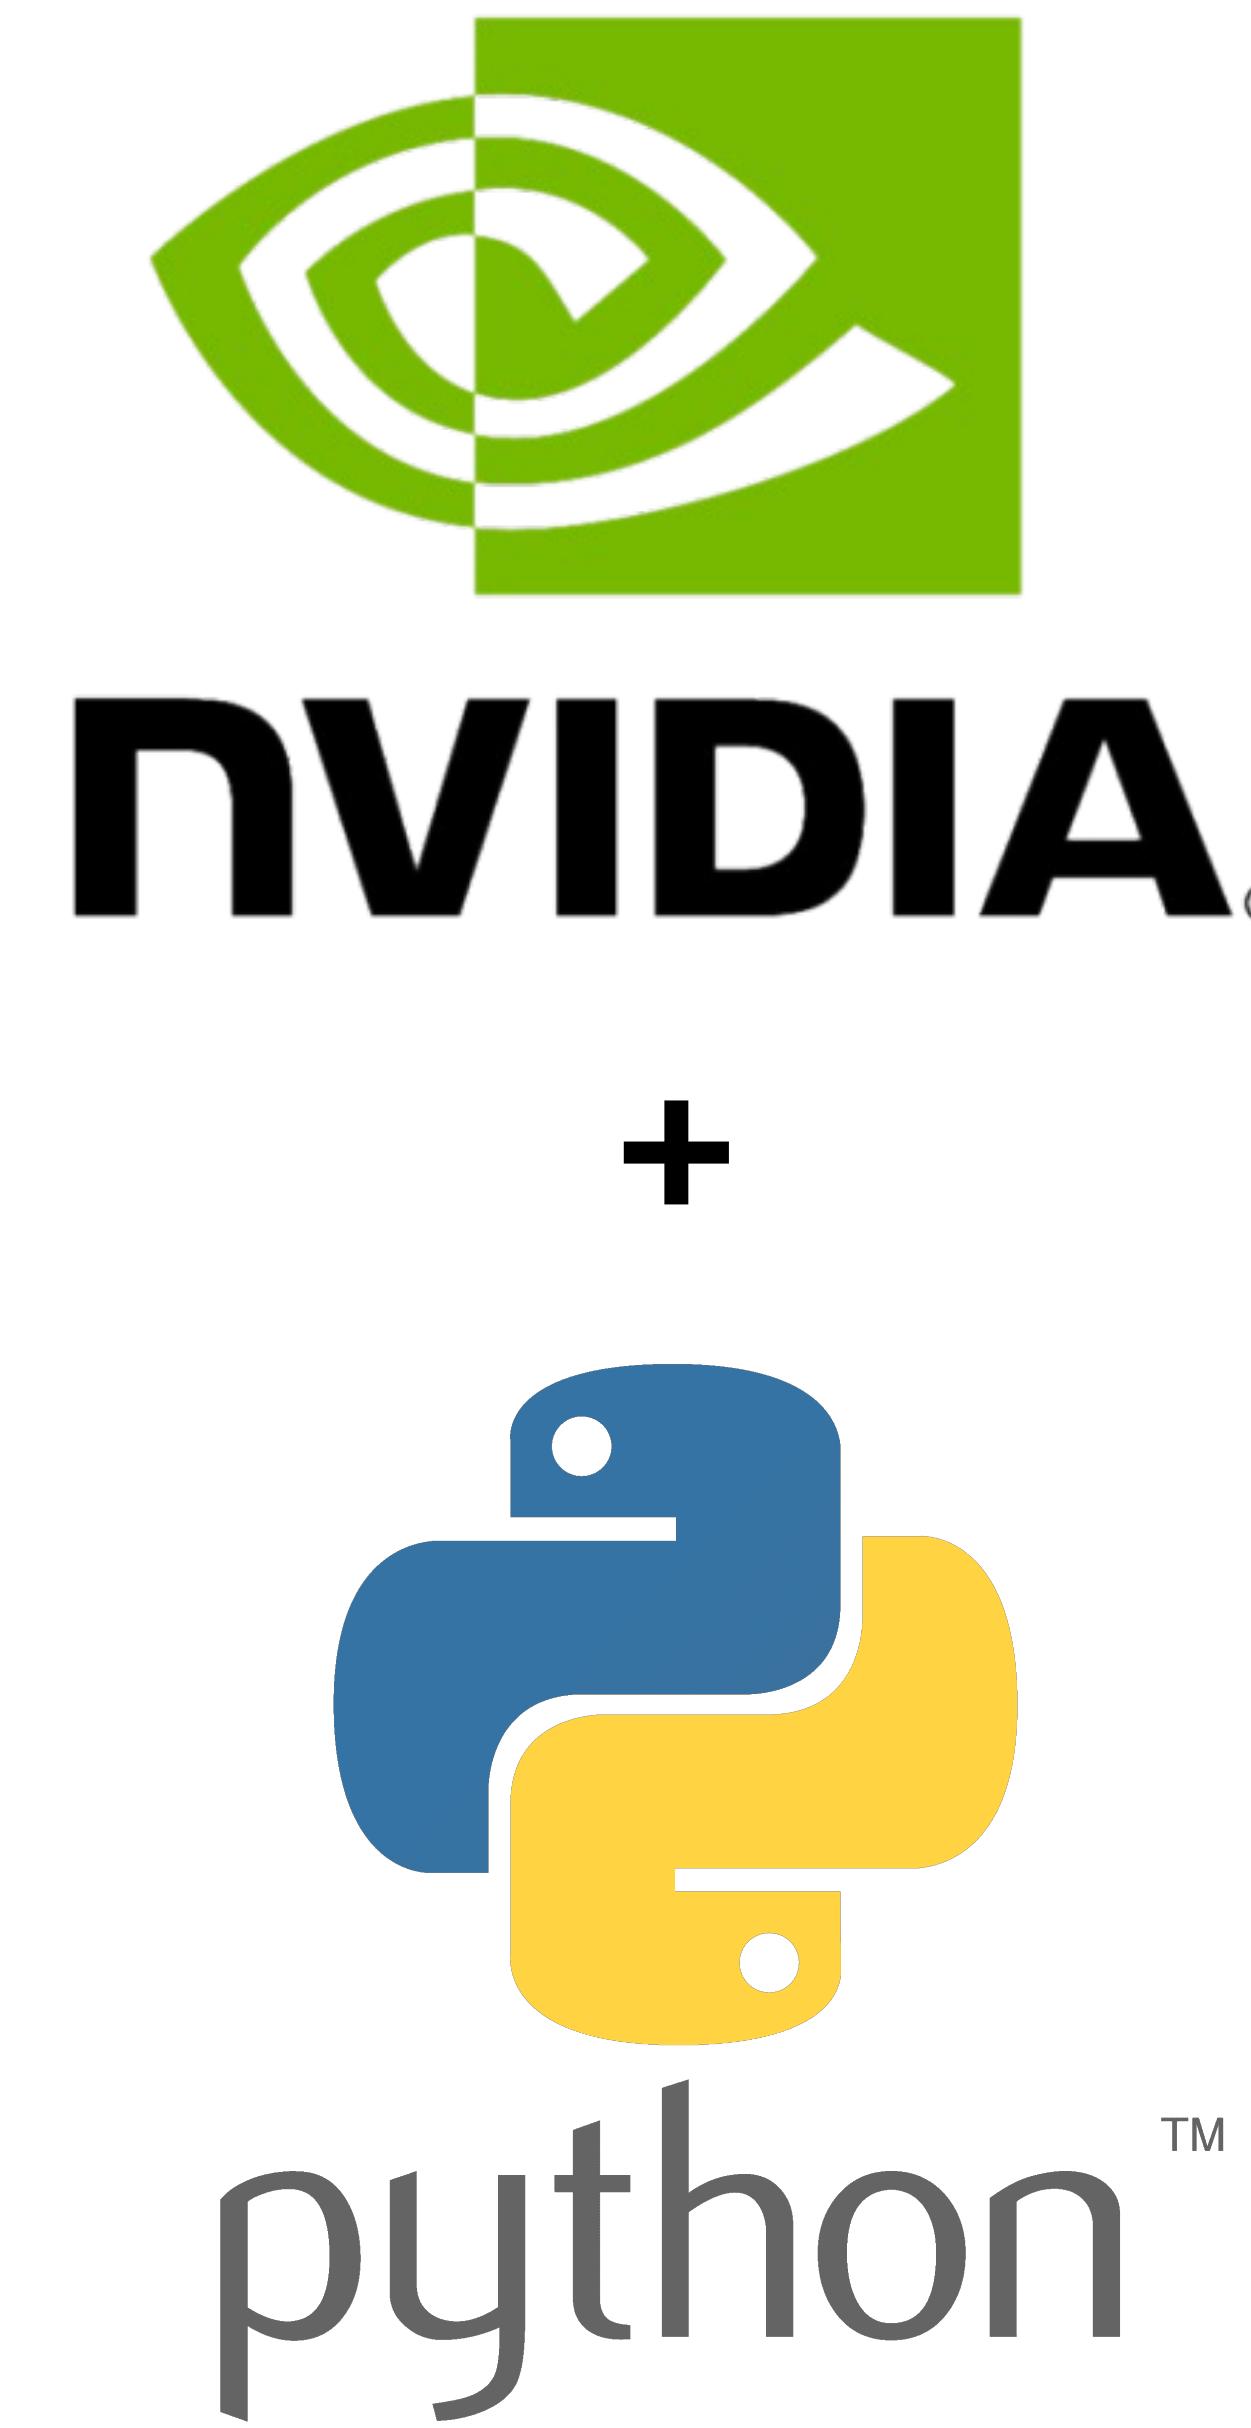


# Agenda

- Introduction
  - SDK Overview
  - Case Studies
  - Road Map
  - Questions
- 
- 
- 
-

# Problem Statement

- CUDA is general purpose, efficient, but device centric:
  - C++ focused
  - Not specialized for simulation
  - Build everything yourself
  - Not differentiable
- Python is the lingua franca for AI
  - Low barrier to entry
  - Fast iteration and deployment
  - Rich ecosystem of DL frameworks
  - Often has poor performance
- How can we bridge the gap between CUDA and Python for simulation developers?



# Introducing Warp

## Differentiable Spatial Computing for Python

```
pip install warp-lang
```

### 1. GPU Kernels in Python

- Write CUDA kernels in 100% Python syntax
- Runtime JIT compilation
- Fast developer iteration

### 2. Spatial Computing

- Rich vector math library
- Mesh processing and queries
- Sparse volumes (OpenVDB)
- Hash grids

### 3. Differentiable Programming

- Auto-differentiation
- Forward + backward kernel generation
- Interop with DL frameworks (e.g.: PyTorch, JAX)

### 4. Omniverse Integration

- OpenUSD import / export
- Runtime extensions for Isaac / Composer

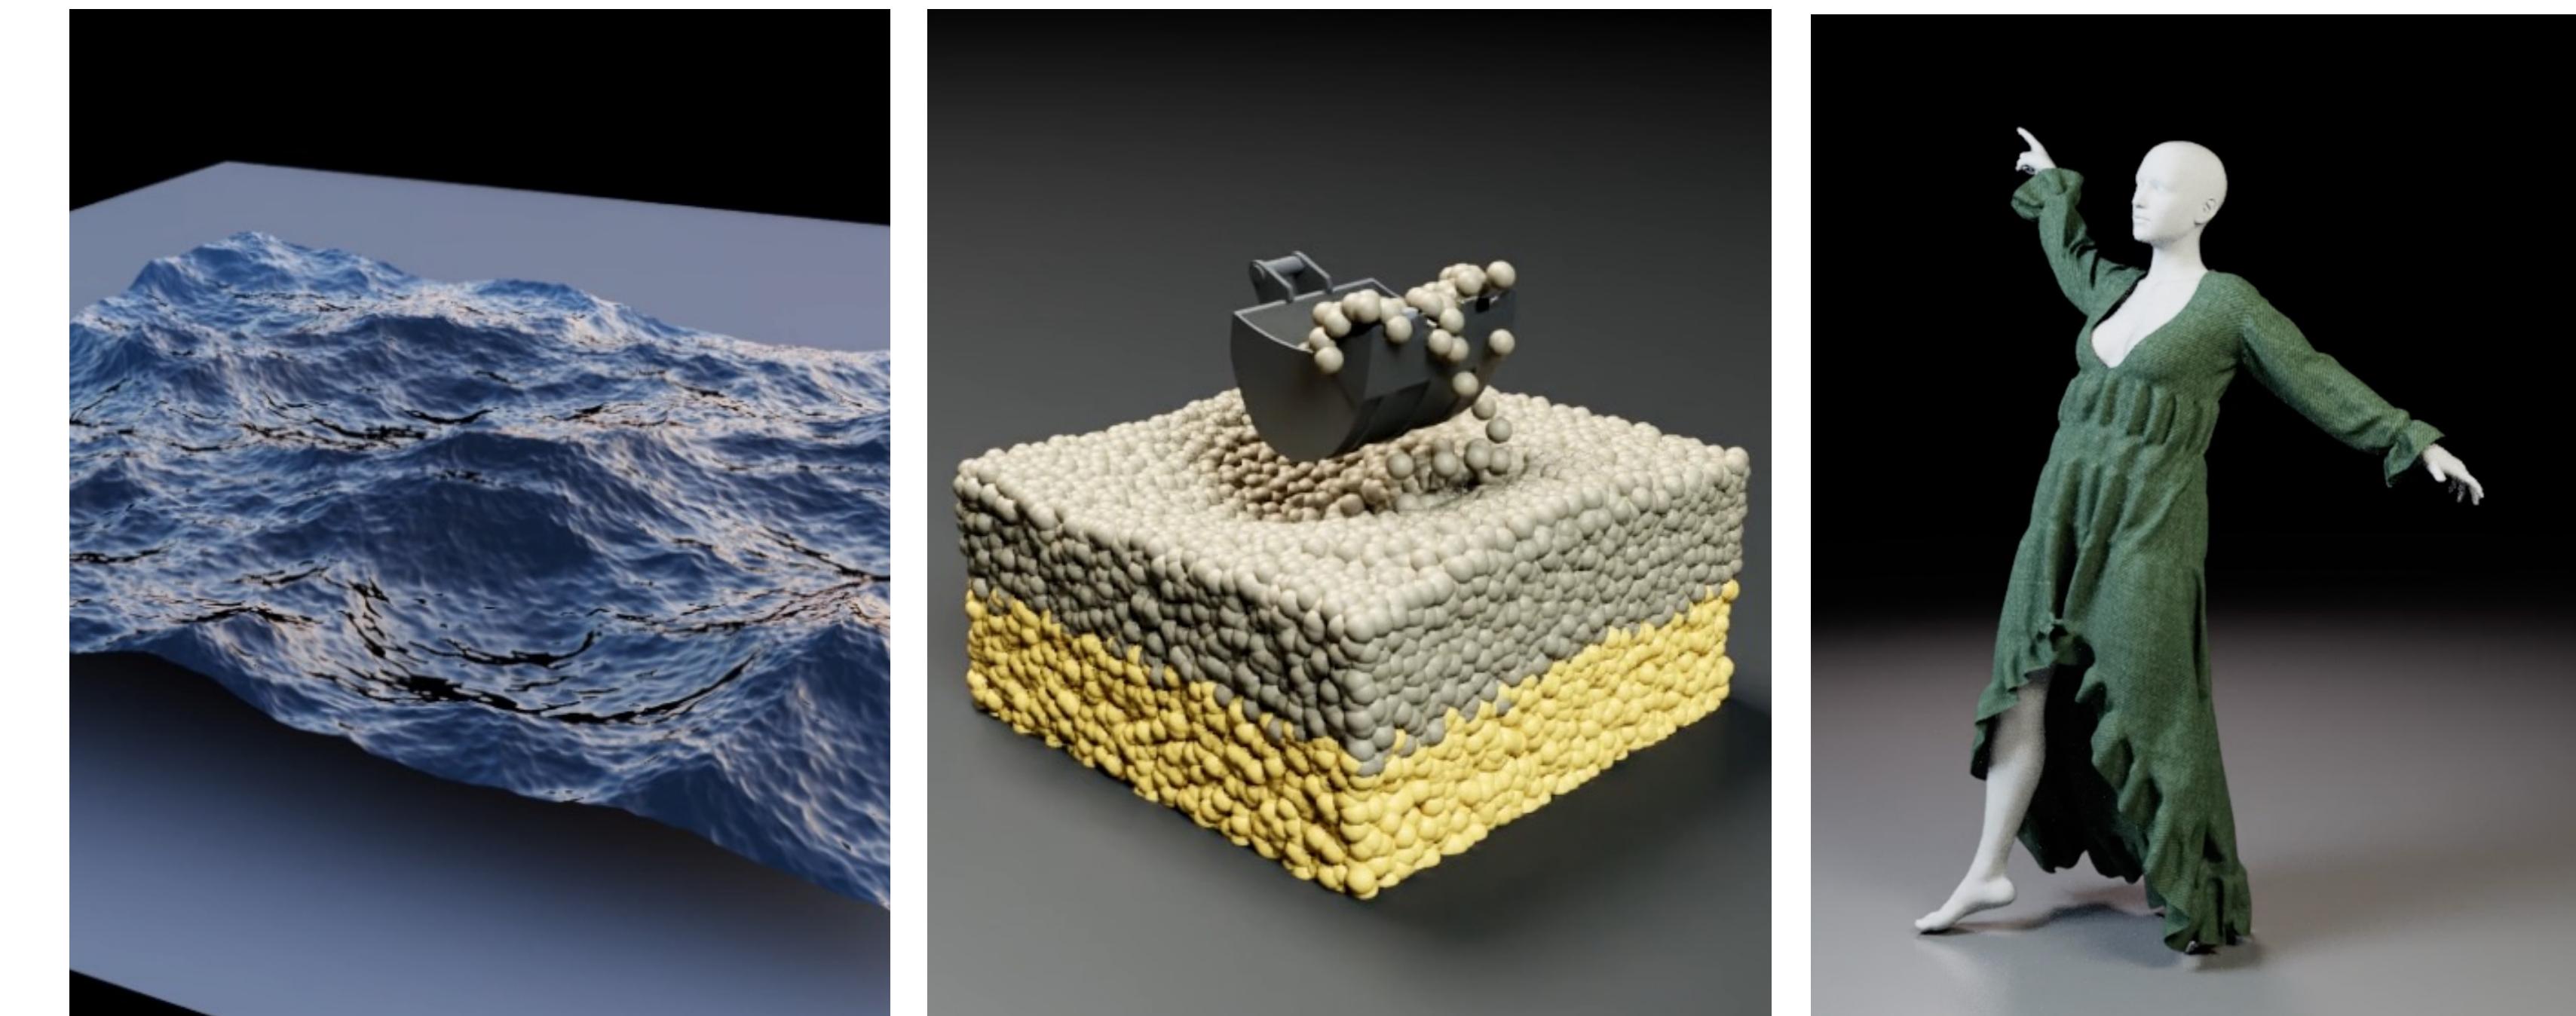
```
import warp as wp

@wp.kernel
def integrate(p: wp.array(dtype=wp.vec3),
              v: wp.array(dtype=wp.vec3),
              f: wp.array(dtype=wp.vec3),
              m: wp.array(dtype=float)):

    # thread id
    tid = wp.tid()

    # Semi-implicit Euler step
    v[tid] = v[tid] + (f[tid] * m[tid] + wp.vec3(0.0, -9.8, 0.0)) * dt
    x[tid] = x[tid] + v[tid] * dt

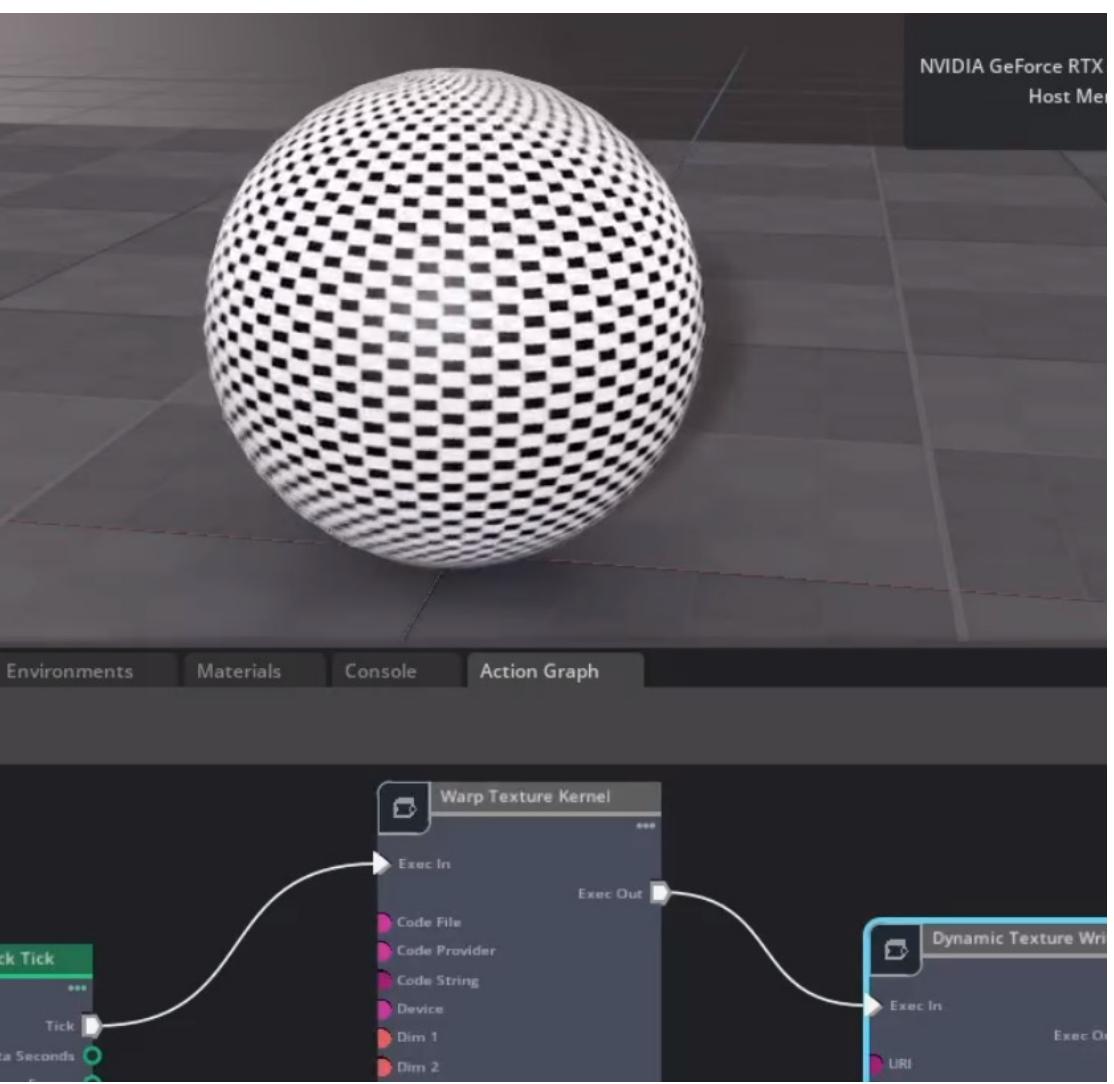
    # kernel launch
    wp.launch(integrate, dim=1024, inputs=[x, v, f, ...], device="cuda:0")
```



# Warp Use Cases

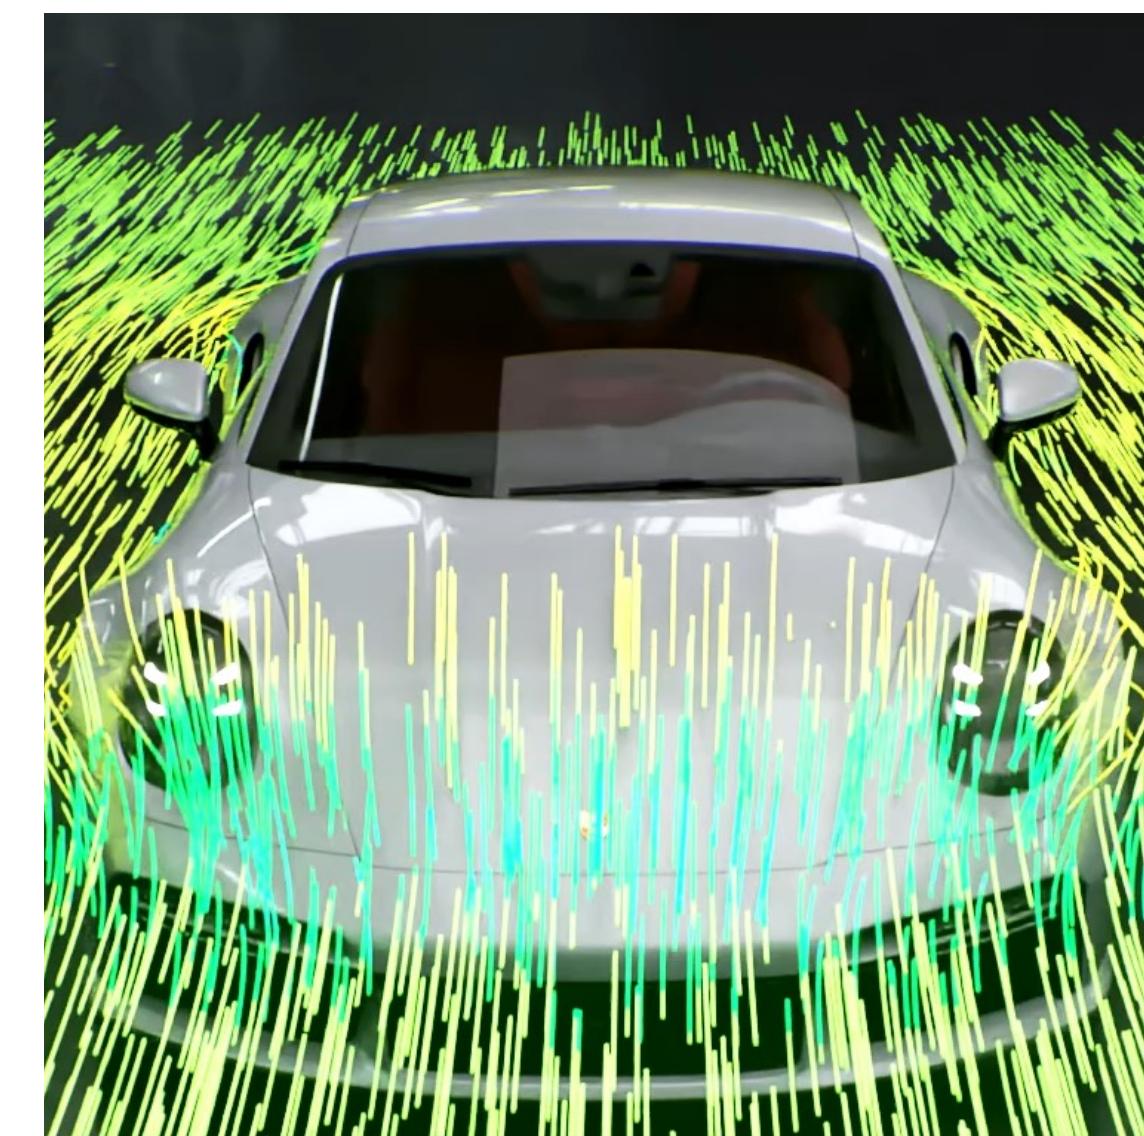
- Warp provides the **building blocks** for developers to create, accelerate and extend their own simulators

## Data Processing



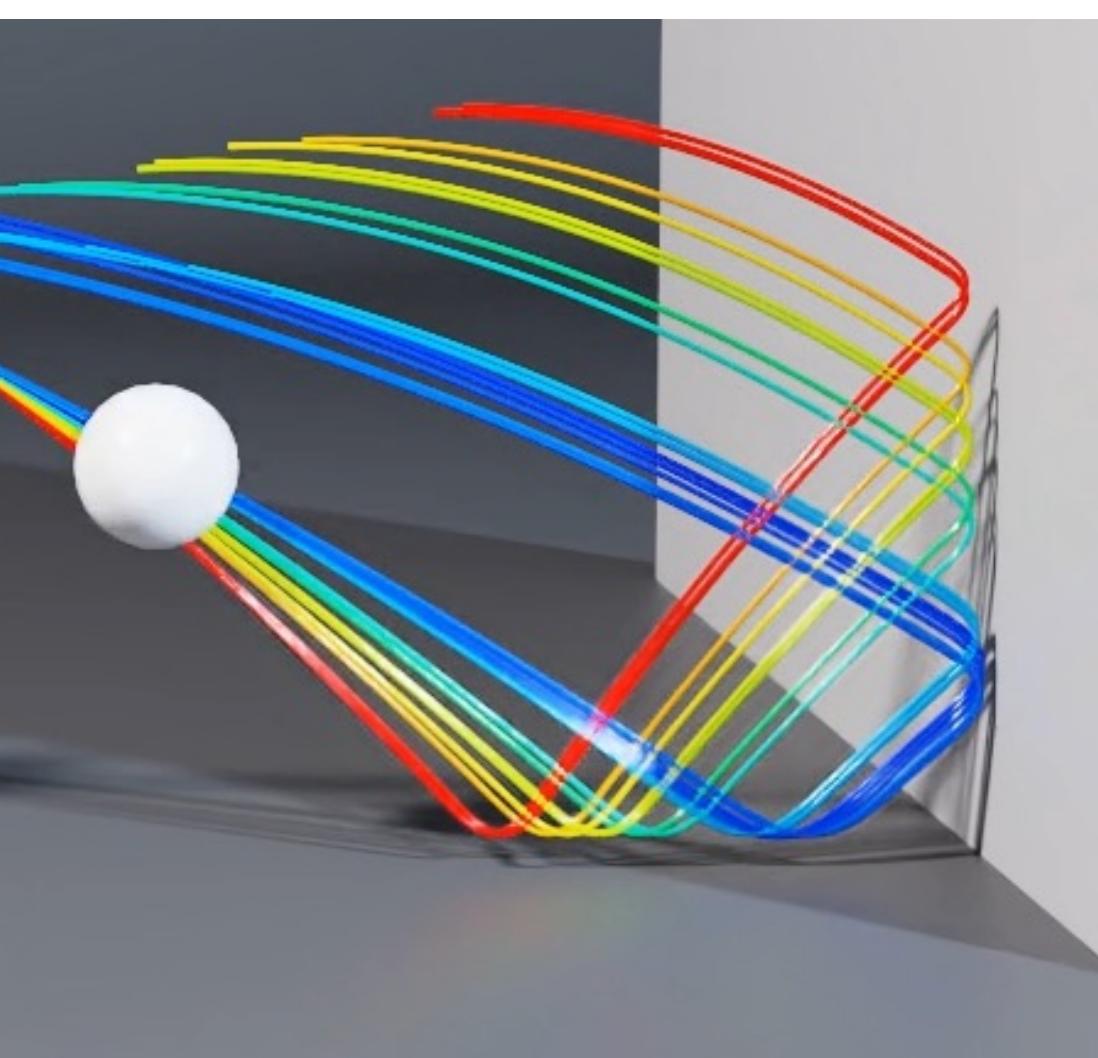
- Mesh processing
- Image processing
- Synthetic data generation

## Simulation



- Rigid body dynamics
- Elasticity
- Fluid flow

## Training



- Neural dynamics
- Parameter estimation
- Trajectory optimization
- Inverse problems

## Scripting

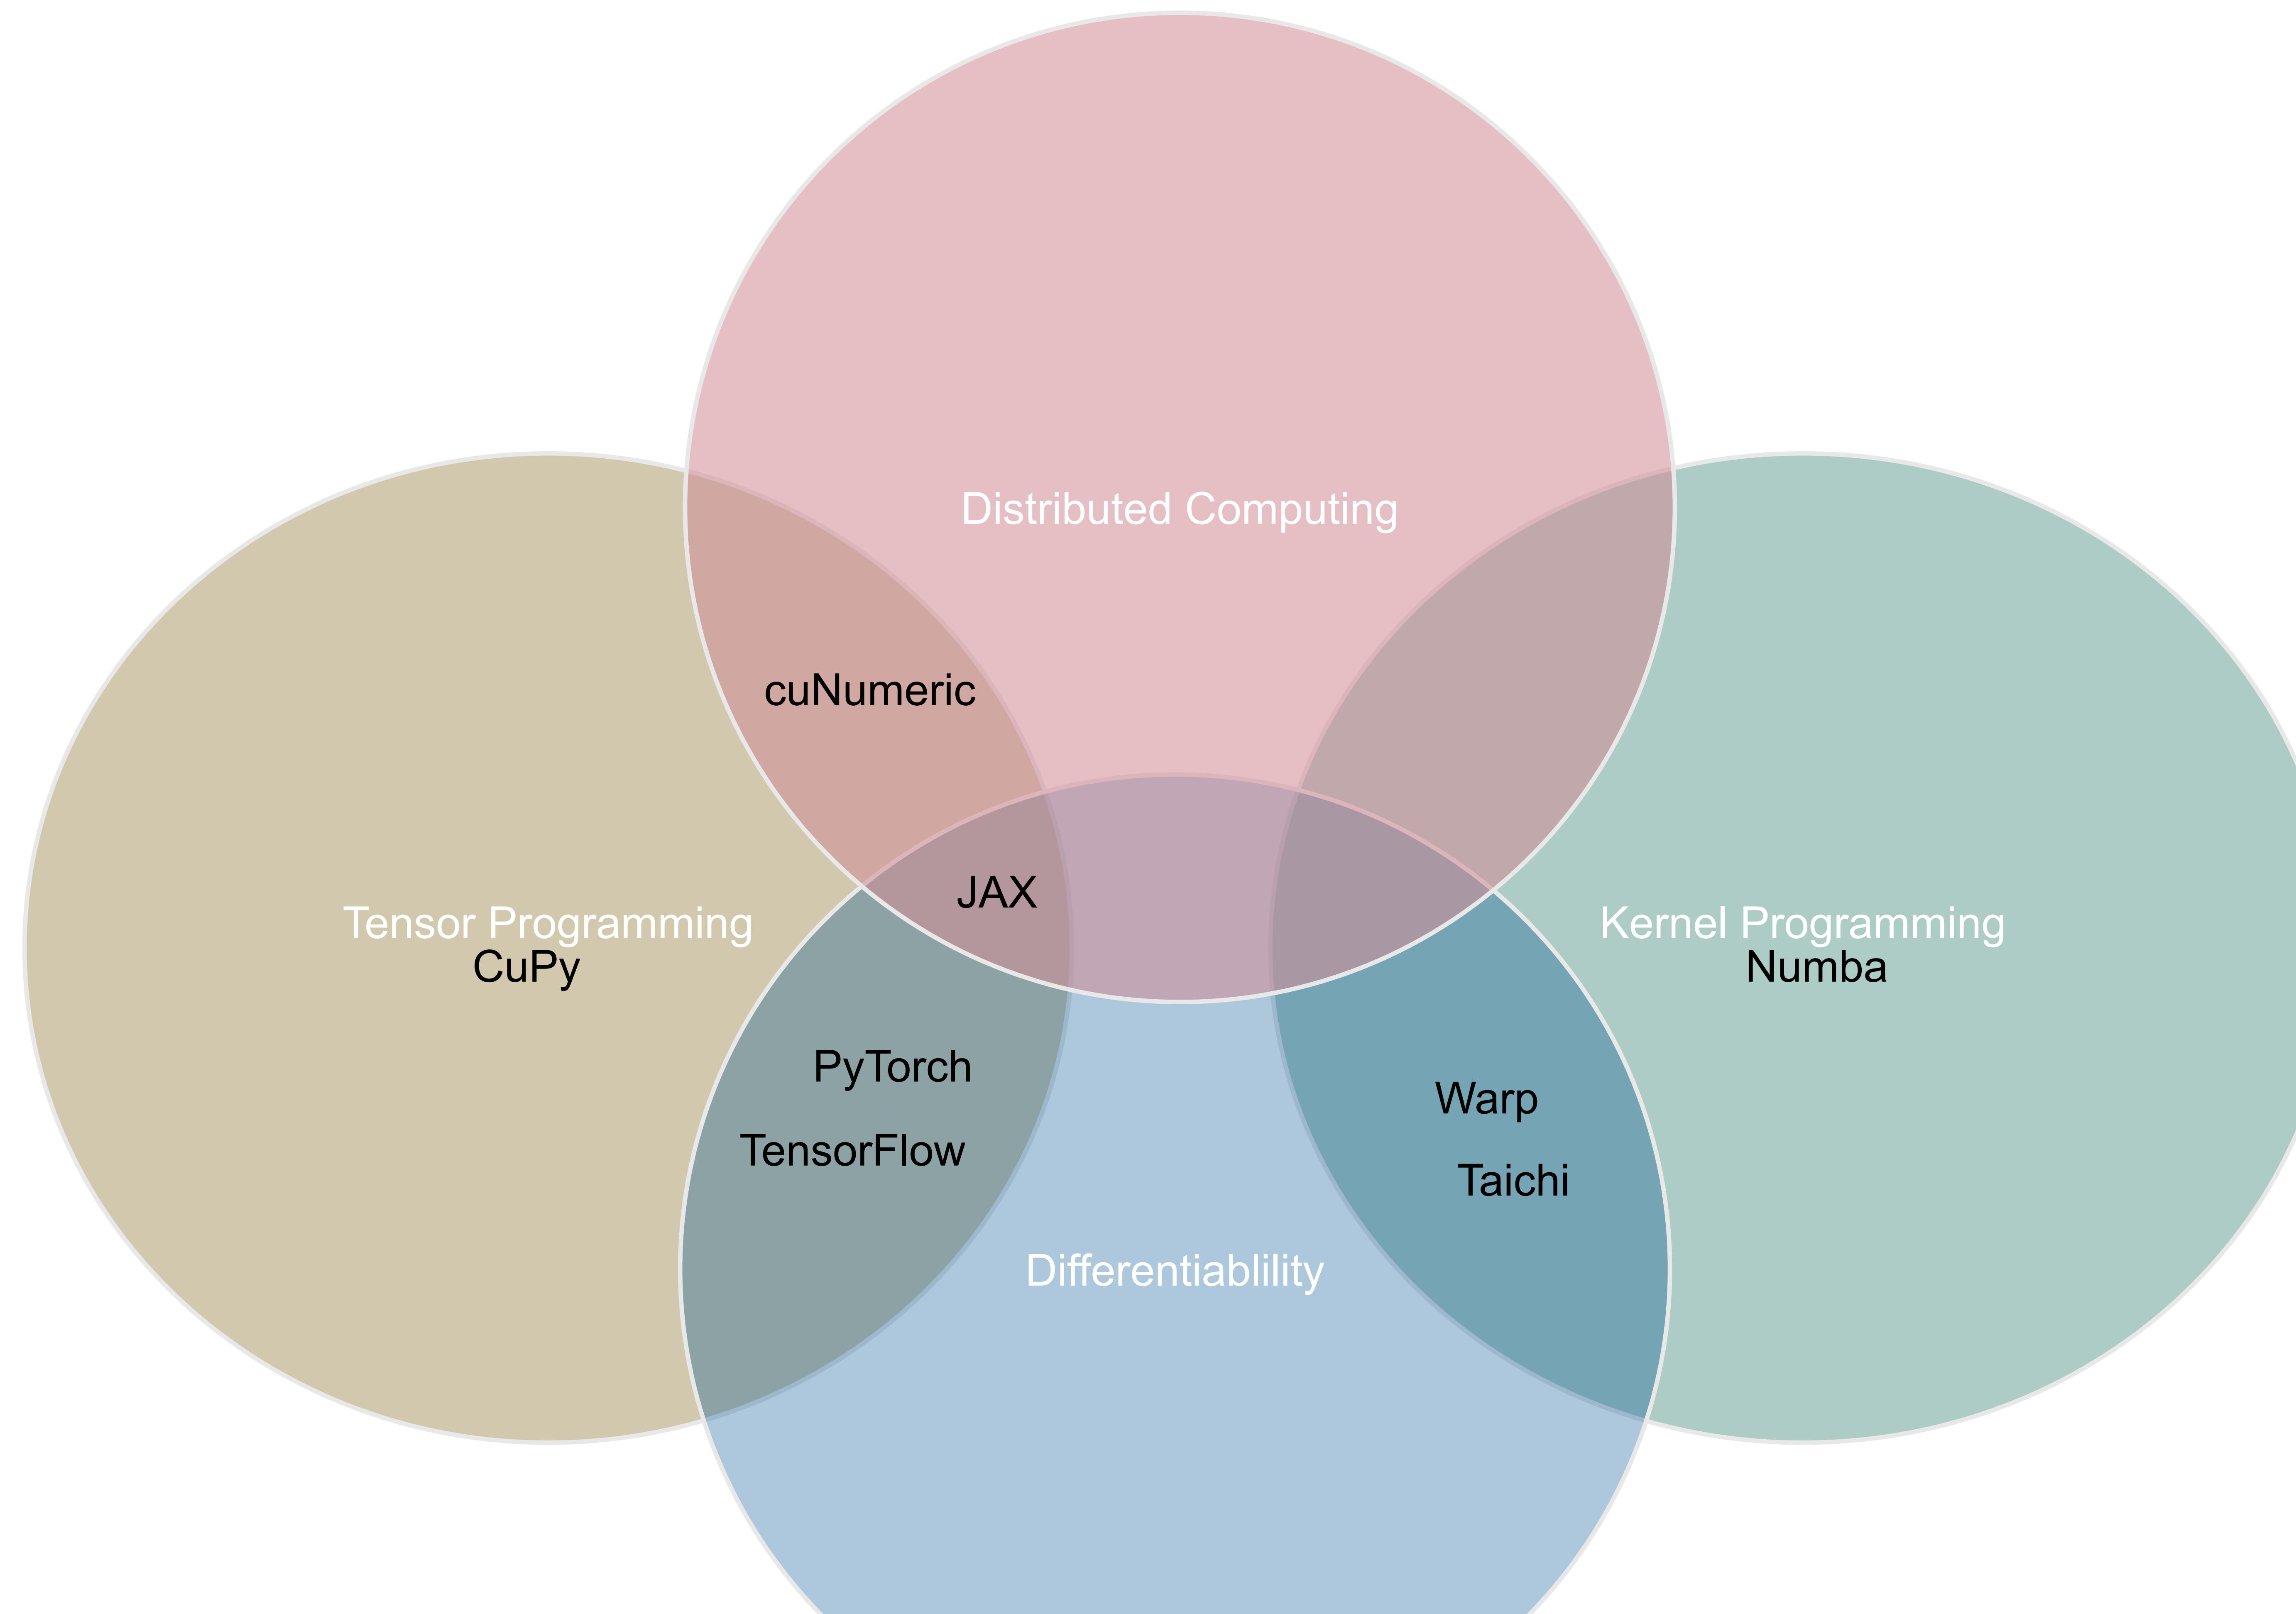
```
@wp.kernel
def initialize_particles(
    particle_x: wp.array(dtype=wp.vec3), sm):
    tid = wp.tid()

    # grid size
    nr_x = wp.int32(width / 4.0 / smoothing)
    nr_y = wp.int32(height / smoothing_length)
    nr_z = wp.int32(length / 4.0 / smoothing)

    # calculate particle position
    z = wp.float(tid % nr_z)
    y = wp.float((tid // nr_z) % nr_y)
    x = wp.float((tid // (nr_z * nr_y)) % nr_x)
    pos = smoothing_length * wp.vec3(x, y, z)
```

- Loss/reward functions
- Custom forces
- Custom behaviors
- Custom boundaries

# Python GPU Ecosystem





# Warp SDK

# Warp Python Modules

## • warp.core

- Differentiable kernel coding for Python
- Math, geometry, vector library

## • warp.sim

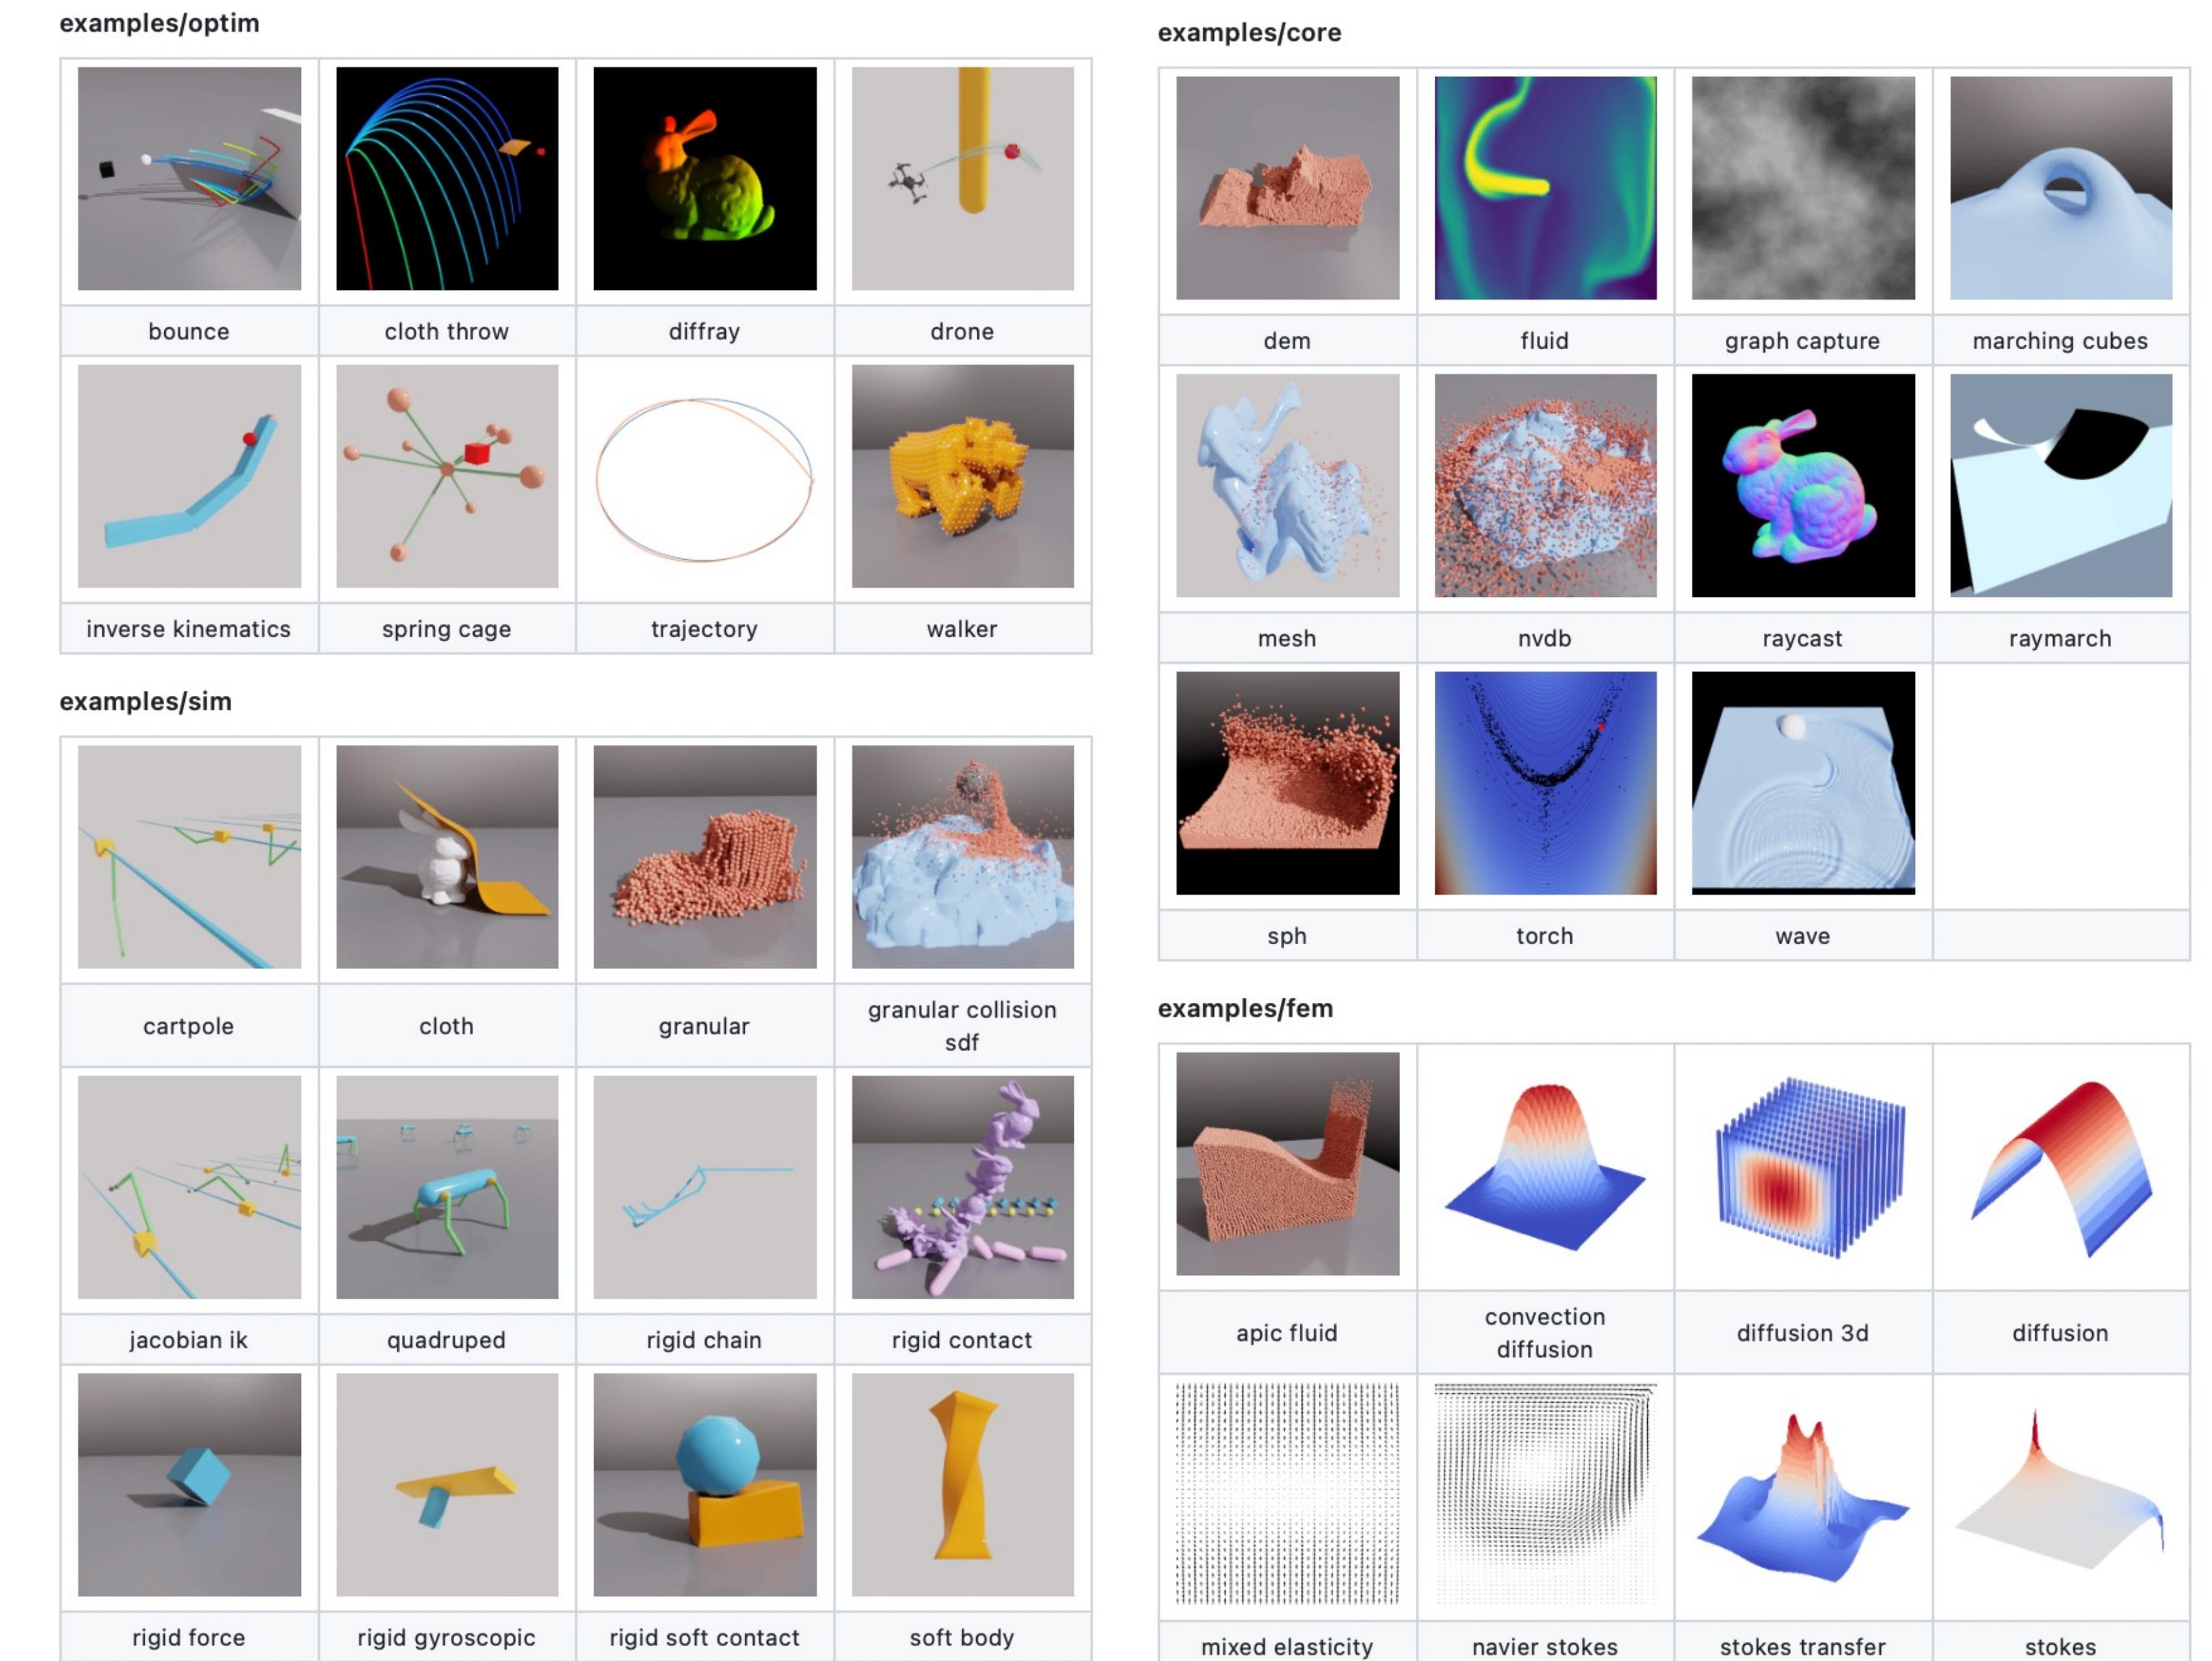
- Differentiable **real-time simulation** for robotic control + prediction
- Rigid bodies, soft bodies, particles, cloth
- URDF, MJCF, UsdPhysics parsers

## • warp.fem (early access)

- Differentiable PDE framework for heat transfer, diffusion, elasticity
- Fast iteration, but **offline** focus
- Not replacing existing codes, but potential to build on them

## • warp.llm (early access)

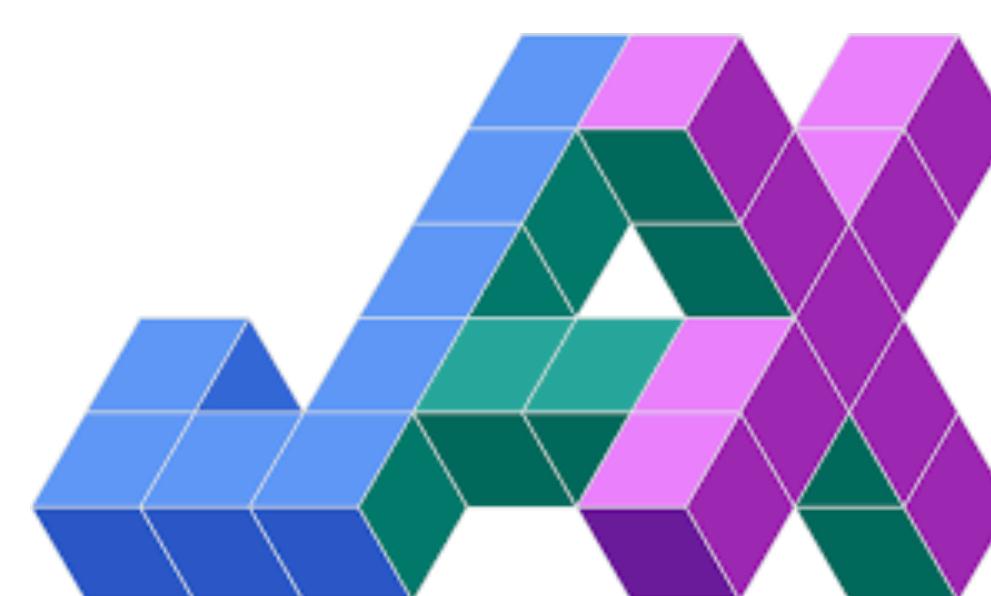
- AI agent specialized for Warp kernel coding
- Generate simulation code directly from prompts
- Generate loss functions from code -> optimization



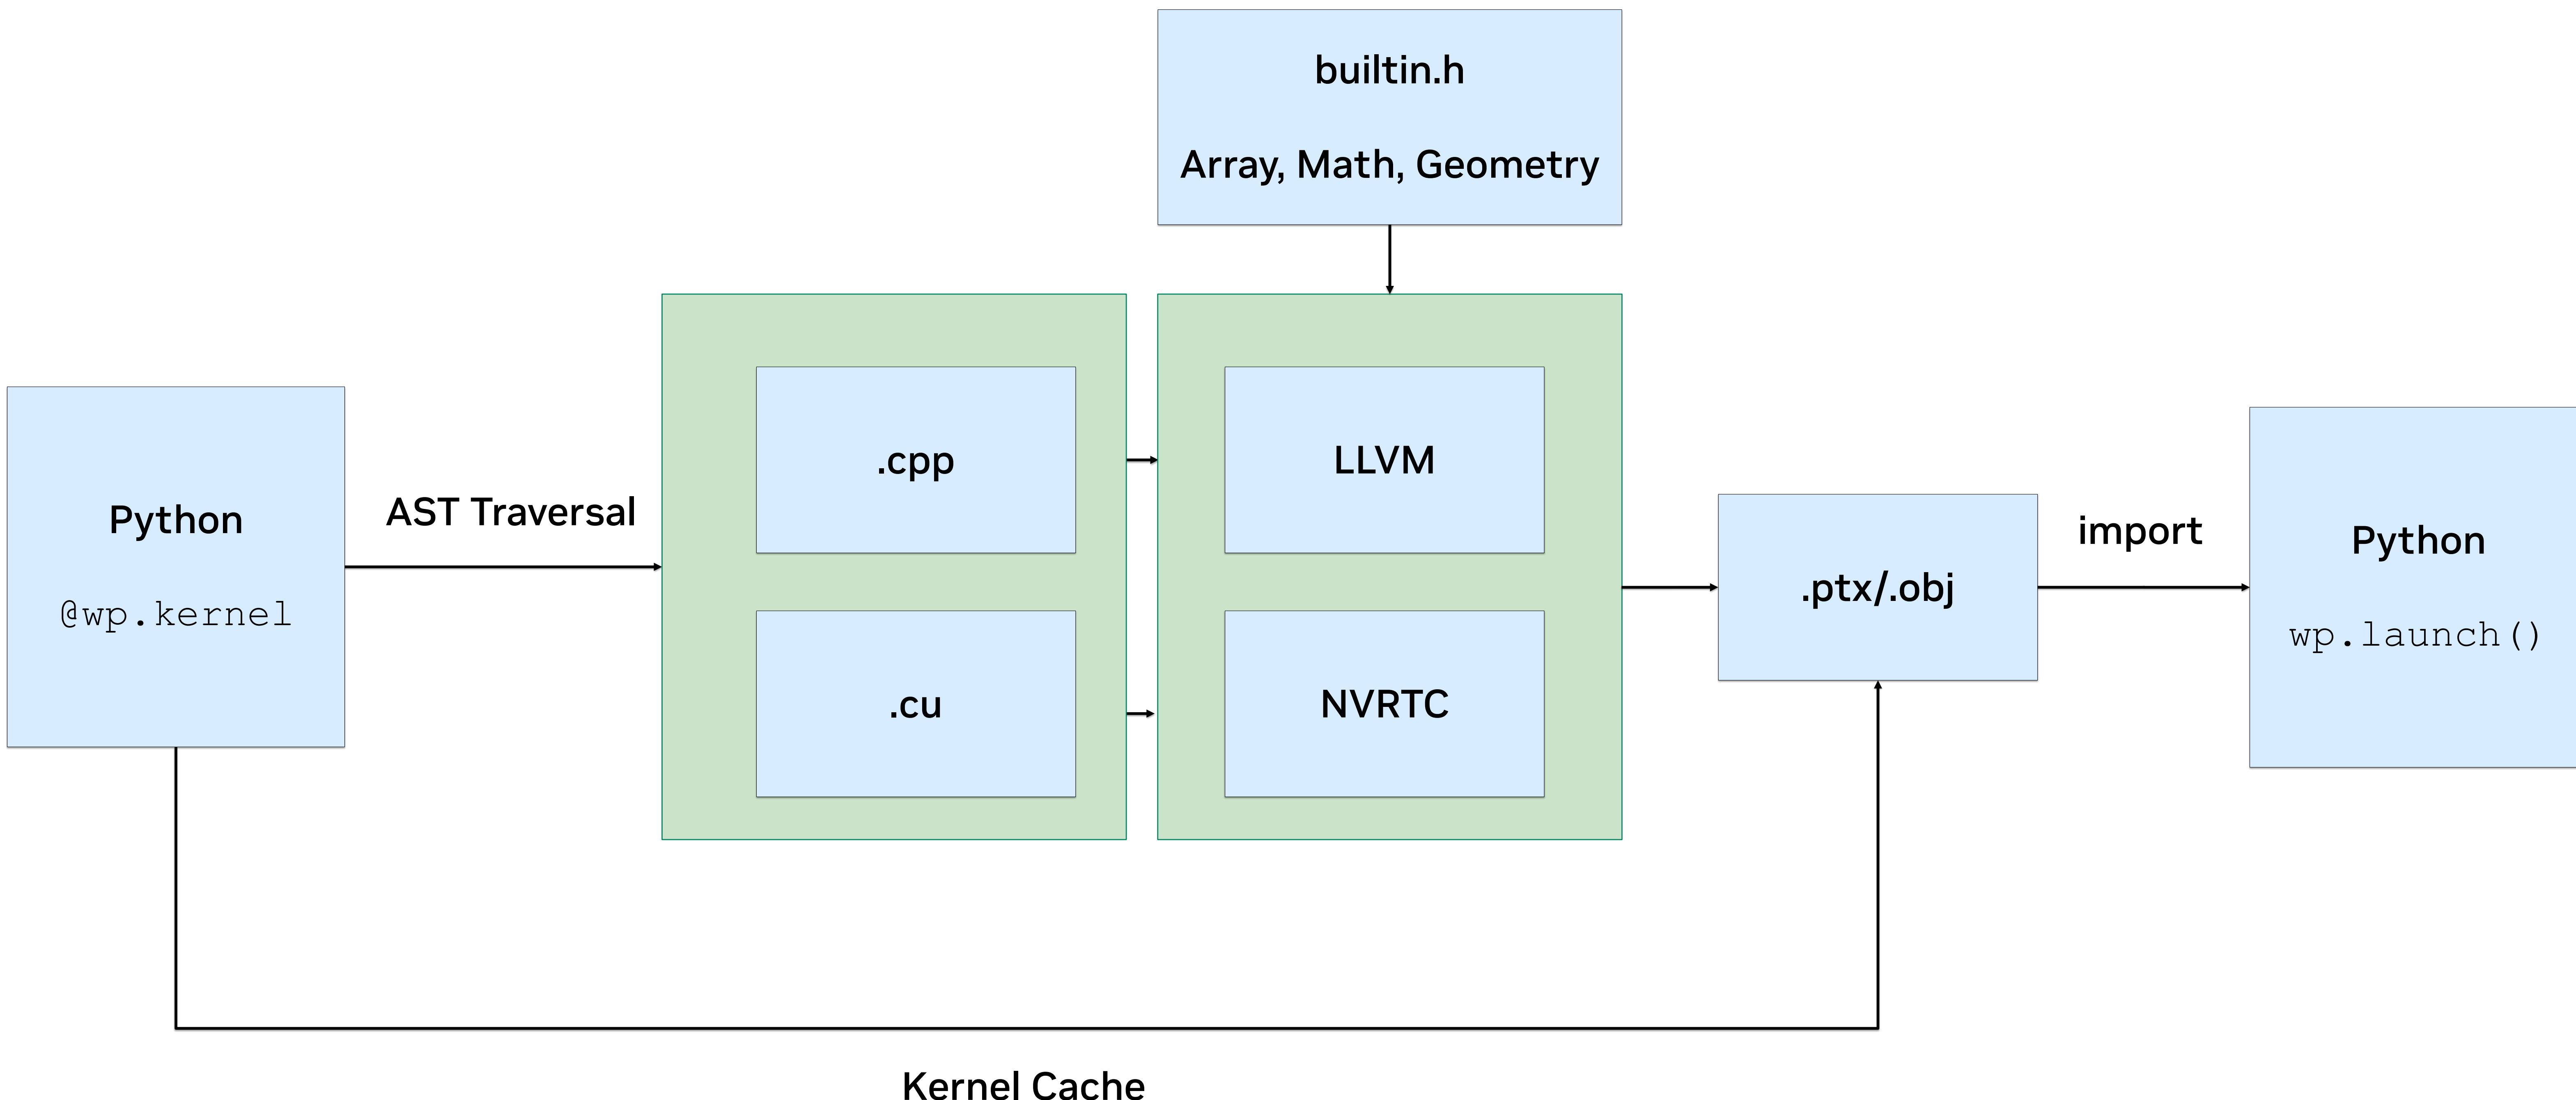
# Warp Data Model

- Host/Device memory managed through `wp.array` type
- Builtin spatial math types similar to OpenCL/HLSL:
  - `vec2`, `vec3`, `vec4`, `mat22`, `mat33`, `mat44`, `quat`, `transform`
- Support for all common array protocols:
  - `__array_interface__`
  - `__cuda_array_interface__`
  - `__dlpack__`
- Zero-copy interop with PyTorch, JAX, NumPy:
  - `wp.from_torch()`, `wp.to_torch()`
  - `wp.from_jax()`, `wp.to_jax()`

```
# 1D array of int8 types:  
a = wp.zeros(shape=[64], dtype=wp.int8)  
  
# 2D array of fp16 vec3 types:  
a = wp.zeros(shape=[64, 64], dtype=wp.vec3h)  
  
# 3D array of fp64 mat44 types:  
a = wp.zeros(shape=[64, 64, 64], dtype=wp.mat44d)
```



# Warp Compilation Pipeline



# Warp Execution Model

- Warp exposes a thin abstraction over CUDA
- Kernels are launched over an **N-dimensional grid of threads**
  - Max 4-dimensional grids
  - Mapping to blocks is handled internally and not exposed to user
  - Grid stride kernels used to scale to large thread counts
    - Max up to  $2^{31}-1$  on each dimension
- Pure SIMT model
  - No shared memory
  - No warp-level primitives, e.g.: `__shfl_sync()`
- Custom native functions
  - C++/CUDA snippets

```
@wp.kernel
def divergence(u: wp.array2d(dtype=wp.vec2),
               div: wp.array2d(dtype=float)):

    # 2D thread indices
    i, j = wp.tid()

    # boundary conditions
    if i == grid_width - 1:
        return
    if j == grid_height - 1:
        return

    # compute divergence
    dx = (u[i + 1, j][0] - u[i, j][0]) * 0.5
    dy = (u[i, j + 1][1] - u[i, j][1]) * 0.5
    div[i, j] = dx + dy

# 2d kernel launch
wp.launch(divergence, dim=[512, 512], inputs=[u, div])
```

Example: 2D Divergence Calculation in Warp

# Custom CUDA Code

## Inserting native code snippets

- Add custom CUDA code snippets to Warp directly from Python via `@wp.func_native` decorator
- Allows easily dropping down to CUDA for access to:
  - Shared memory
  - Fine grained synchronization
  - Cooperative operations
- Custom native backward functions
- Ability to include larger C++/CUDA codes (header-only) coming soon

```
snippet = """
out[tid] = a * x[tid] + y[tid];
"""

adj_snippet = """
adj_a = x[tid] * adj_out[tid];
adj_x[tid] = a * adj_out[tid];
adj_y[tid] = adj_out[tid];
"""

@wp.func_native(snippet, adj_snippet)
def saxpy(
    a: wp.float32,
    x: wp.array(dtype=wp.float32),
    y: wp.array(dtype=wp.float32),
    out: wp.array(dtype=wp.float32),
    tid: int,
):
    ...
    ...
```

Example: Custom CUDA snippet as a Warp function

# Automatic Differentiation

## Differentiable Programming in Warp

- Reverse-mode automatic differentiation
- Kernel adjoint codes are generated automatically
- Store and reply kernels using `wp.Tape()`
- User-provided custom gradients
- Limitations:
  - First-order derivatives only
  - Dynamic loops and side-effects
- Capture entire backward pass in a CUDA graph

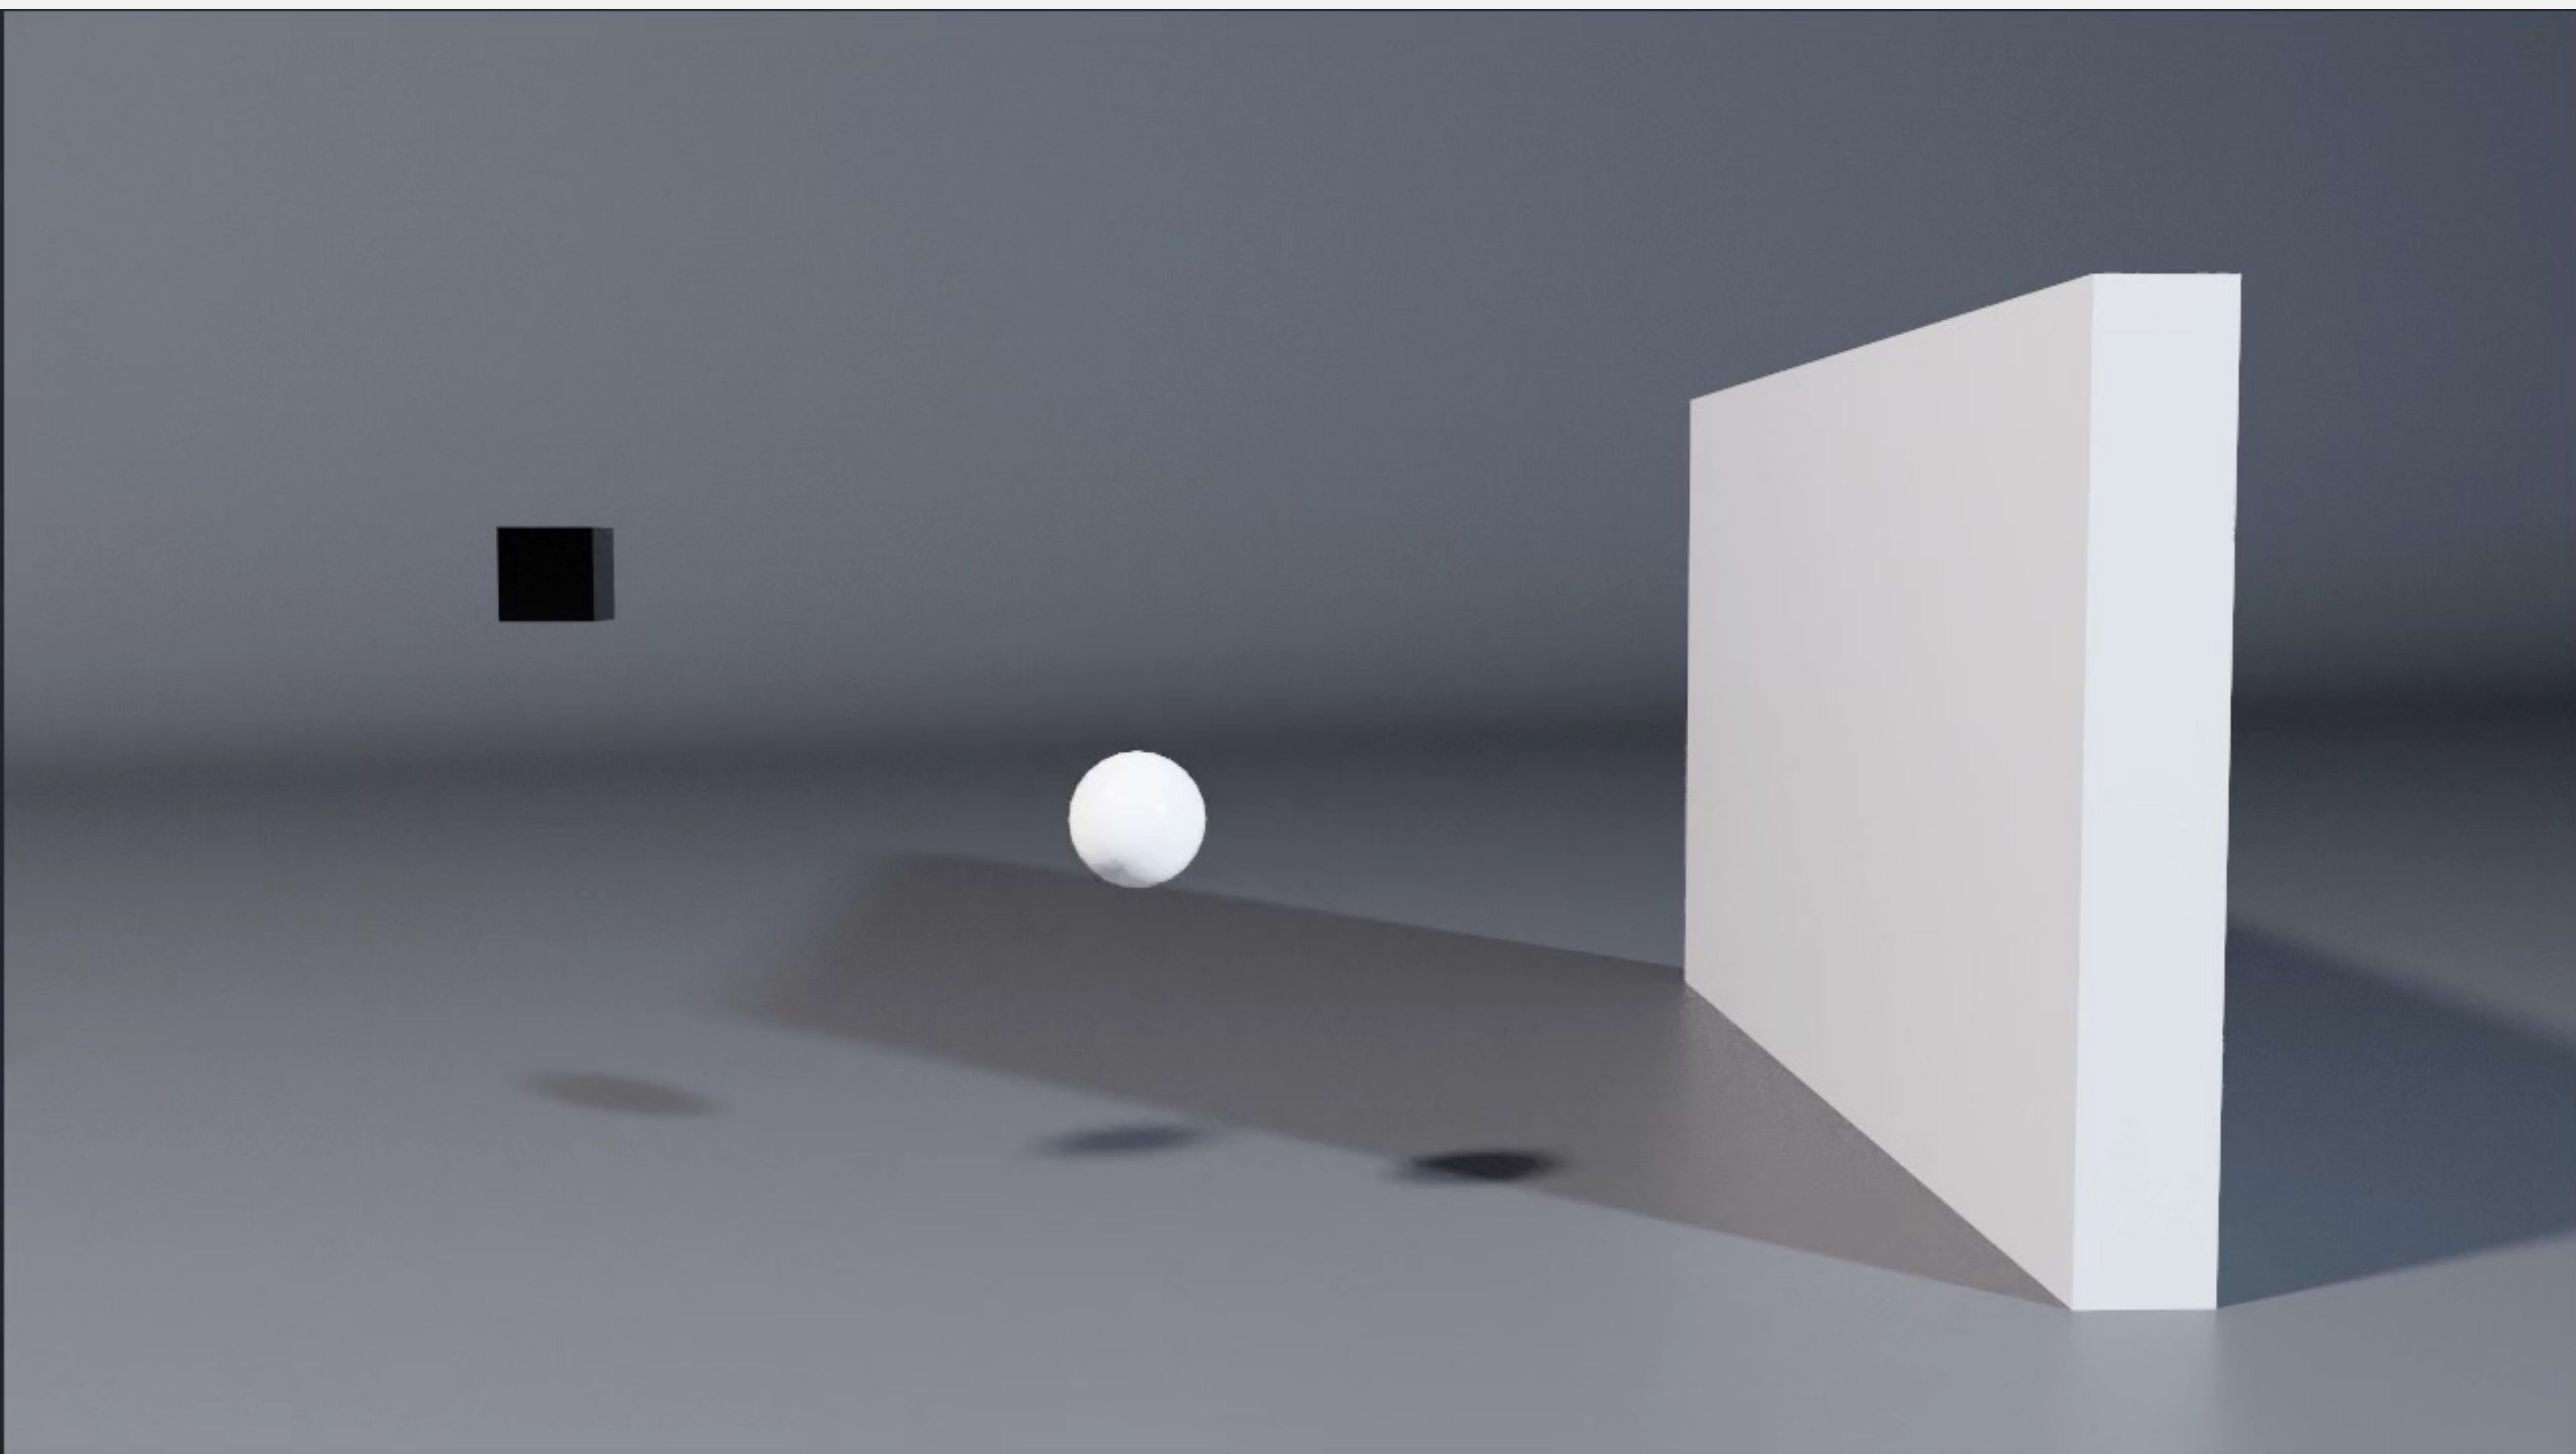
```
tape = wp.Tape()

with tape:
    wp.launch(kernel=kernel_1, dim=dim, inputs=[a], outputs=[b])
    wp.launch(kernel=kernel_2, dim=dim, inputs=[b], outputs=[c])
    wp.launch(kernel=kernel_3, dim=dim, inputs=[c], outputs=[loss])

    # run backward
    tape.backward(loss)

    # gradients of loss w.r.t inputs
    print(a.grad)
    print(b.grad)
    print(c.grad)
```

Example: Auto-differentiation through multiple kernel launches



# Mesh Data Structure

## Accelerated Mesh Queries

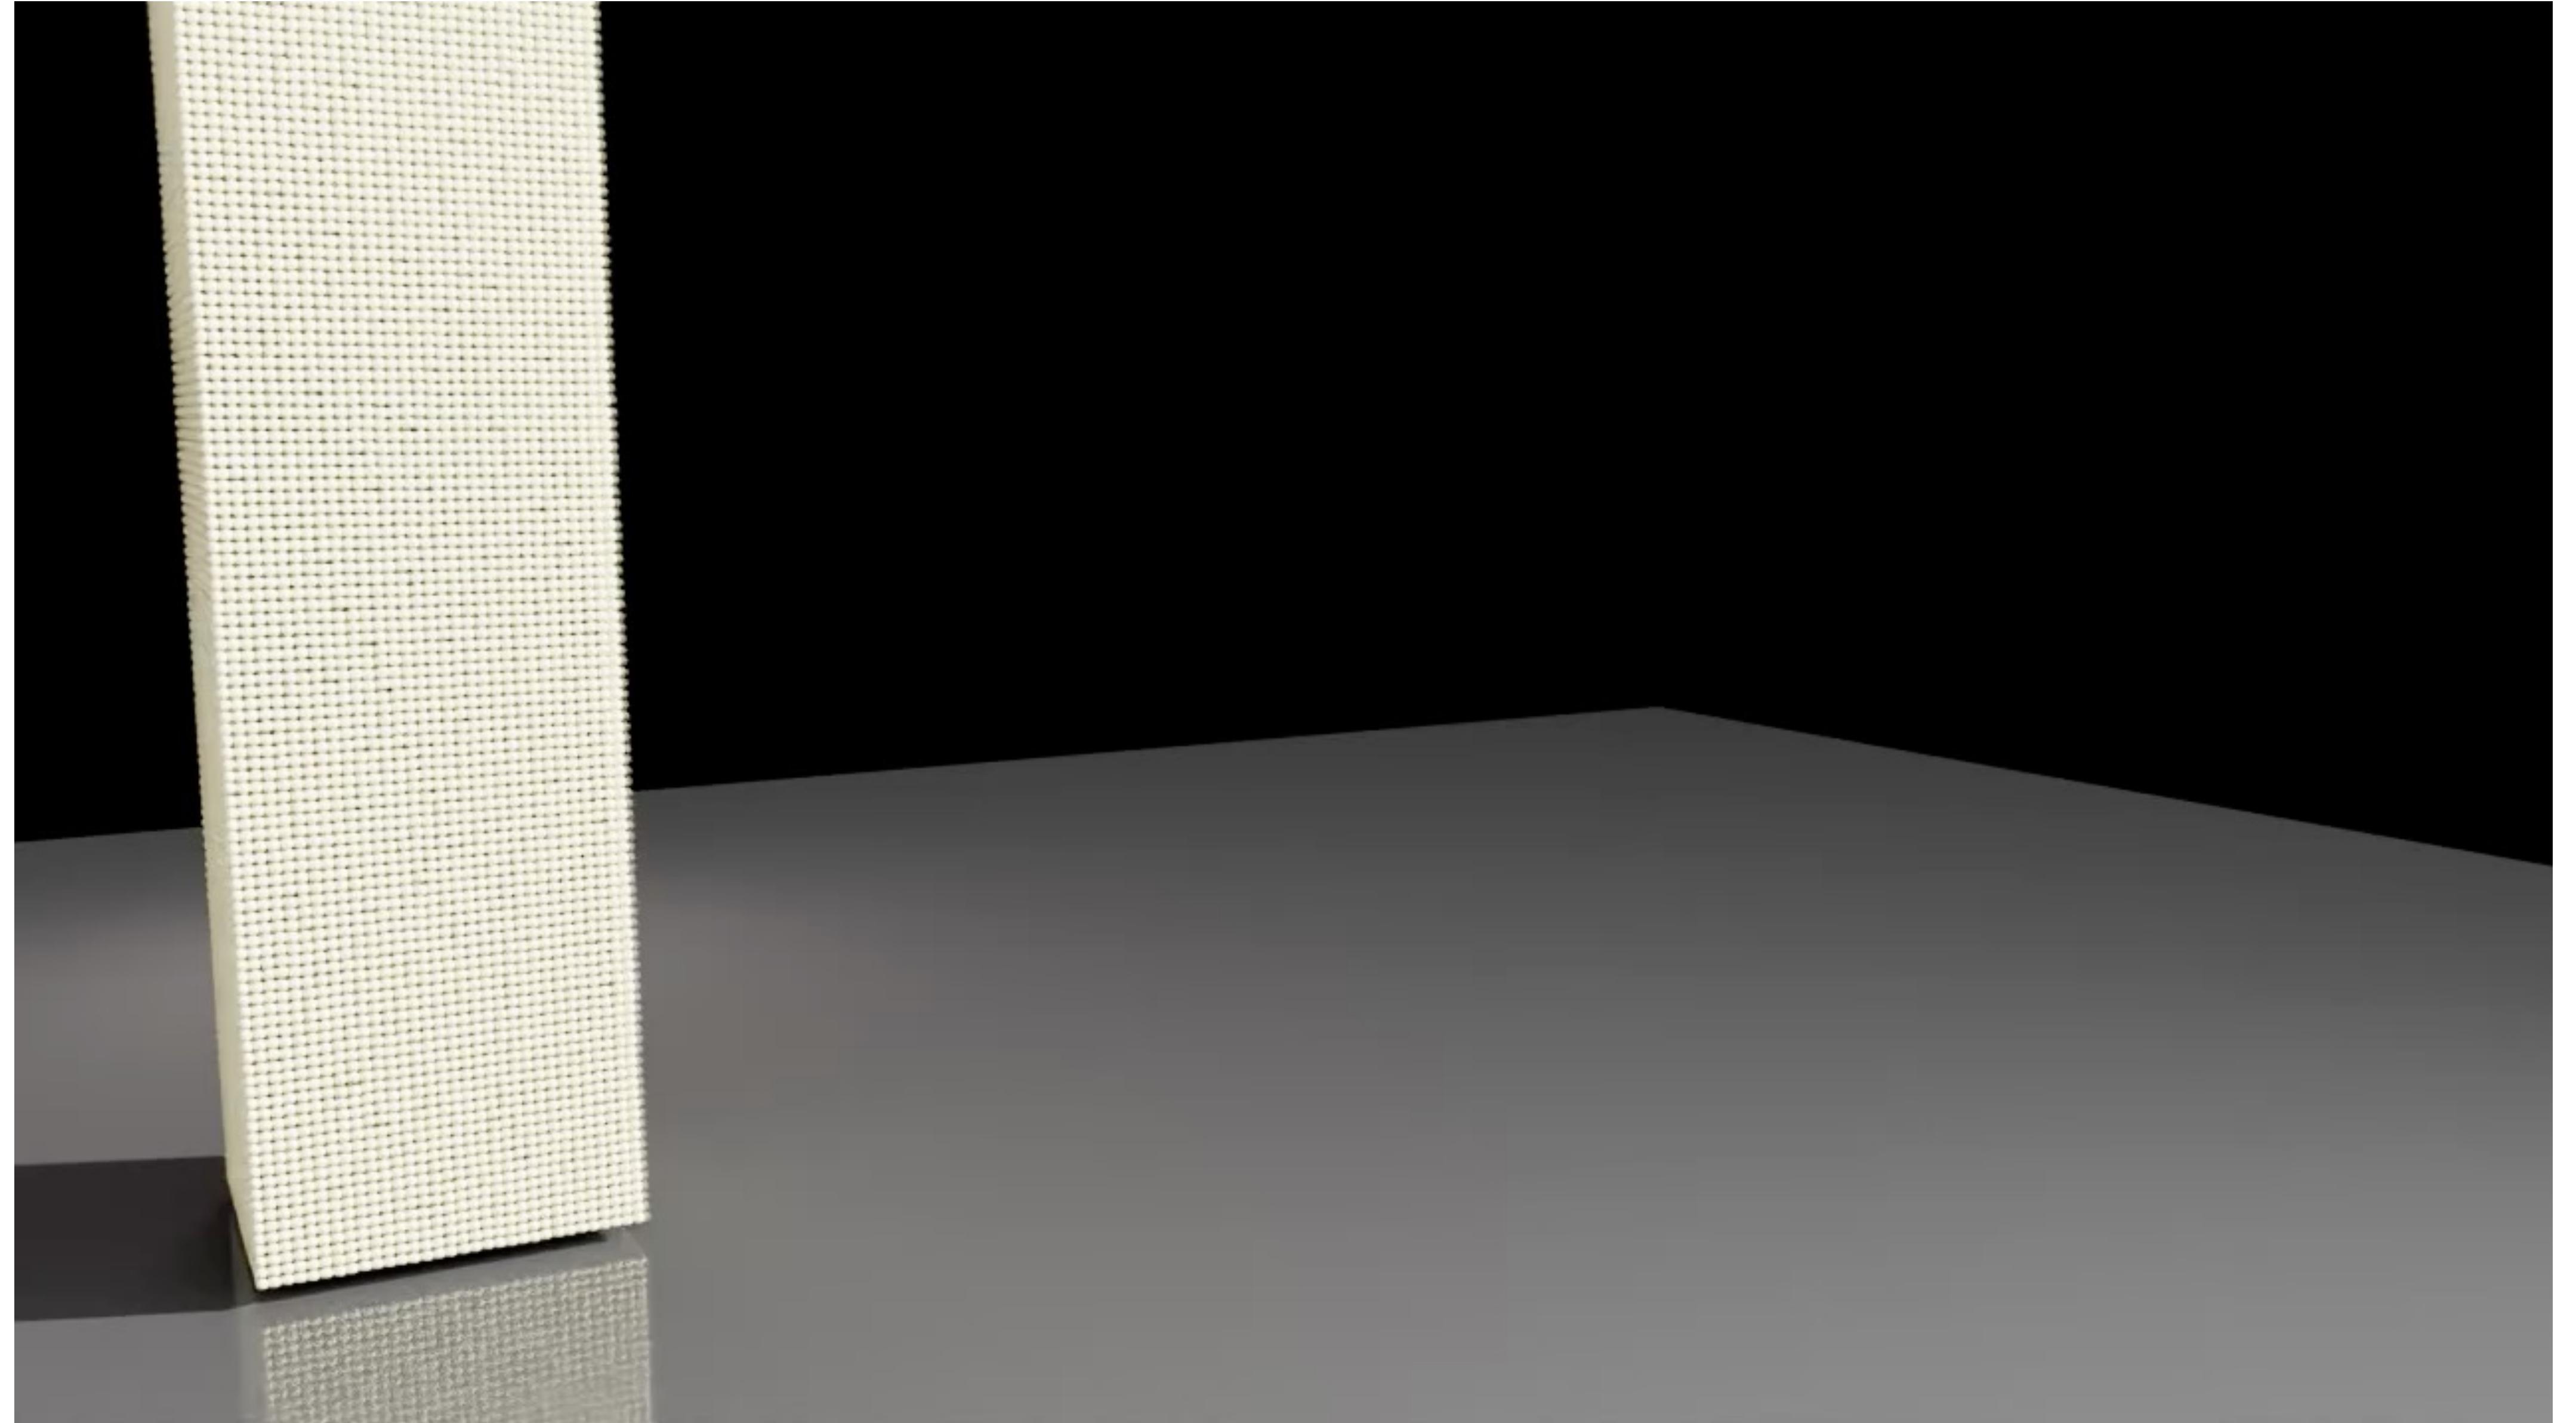
- Warp mesh data structure:
  - `wp.Mesh`
- Built in mesh queries:
  - `wp.mesh_query_point()`
  - `wp.mesh_query_ray()`
  - `wp.mesh_query_aabb()`
- Fast inside/outside (sign) determination
- Fast GPU bounding volume hierarchy (BVH) refits:
  - `mesh.refit()`
- Example:
  - Neo-Hookean FEM elastic model
  - 60k particles versus 32k tris in 0.25ms
  - Dynamic signed-distance field (SDF) contact against meshes



# Hash Grid Data Structure

## Nearest Neighbor Point Queries

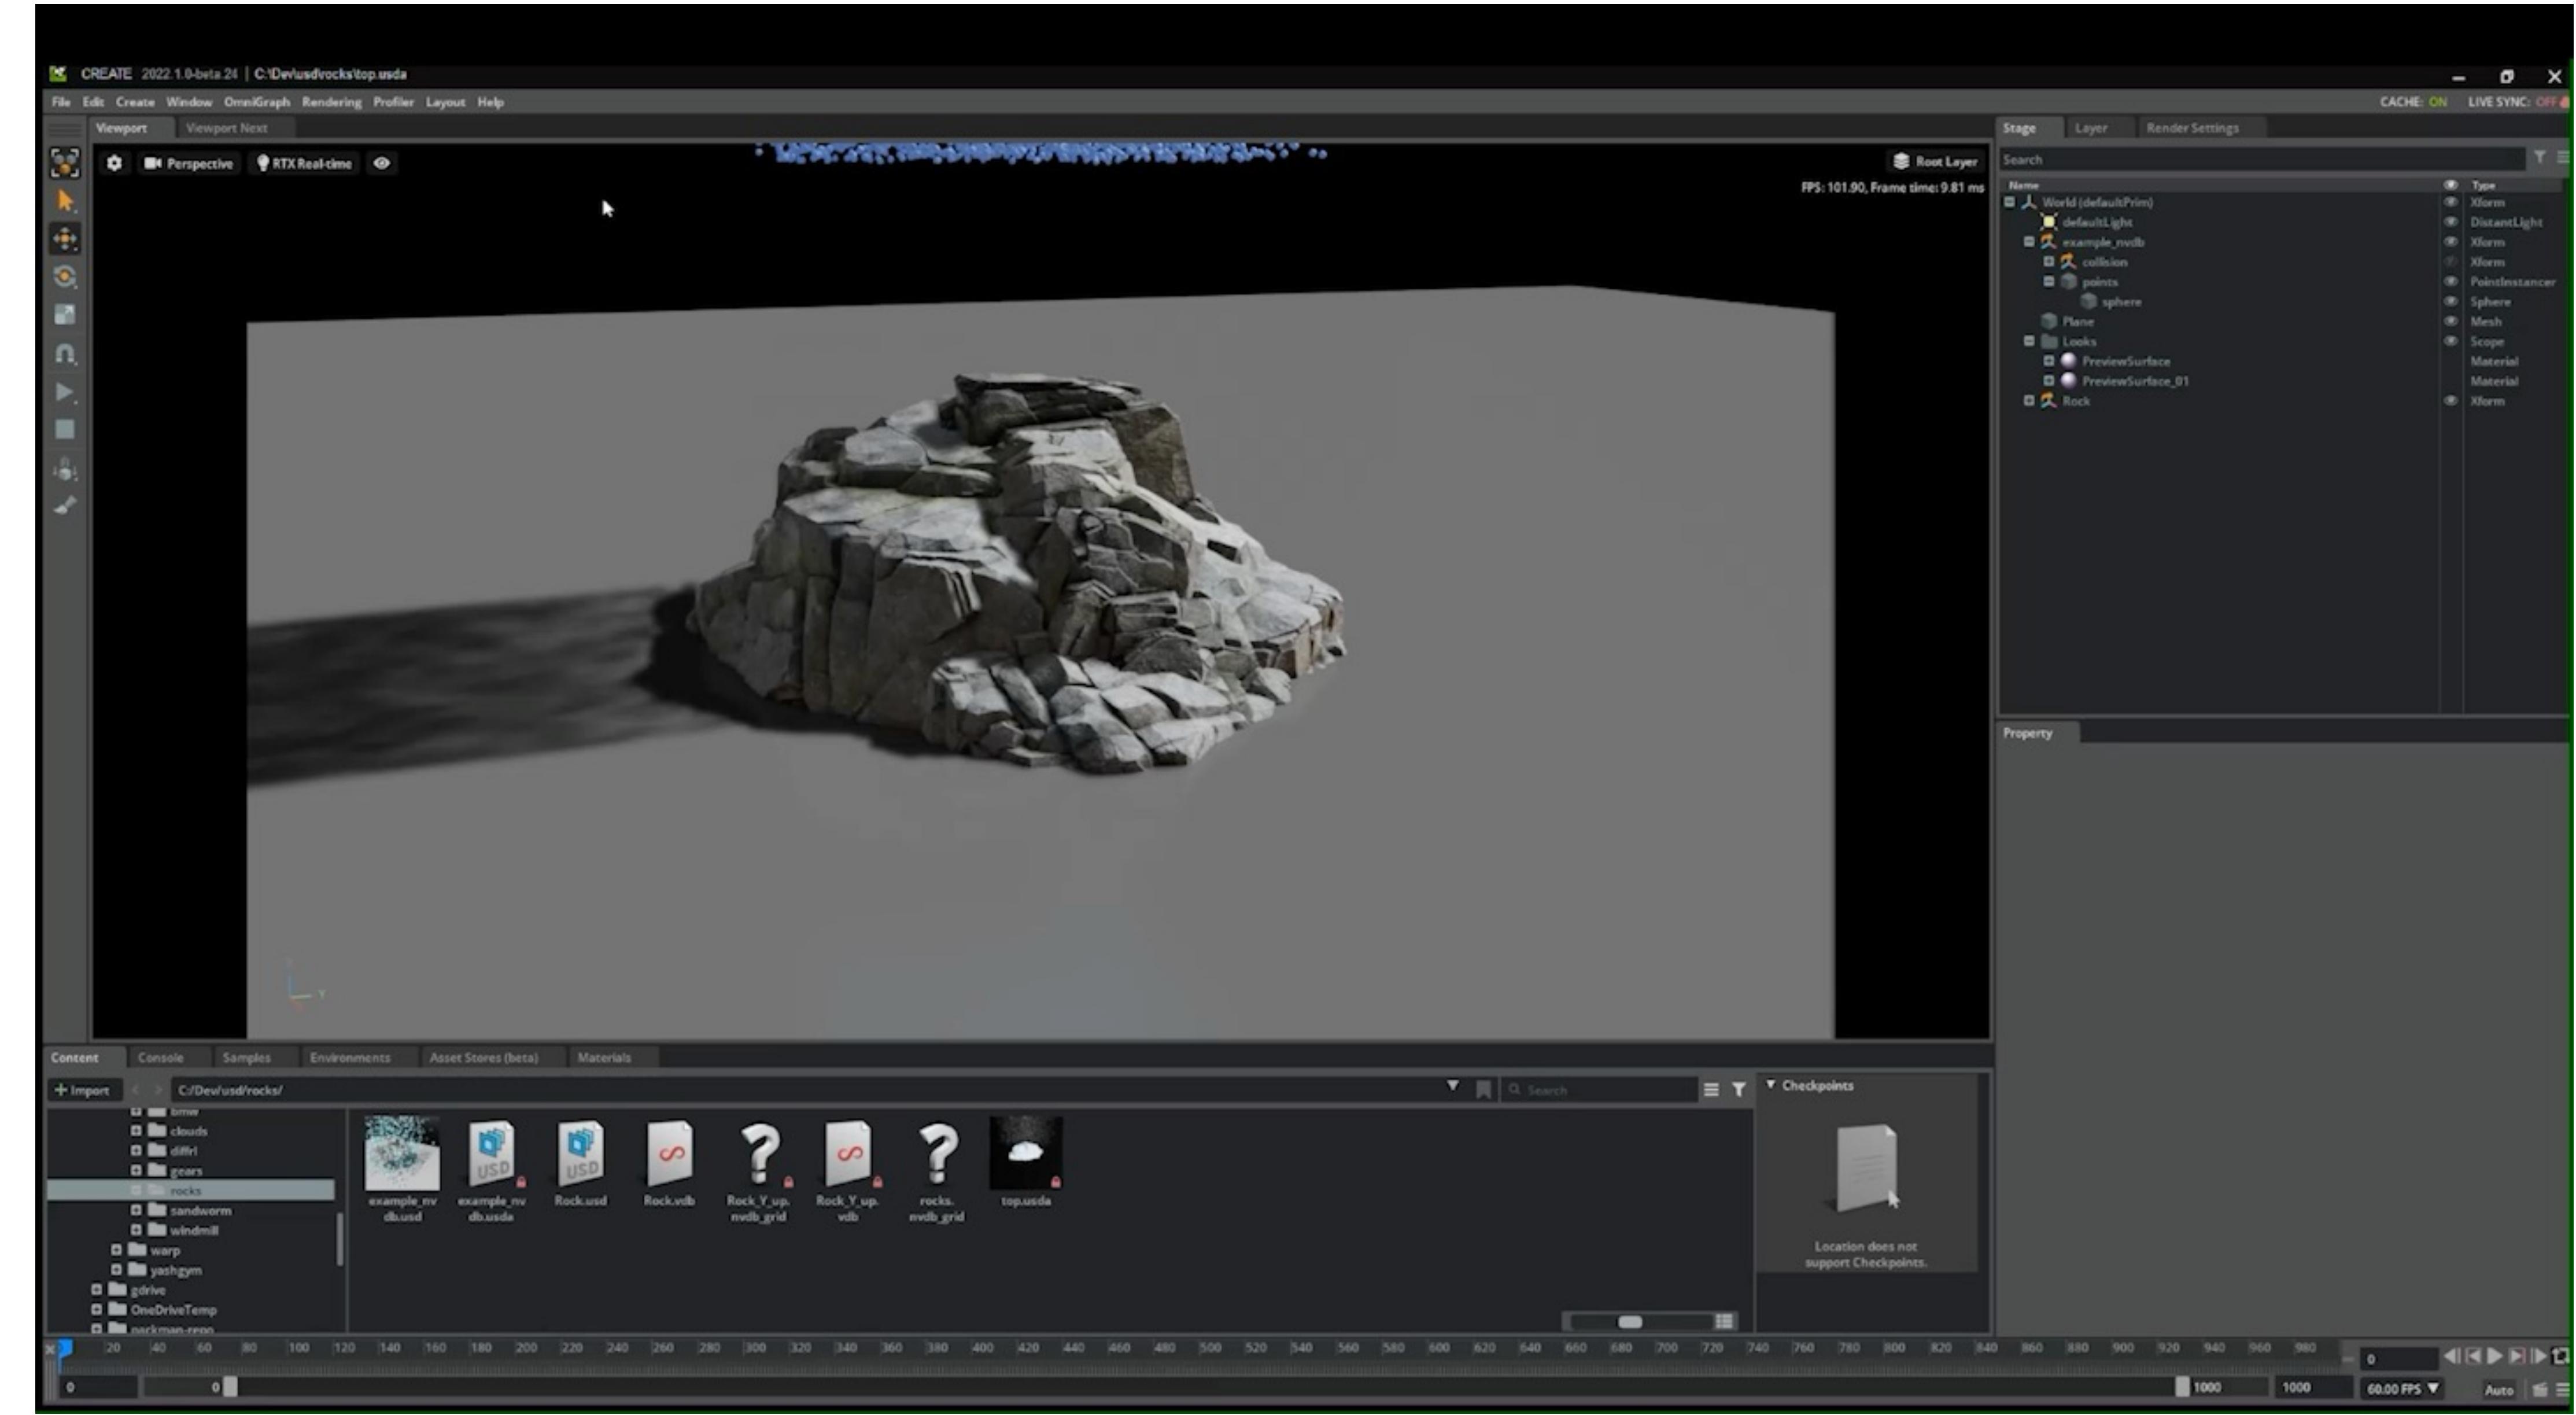
- Built-in hash-grid primitive
  - `wp.HashGrid`
- Fast neighbor finding / radius queries:
  - `wp.hash_grid_query_point()`
  - `wp.hash_grid_query_next()`
- GPU-side rebuild:
  - `hashgrid.build()`
- 10x faster than Open3D
- Example:
  - Discrete Element Method (DEM) in 200 lines of Python
  - 128k particles ~10ms/frame



# Sparse Grid Data Structure

## OpenVDB

- Support for NanoVDB sparse volumes
  - `wp.Volume`
- Great for narrow-band SDF collision
- Simple to use API:
  - `wp.volume_sample_world(vol, xyz)`
  - `wp.volume_sample_local(vol, uvw)`
  - `wp.volume_lookup(vol, ijk)`
  - `wp.volume_transform(vol, xyz)`
  - `wp.volume_transform_inv(vol, xyz)`



# Warp Sim

## Differentiable Simulation for Robotics

- A framework for real-time, differentiable, robotic simulation
- Physical Models
  - Rigid bodies
  - Particles
  - Constraints
  - Geometry
  - Forces
- Physical State
  - Position ( $q$ )
  - Velocity ( $qd$ )
- Integrators
  - Symplectic Euler (semi-implicit)
  - XPBD (implicit)
  - Featherstone



```
import warp as wp
import warp.sim

sim_dt = 1/60

builder = wp.sim.ModelBuilder()
wp.sim.parse_urdf("cartpole.urdf", builder)
model = builder.finalize()

integrator = wp.sim.SemiImplicitIntegrator()
state, next_state = model.state(), model.state()

for i in range(100):
    state.clear_forces()
    state = integrator.simulate(model, state, next_state, sim_dt)
    state, next_state = next_state, state
```

Example: Creating a cartpole simulation using warp.sim

# Warp FEM

## Differentiable Framework for PDEs

- Flexible finite element-based (FEM/DG) framework for
  - Diffusion
  - Convection
  - Fluid flow
  - Elasticity
- Define your weak-form PDE in high-level Python syntax:
  - `wp.fem.grad()`
  - `wp.fem.div()`
  - `wp.fem.curl()`
- Combine with geometry + shape functions:
  - `wp.fem.TetMesh()`
  - `wp.fem.HexMesh()`
  - `wp.fem.make_polynomial_space()`
- Integrate to assemble linear system:
  - `wp.fem.integrate()`
- Solve with block-sparse (BSR) solvers:
  - `warp.sparse.BsrMatrix`

### Define PDE

```
import warp.fem as fem

@fem.integrand
def diffusion_form(s: fem.Sample, u: fem.Field, v: fem.Field, nu: float):
    return nu * wp.dot(
        fem.grad(u, s),
        fem.grad(v, s),
    )
```

### Define Geometry

```
geo = fem.Tetmesh(positions=positions, tet_vertex_indices=indices)
scalar_space = fem.make_polynomial_space(
    geo, degree=2, element_basis=fem.ElementBasis.SERENDIPITY)
```

### Assemble System

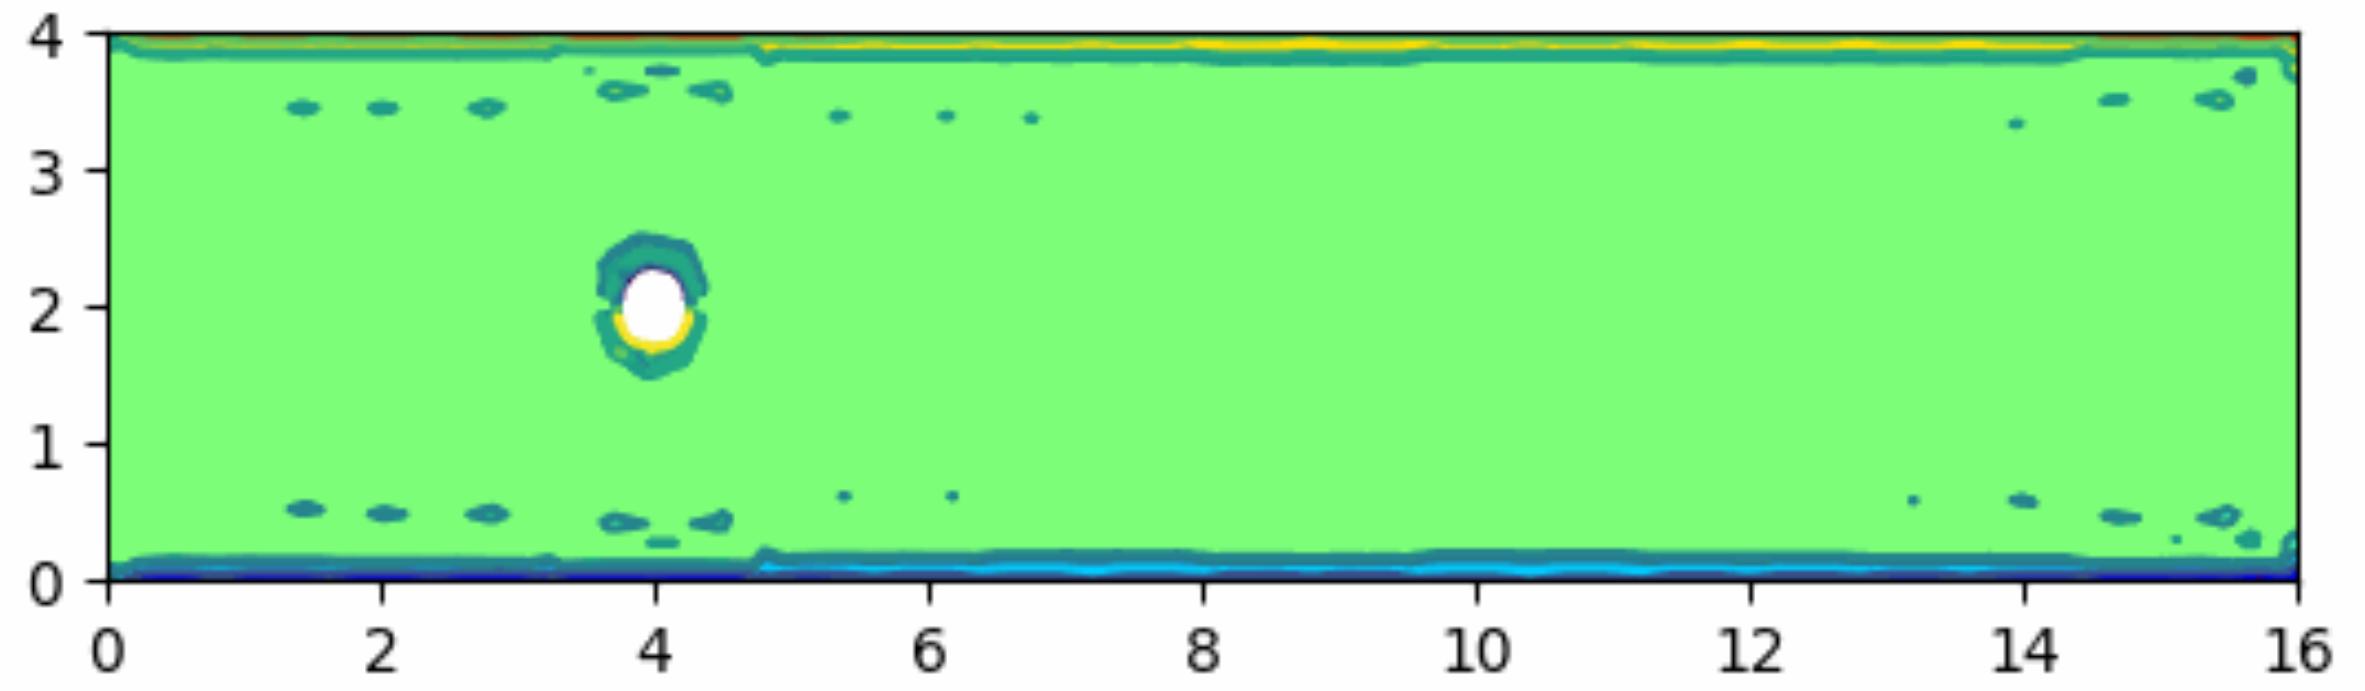
```
trial = fem.make_trial(scalar_space)
test = fem.make_test(scalar_space)

matrix = fem.integrate(diffusion_form,
    fields={"u": trial, "v": test},
    values={"nu": 0.1})
```

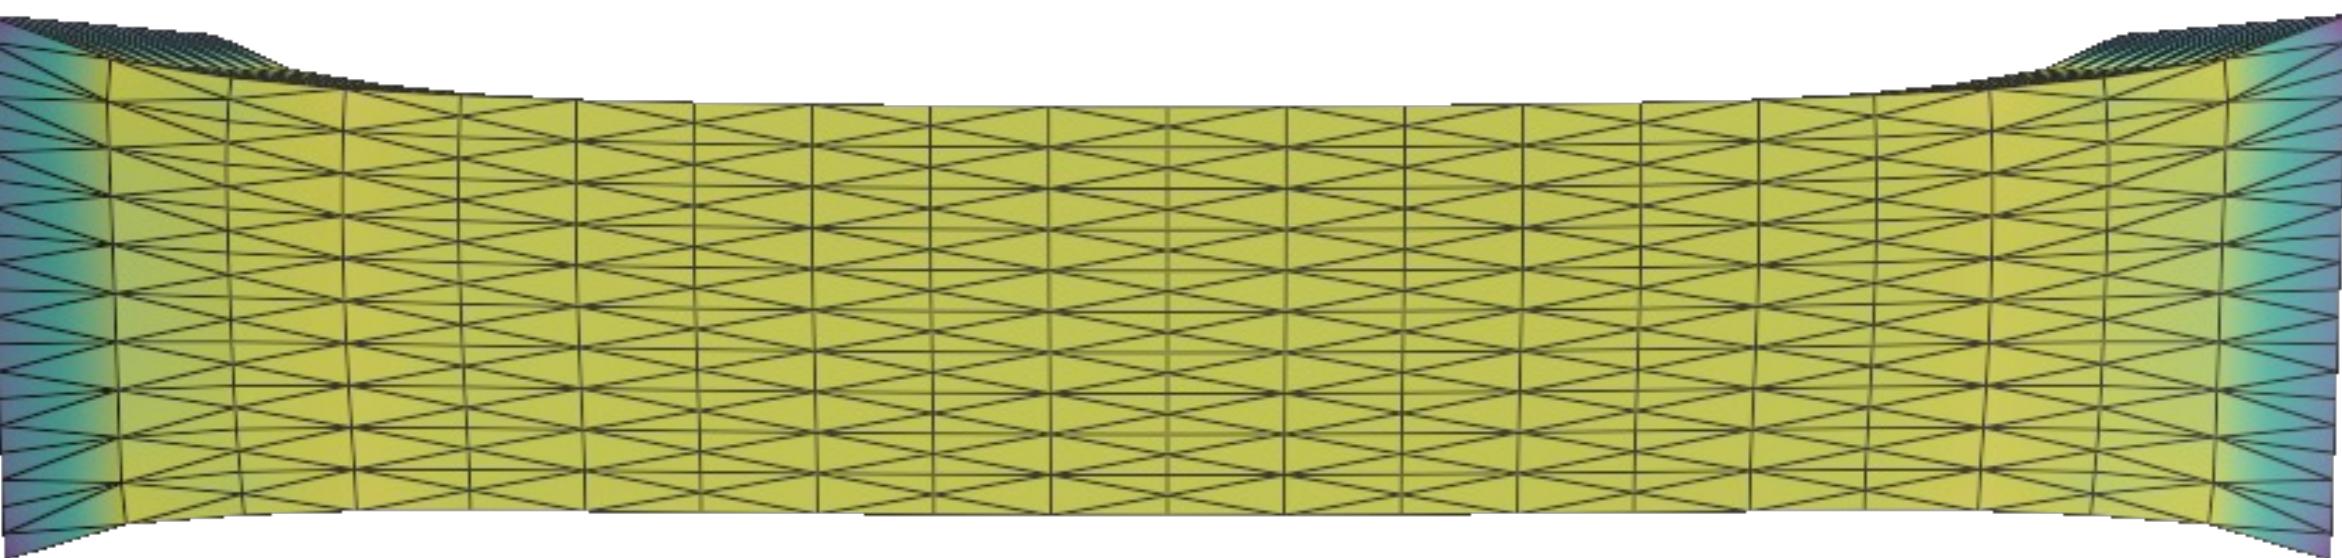
Example: Diffusion PDE bilinear form

# Warp FEM

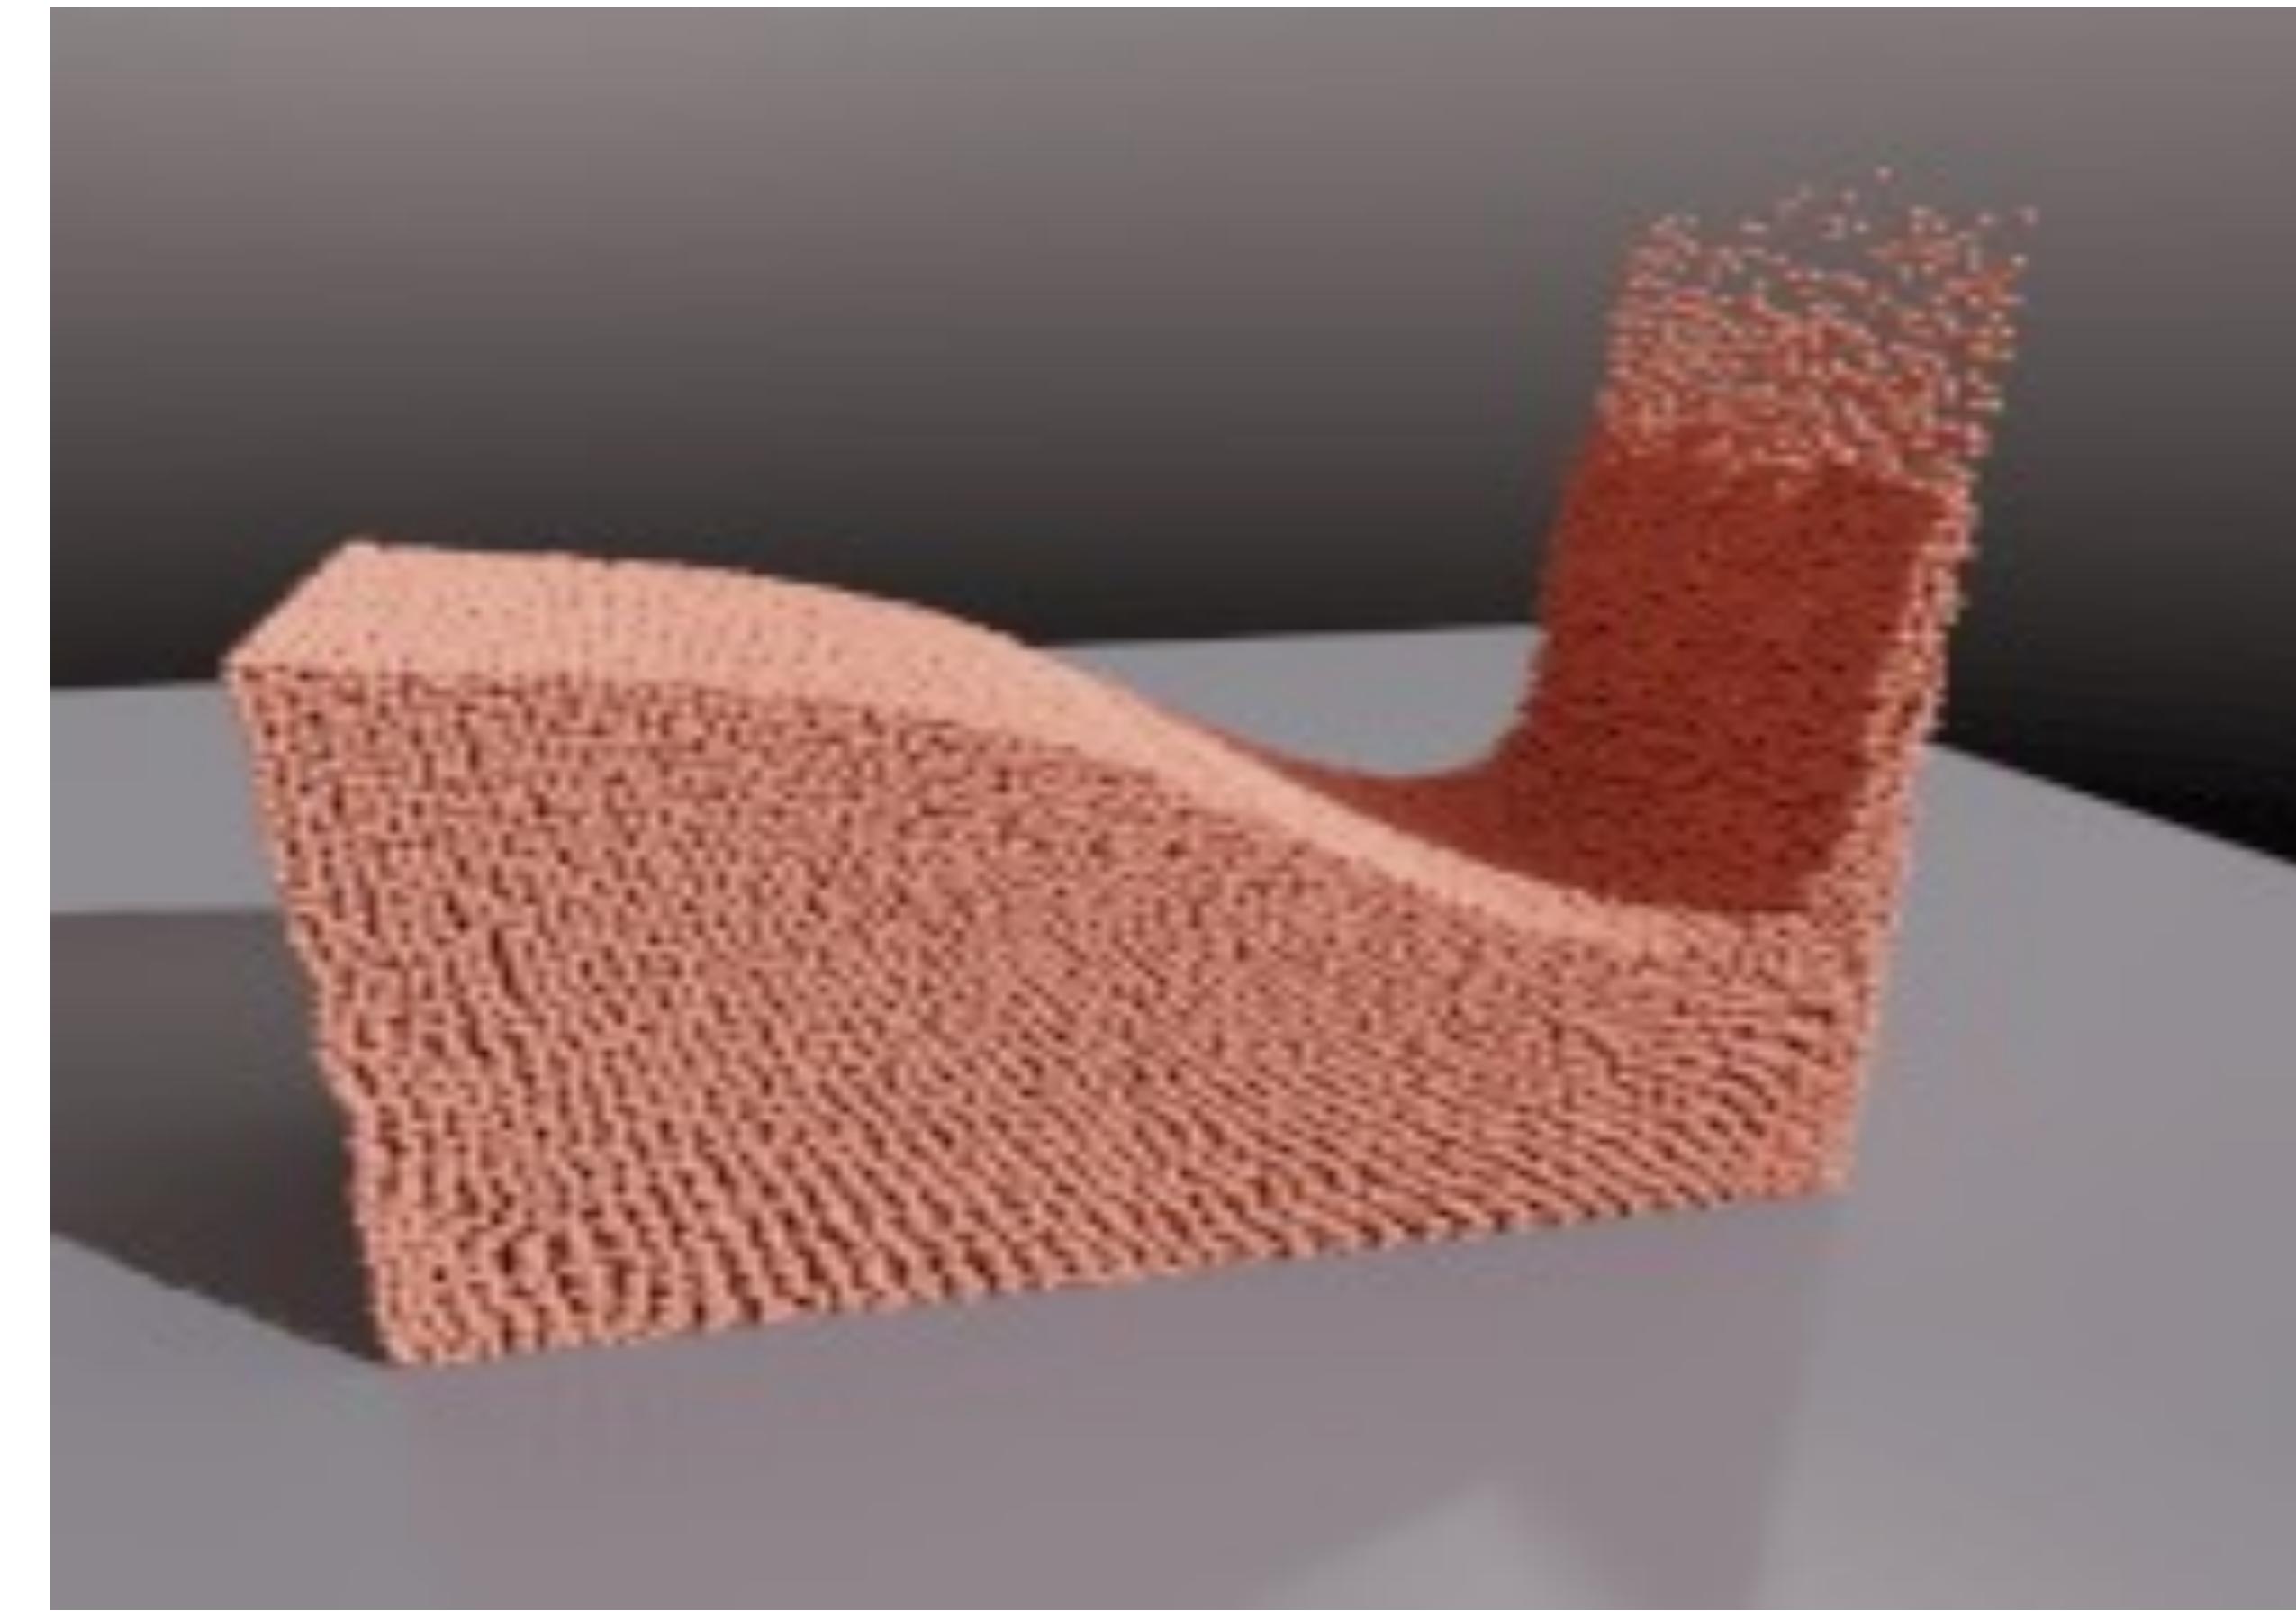
## Examples



Example: Navier—Stokes flow simulation



Example: Neo-Hookean Elasticity



Example: Hybrid Eulerian—Lagrangian fluids

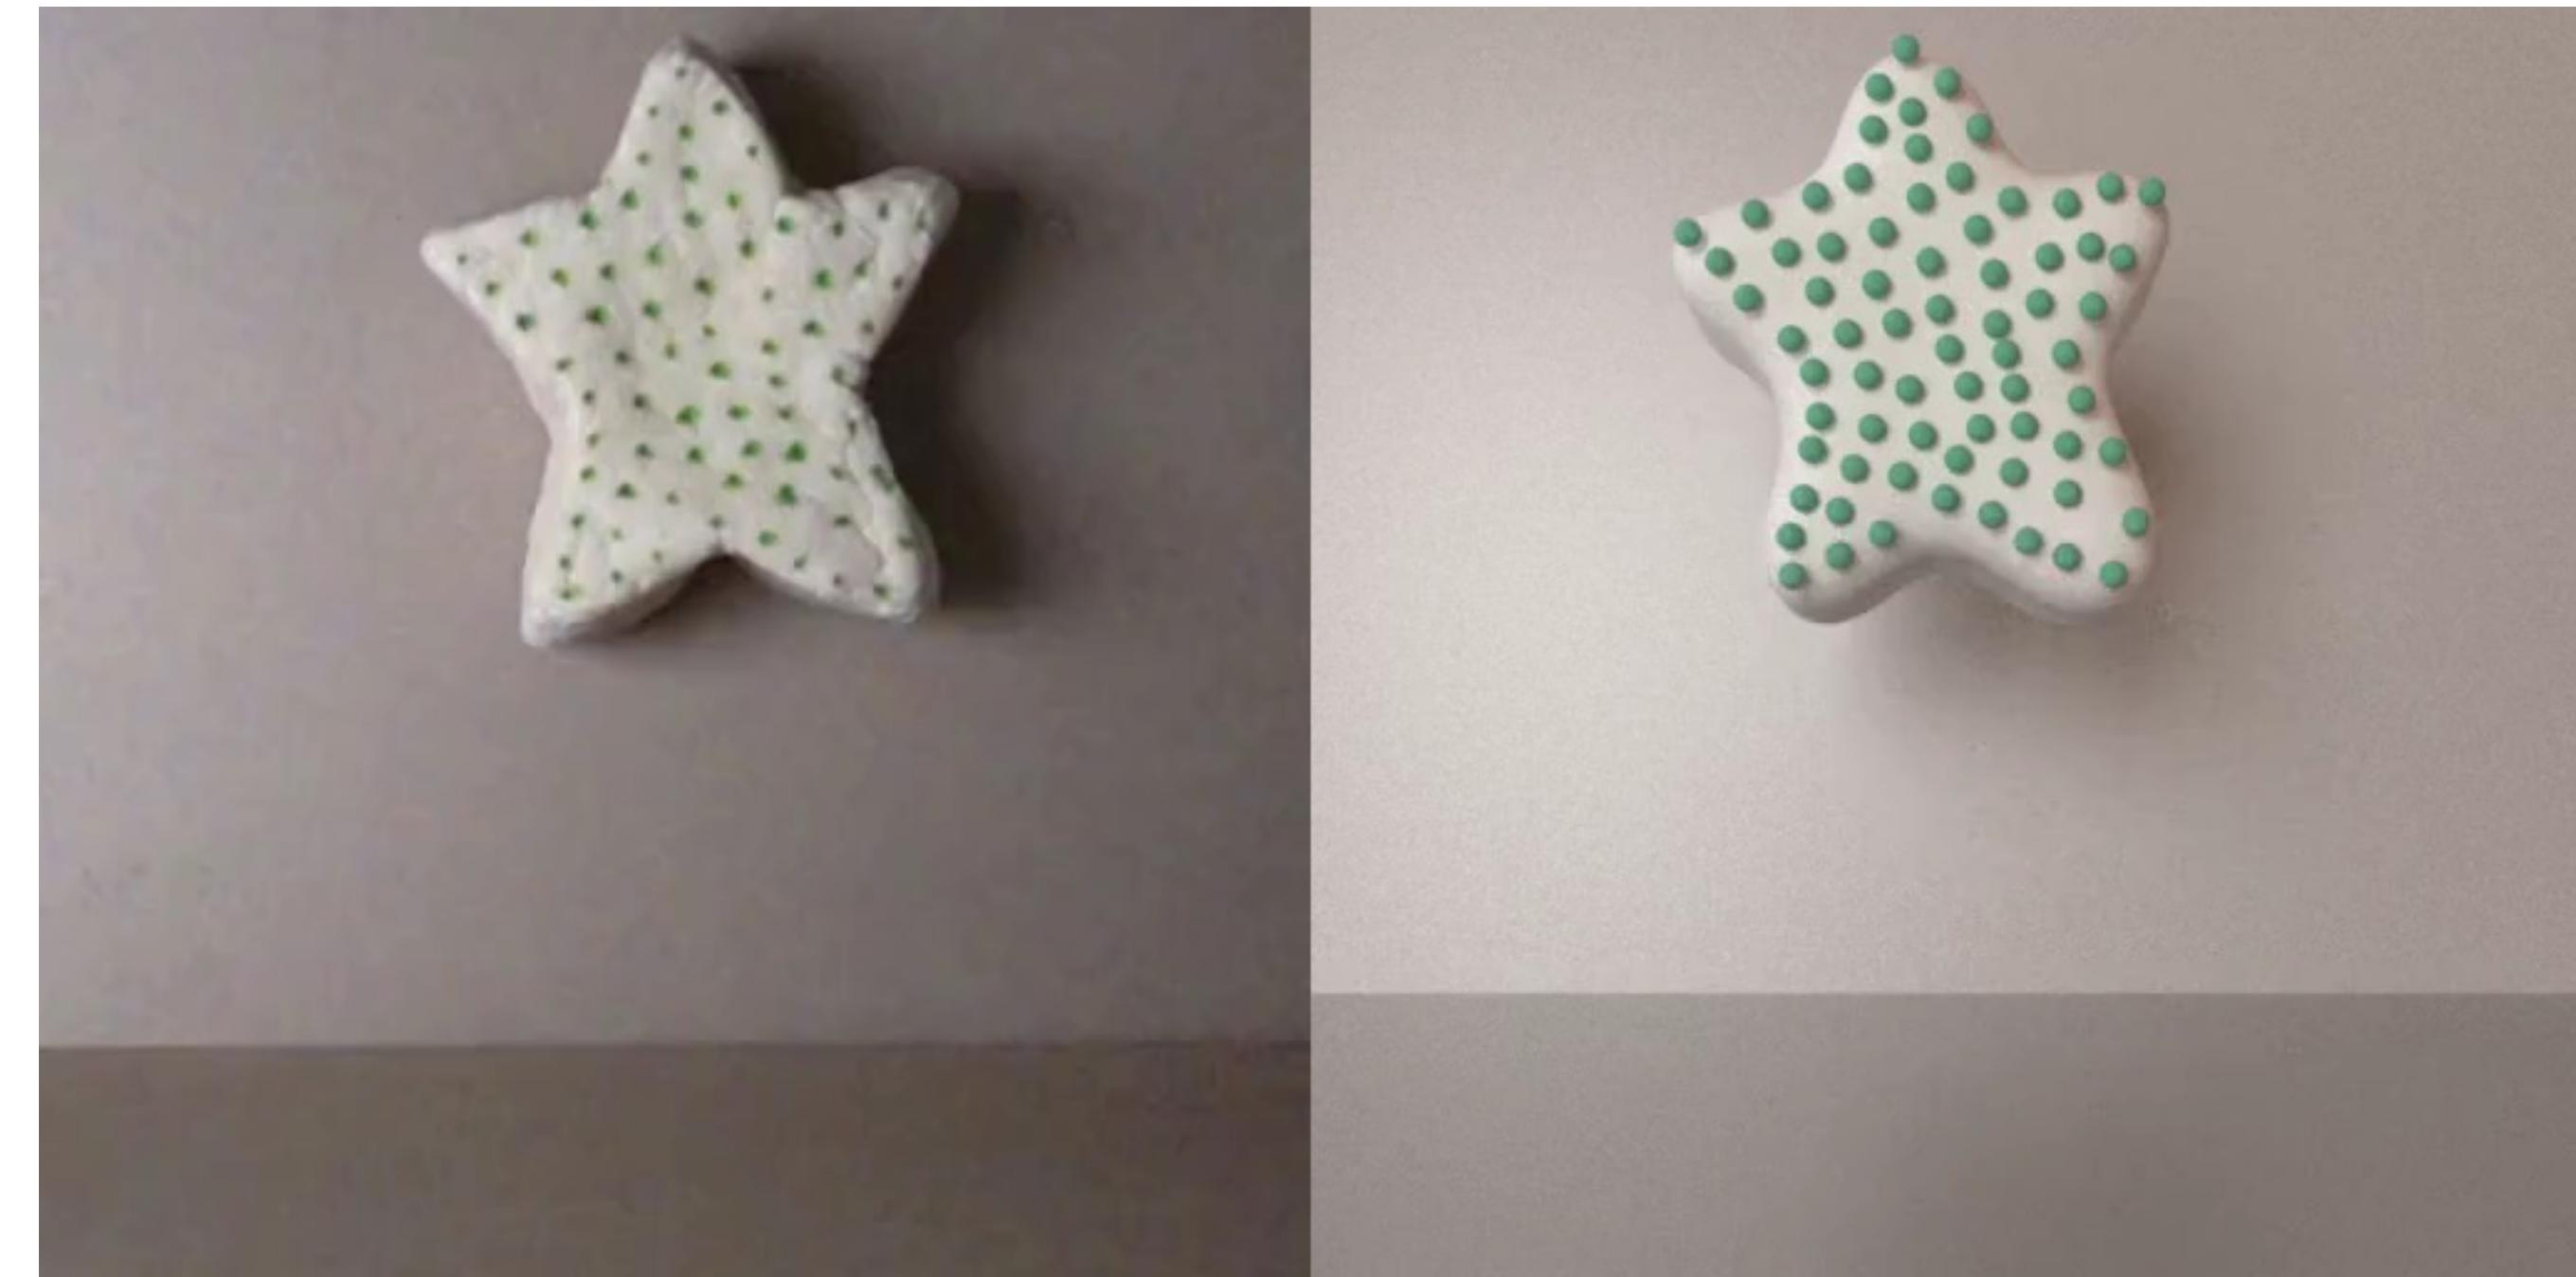


# Case Studies

# Neural Constitutive Laws

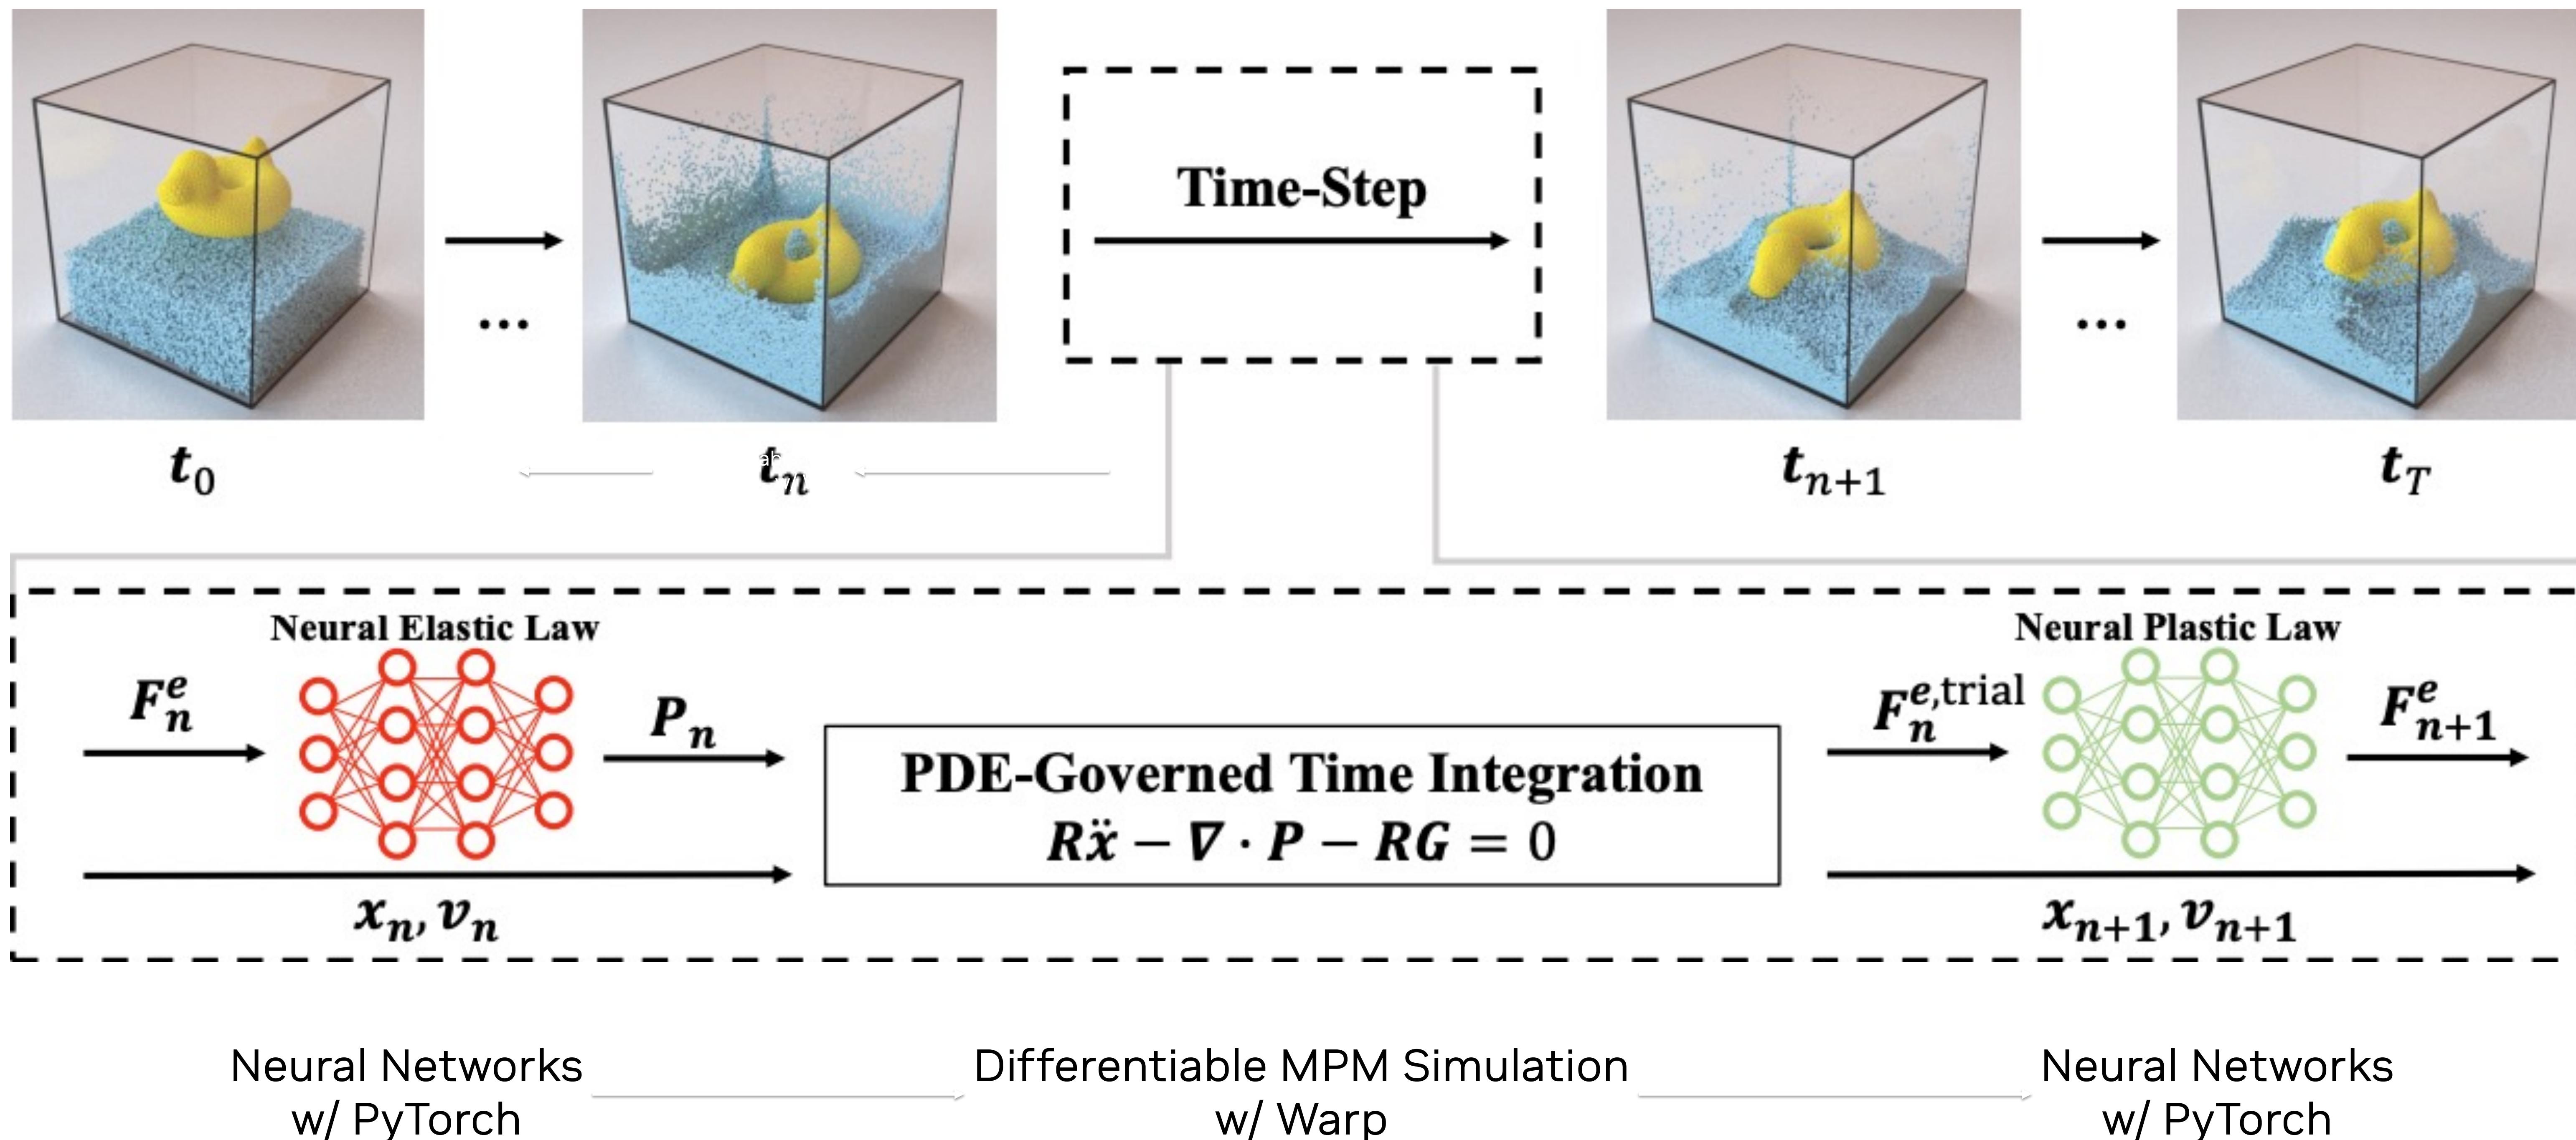
## Learned Elastic Modeling

- Goal: Learn constitutive material models from video
- Research from MIT CSAIL @ ICML 2023
- Differentiable MPM solver written in Warp
- Use to train DNN constitutive model to match captured deformation
- Generalizes to 30x larger environments, unseen geometries and trajectories
- Handles multi-physics environments: coupled rigid body–fluid simulations, phase transitions
- <https://github.com/PingchuanMa/NCLaw>
- [Ma et al. 2023]



Learning Neural Constitutive Laws

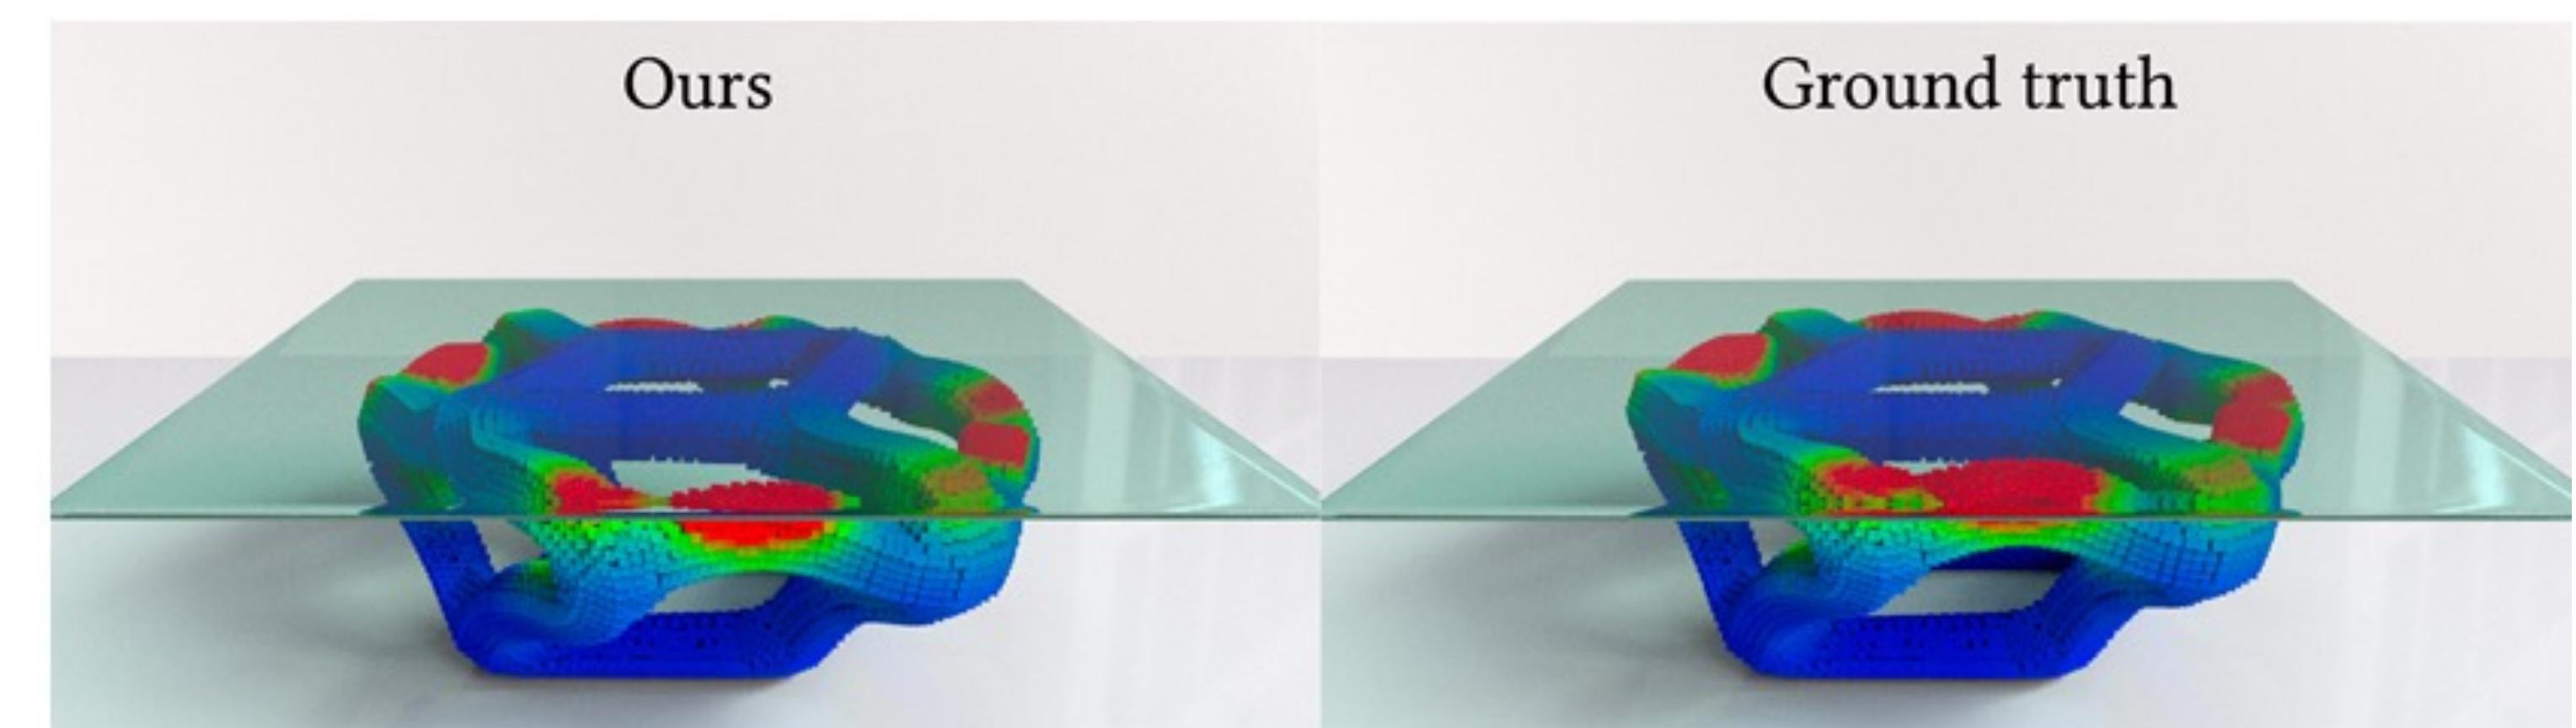
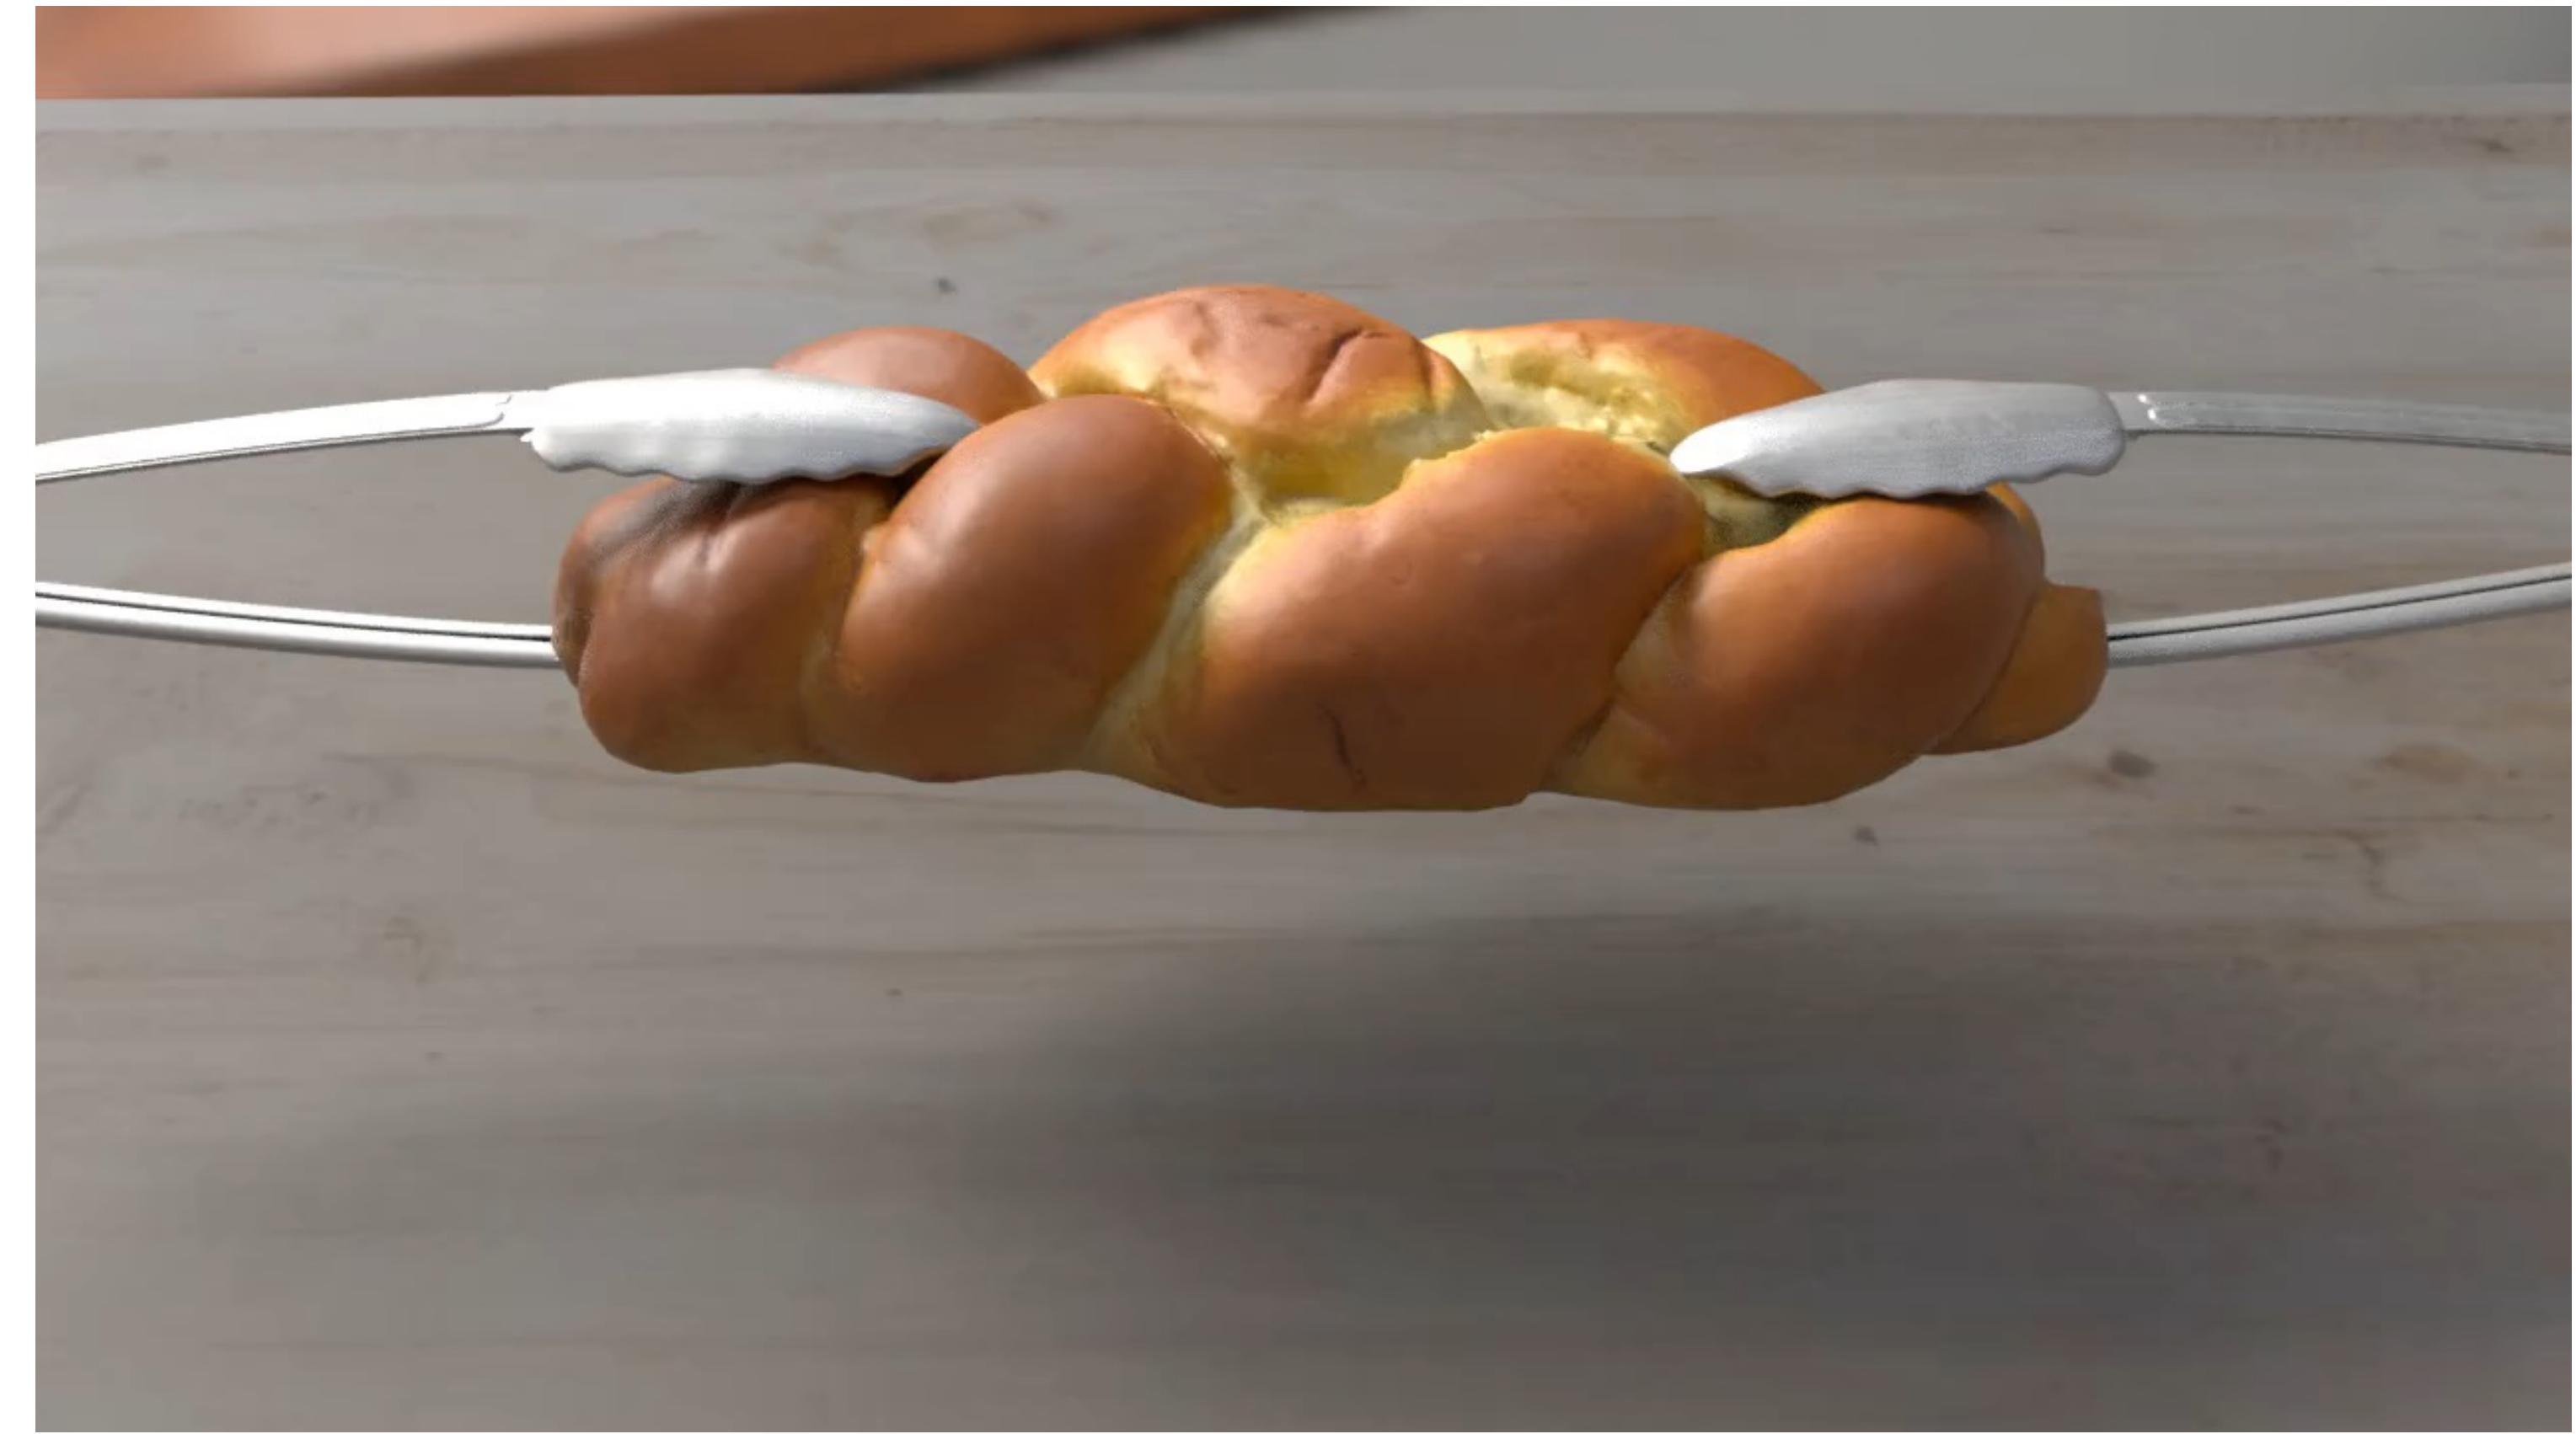
# Neural Constitutive Laws Pipeline



# Neural Stress Fields

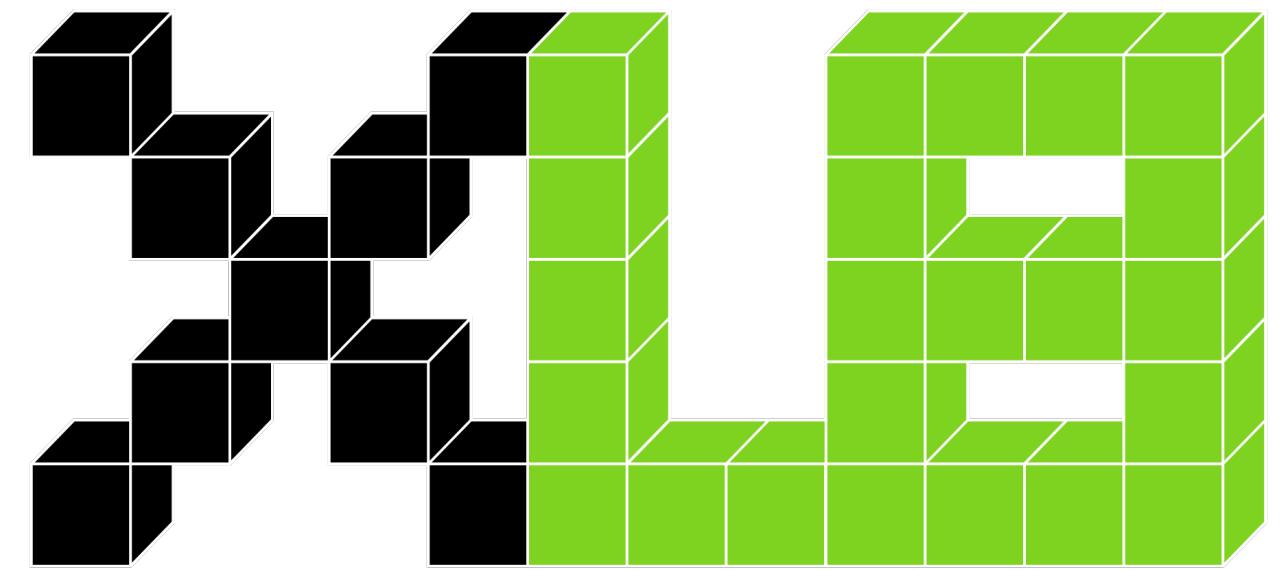
## Reduced-Order Elasto-plasticity and Fracture

- Goal: Accelerated MPM solvers for elastoplastic simulation
- Created differentiable MPM solver in Warp
- Learn a reduced neural basis to represent stress fields
- Dimension reduction of up to 100,000x
- 10x faster simulation
- Generalizes to a broad range of materials
- Joint work from UCLA, CMU, MIT, UToT
- <https://zeshunzong.github.io/reduced-order-mpm/>
- [Zhong et al. 2023]



# Large Scale Fluid Simulation

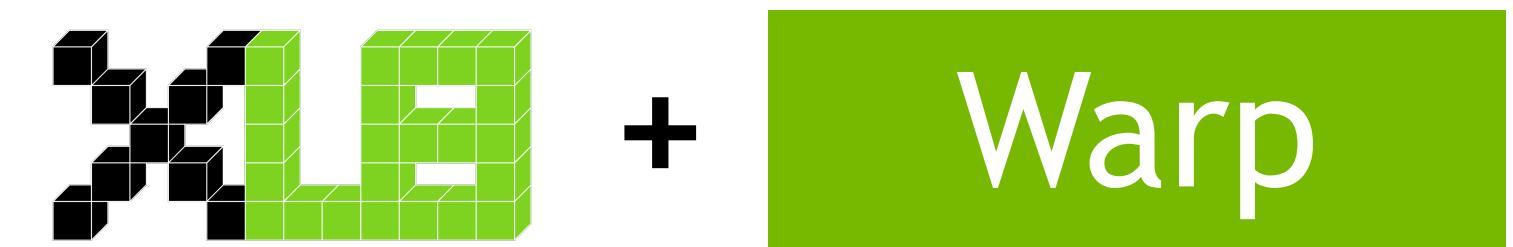
## Modulus



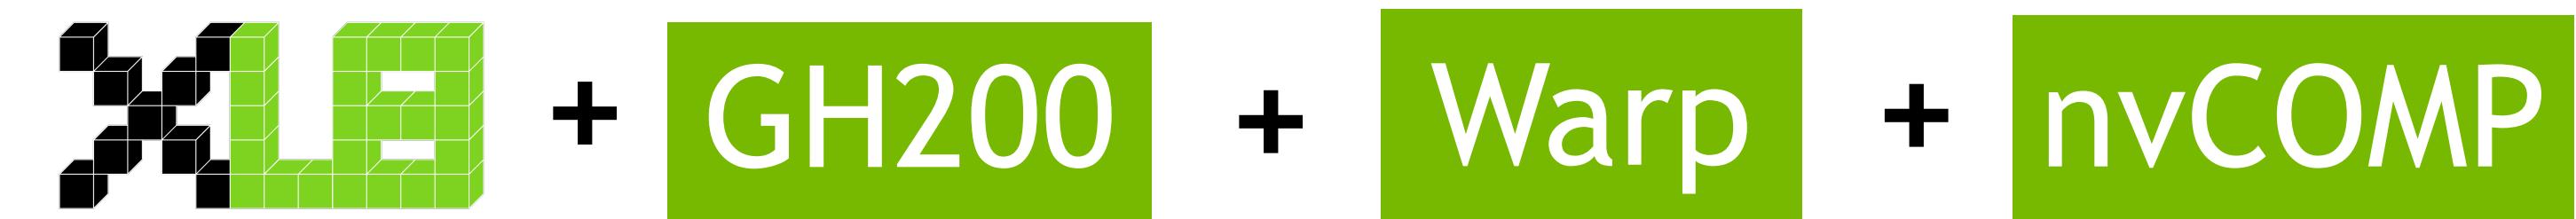
XLB is a Lattice Boltzmann fluid dynamics solver developed by Autodesk Research.

Developed in Python using JAX with key features:

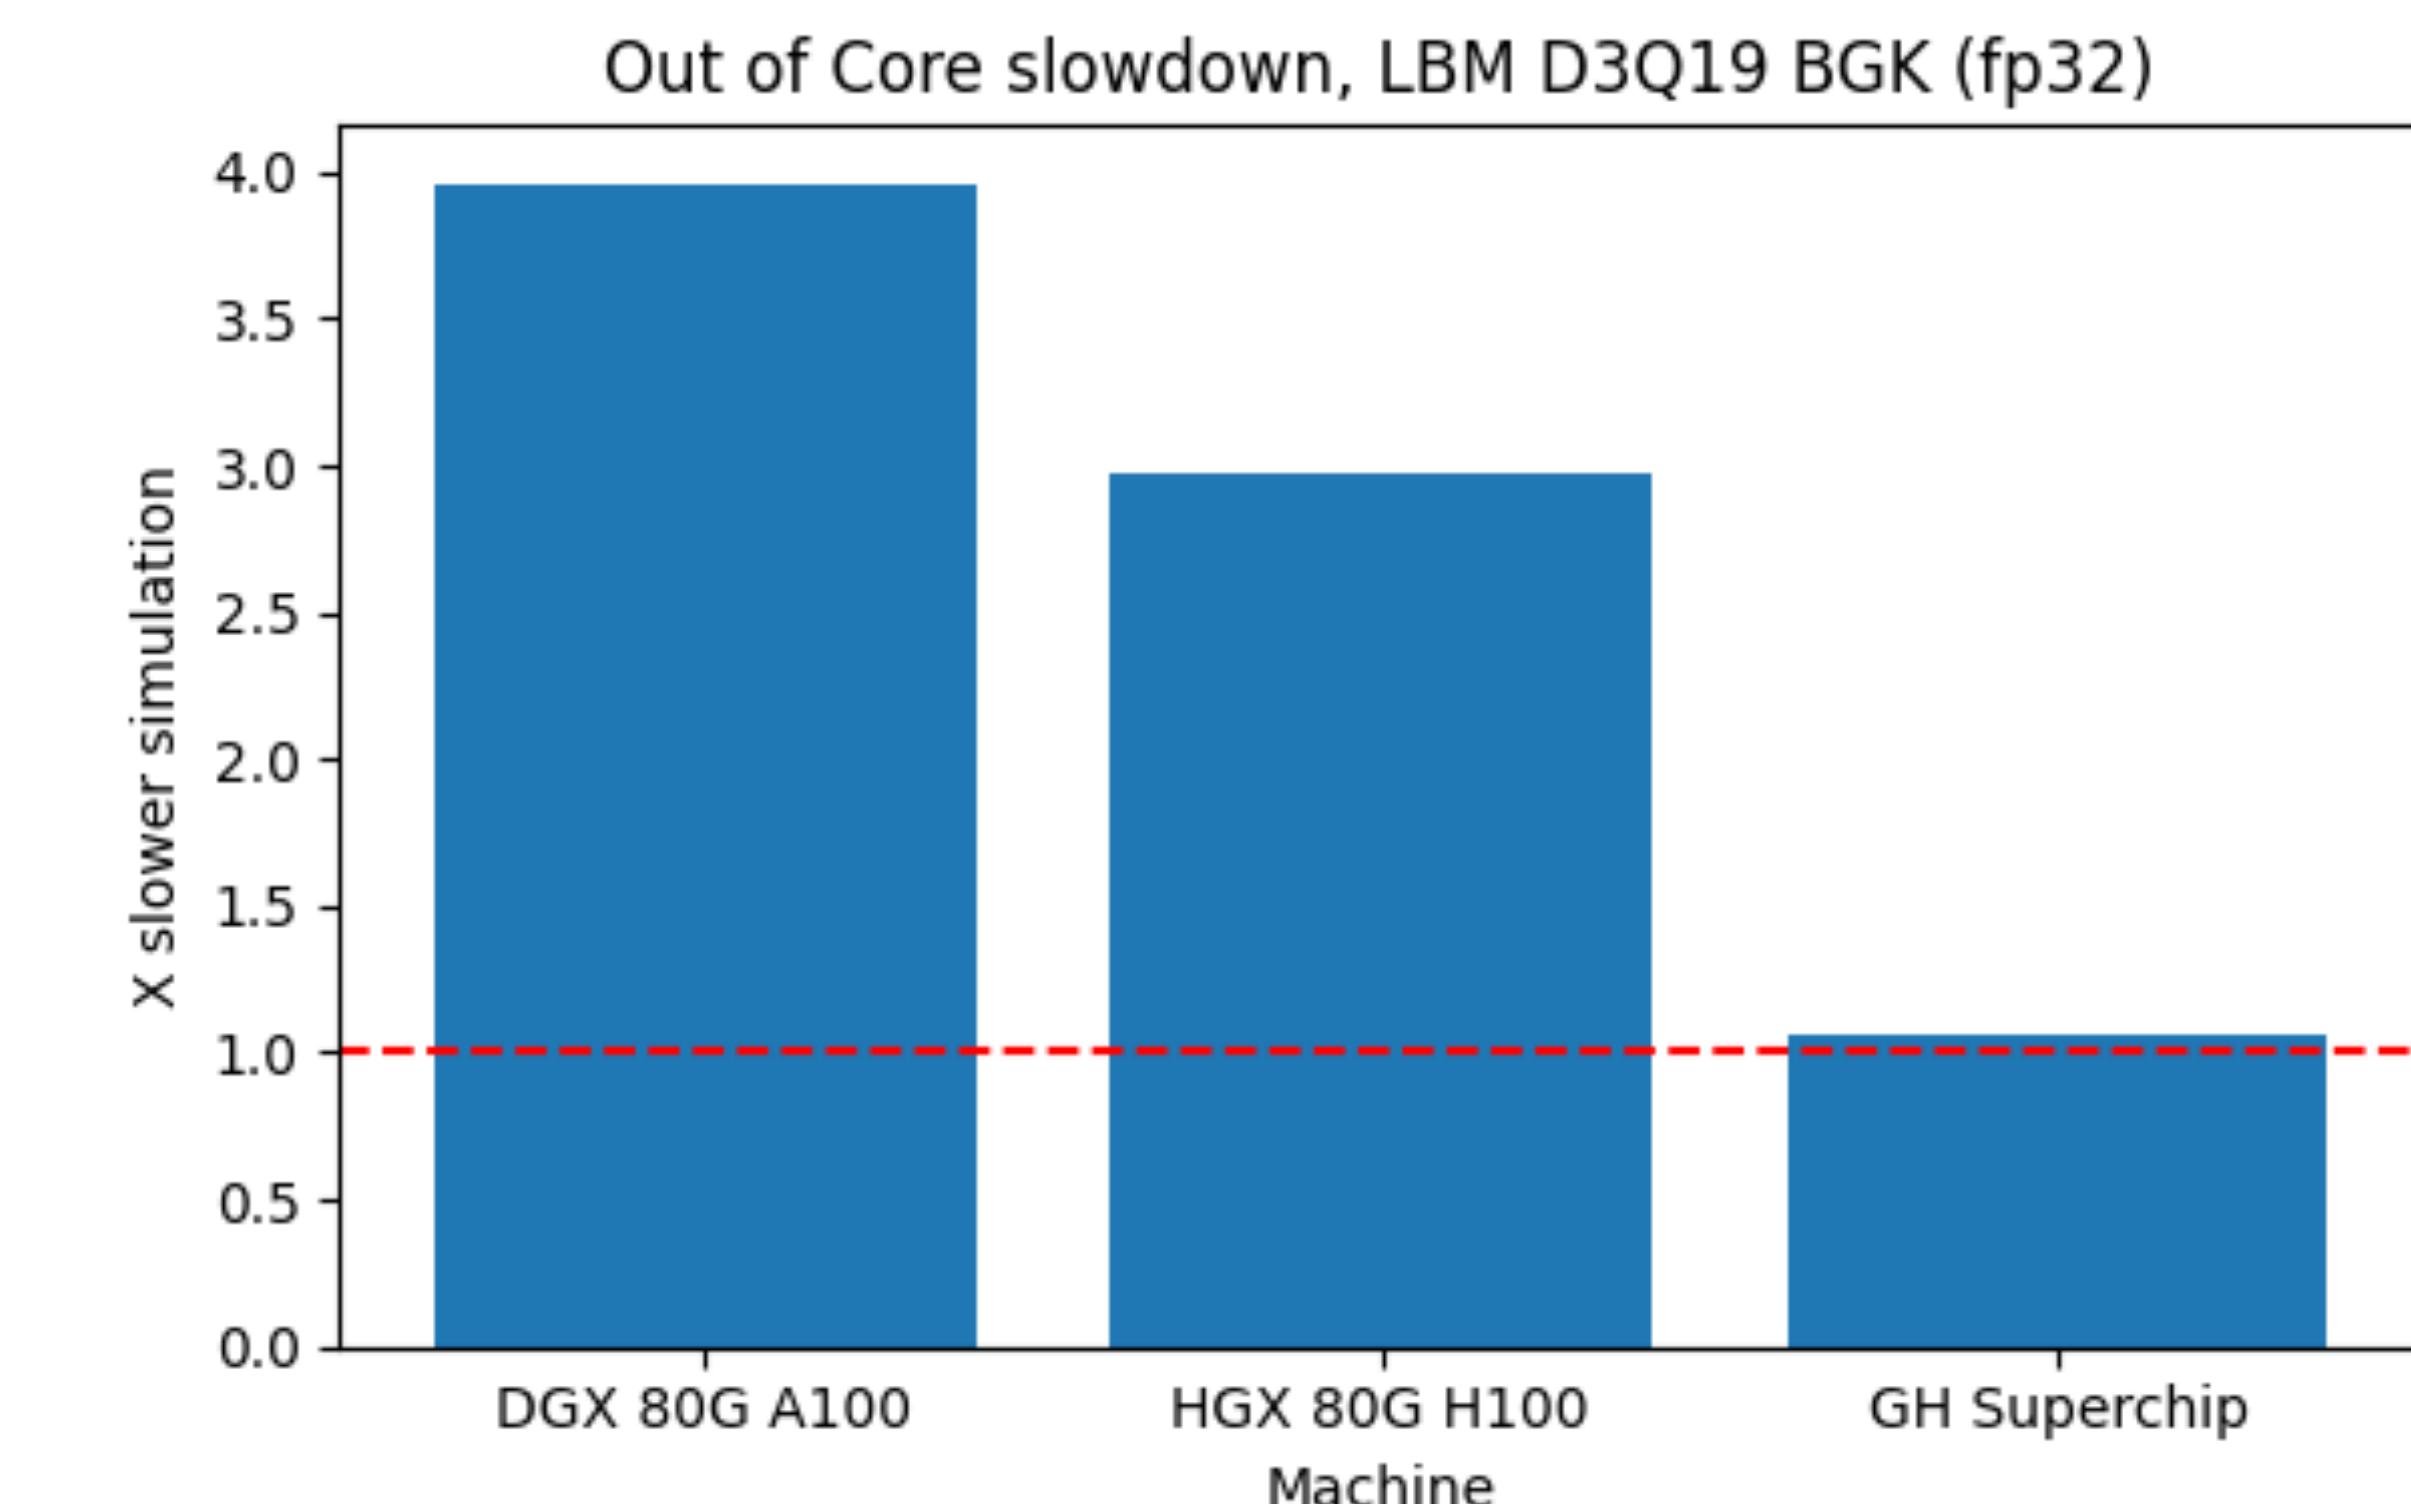
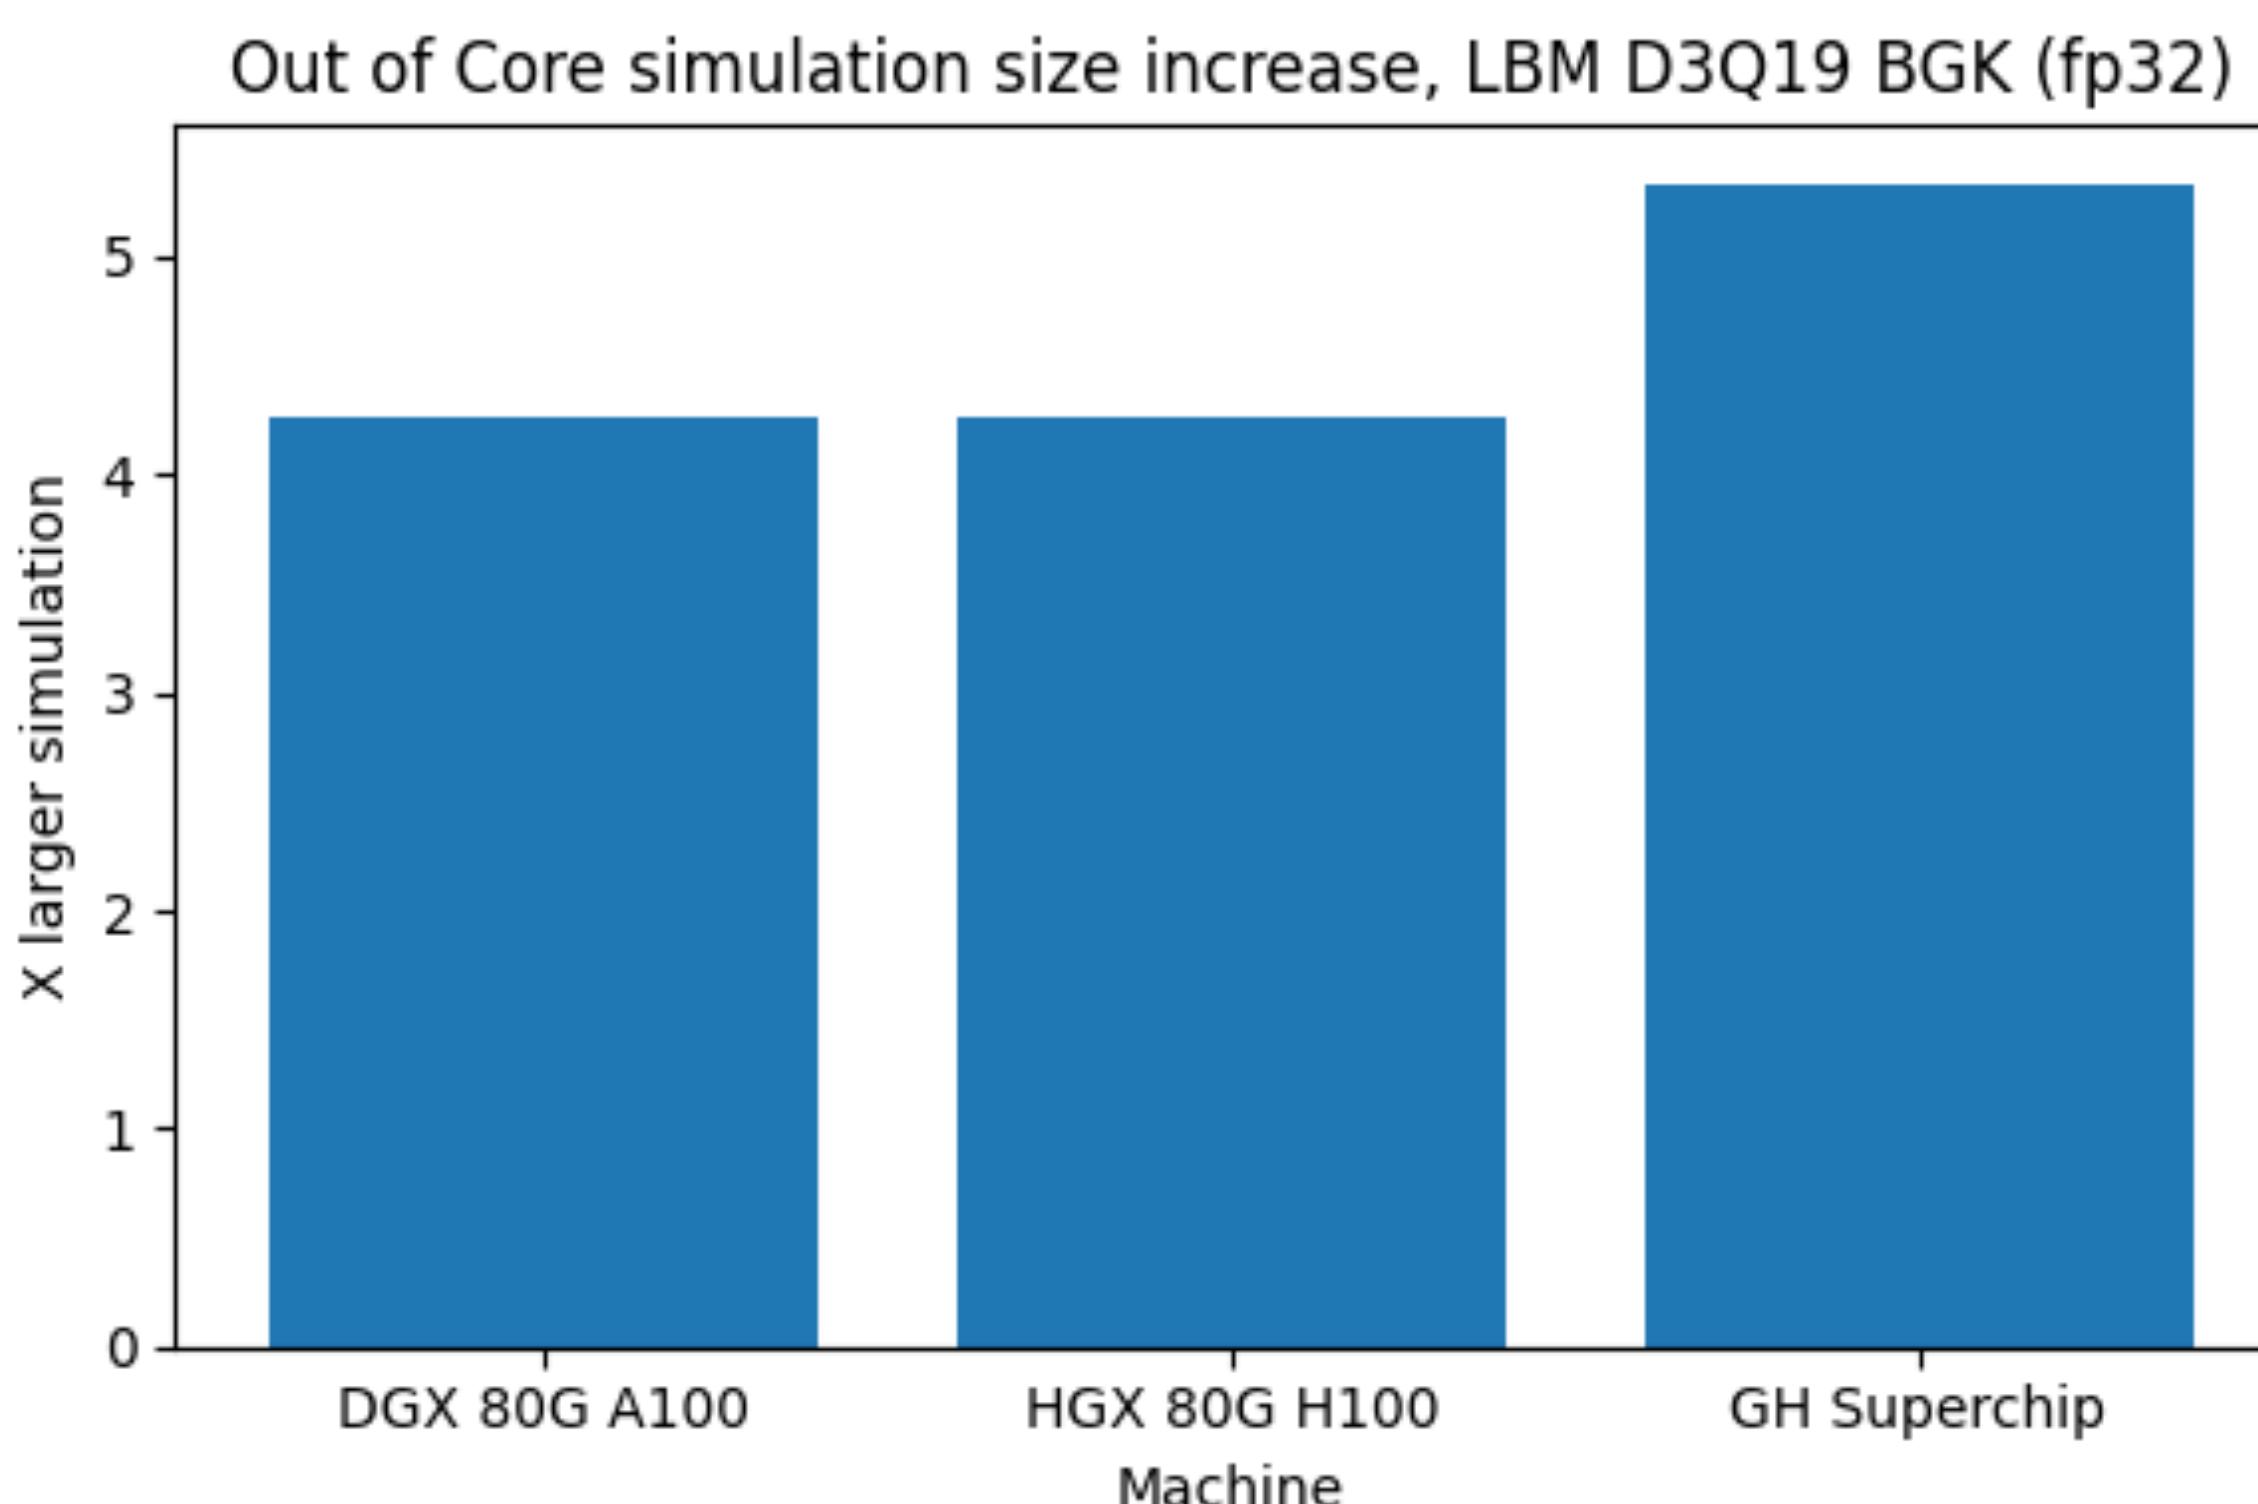
- Easily extendable to add new methods
- Scalable across distributed systems
- ML friendly allowing integration with physics-based ML models
- Topological optimization and inverse problems



XLB is being restructured to allow Warp as a compute backend. This results in ~3.8x speed improvement and ~3.1x larger simulations

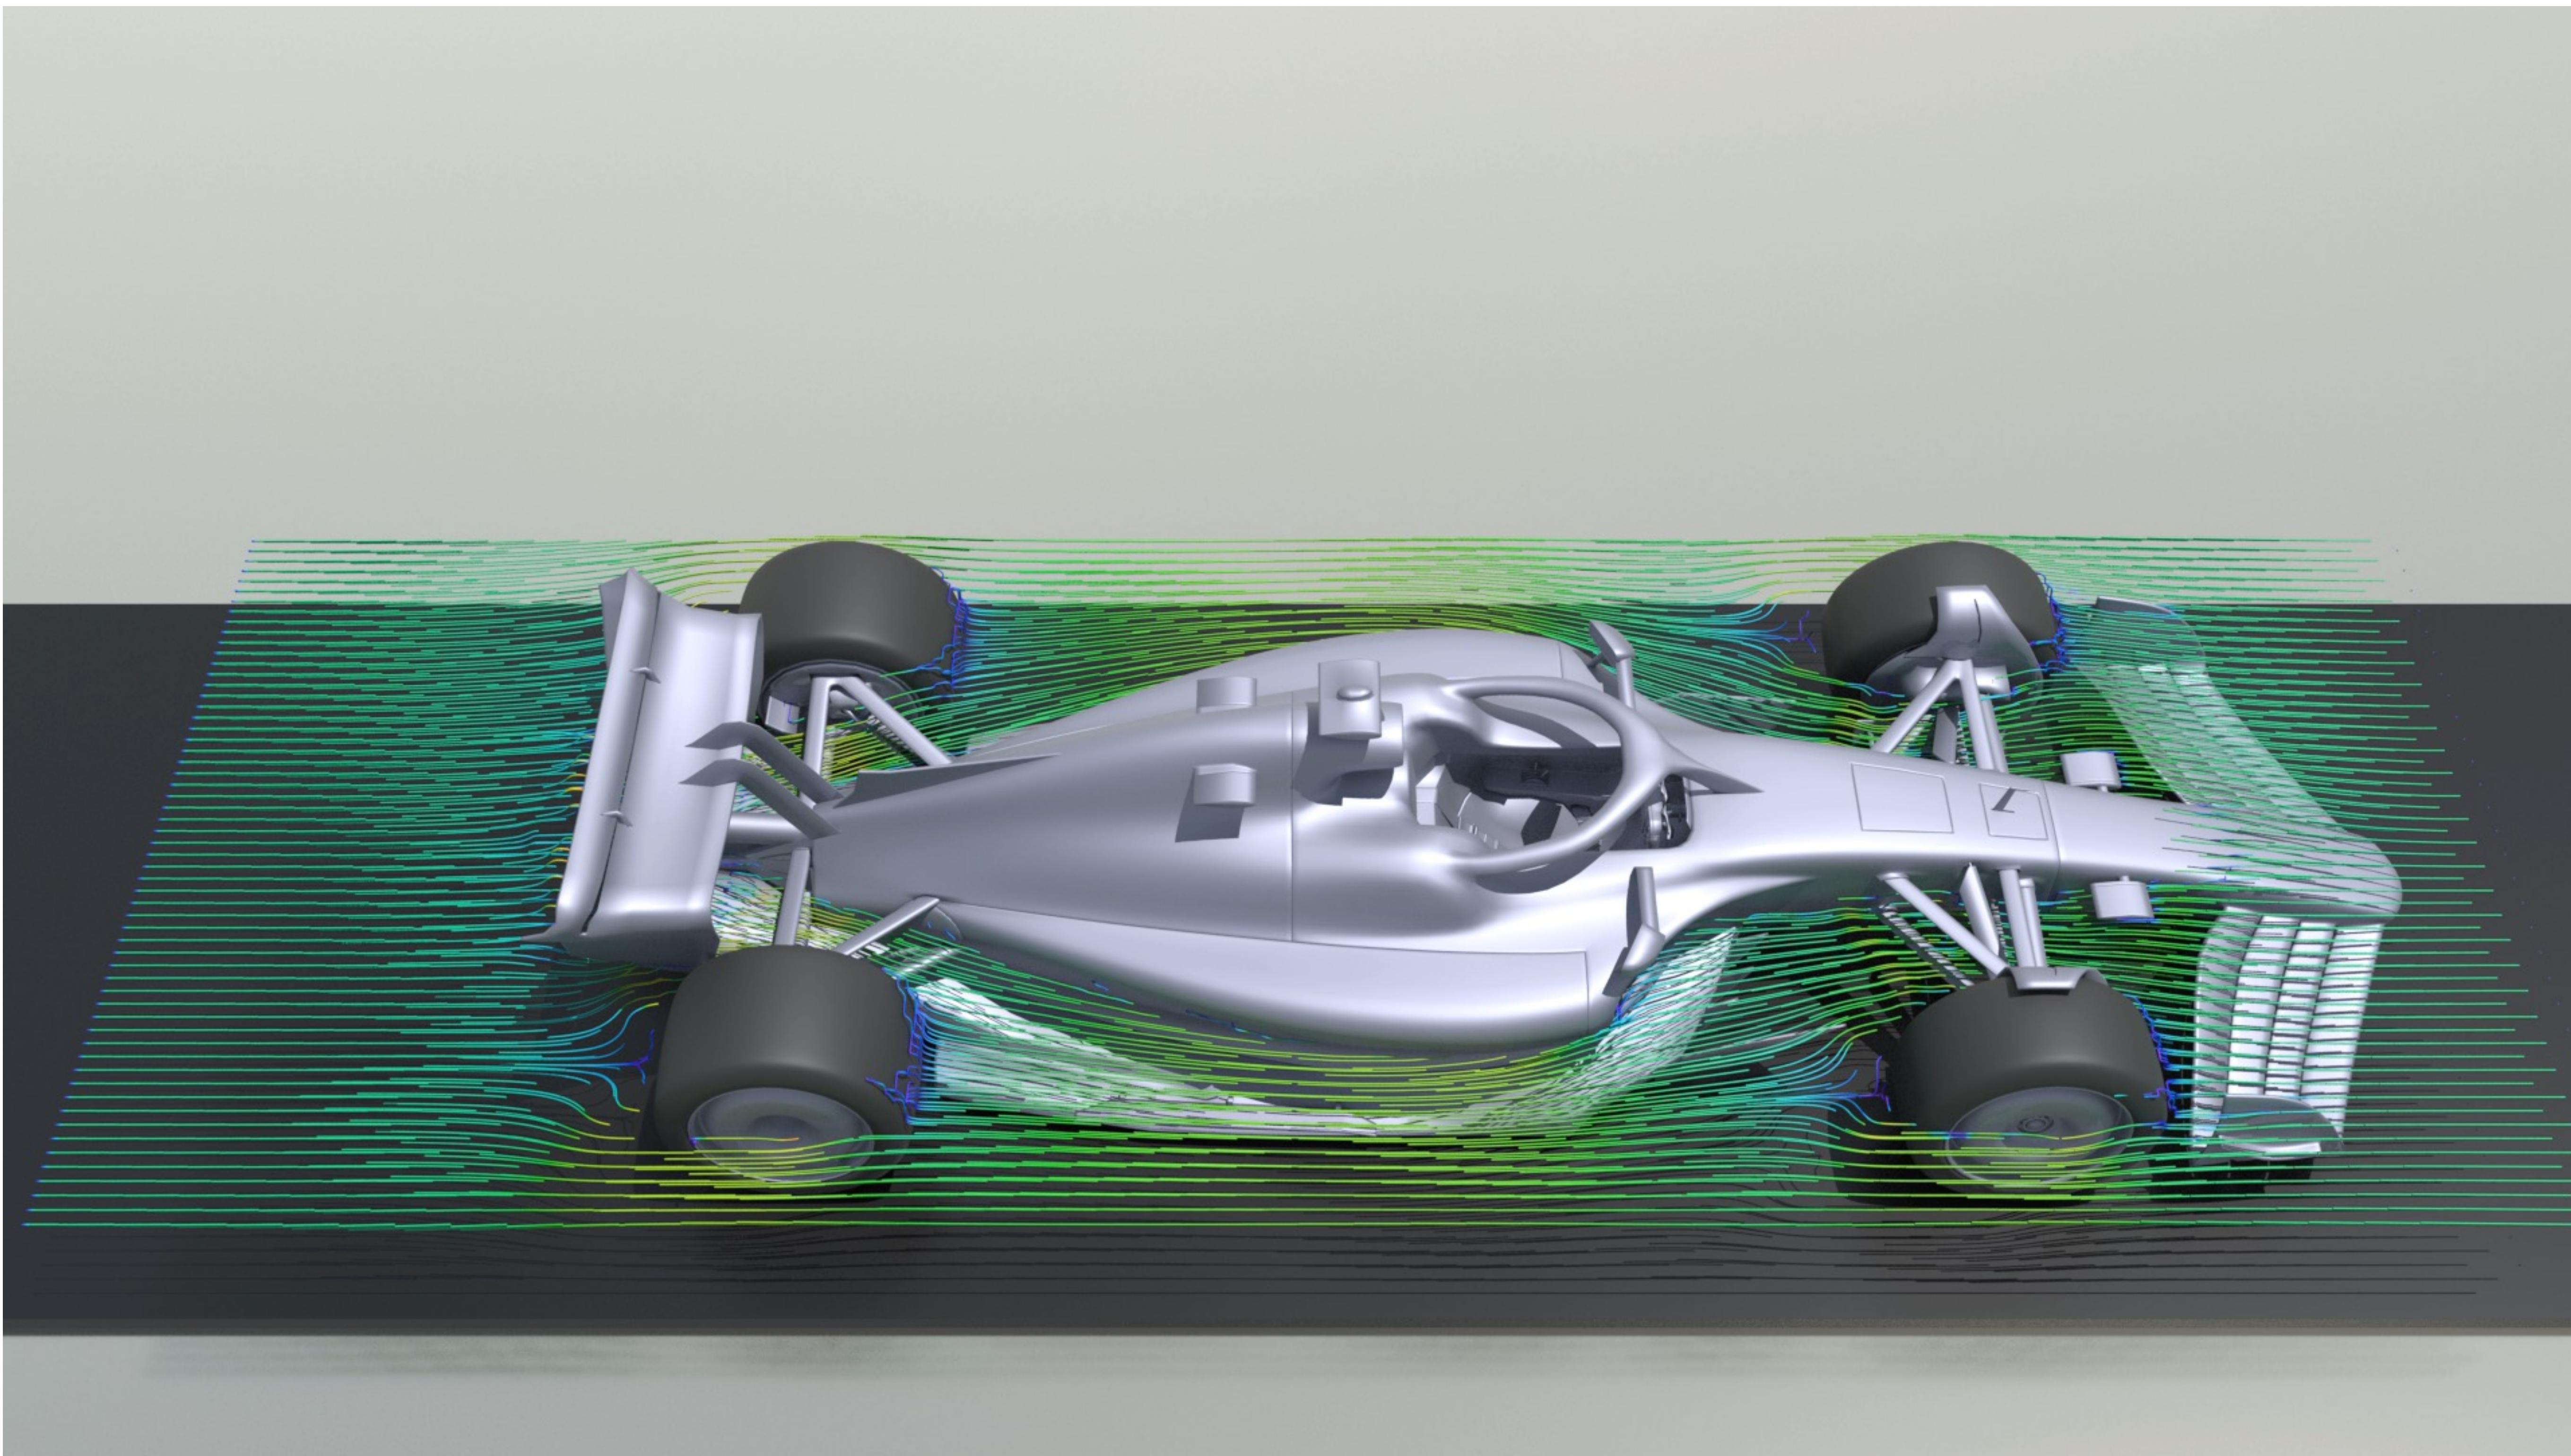


Combining XLB, Warp, nvCOMP and our Grace Hopper super chips allows for efficient out of core memory simulations. This allows for ~5x larger simulations with minimal impact on performance.



# Large Scale Fluid Simulation

## Aerodynamic Simulation

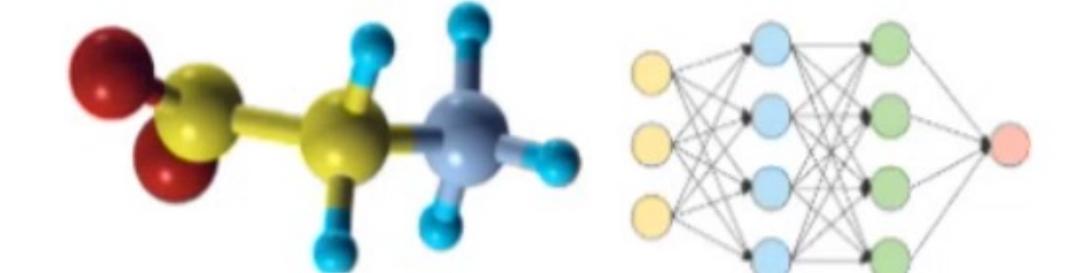
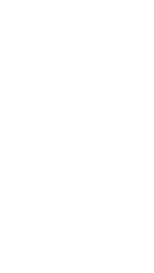


# Differentiable Molecular Dynamics

## Protein Docking

- Goal: Accelerate protein docking simulation
- Implement equivariant operators
  - Rotation / translation invariant
- Fastest ML framework implementation in Warp:
  - 10x faster than Torch
  - Near native CUDA performance at a fraction of the effort
- Create customizable differentiable kernels at runtime for specific models, e.g.: DiffDock\*
  - Runtime kernel specialization for different datatypes, dimensions, etc.
  - Faster than fixed-parameter hand rolled CUDA
- \* Corso et al. [2023]

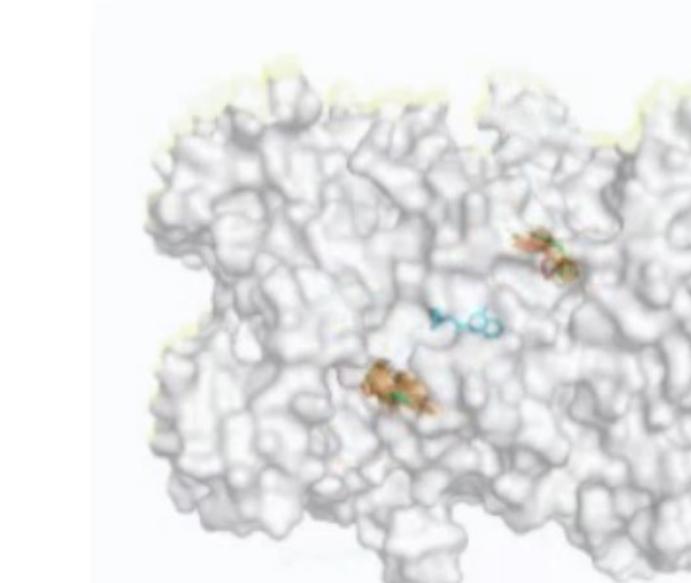
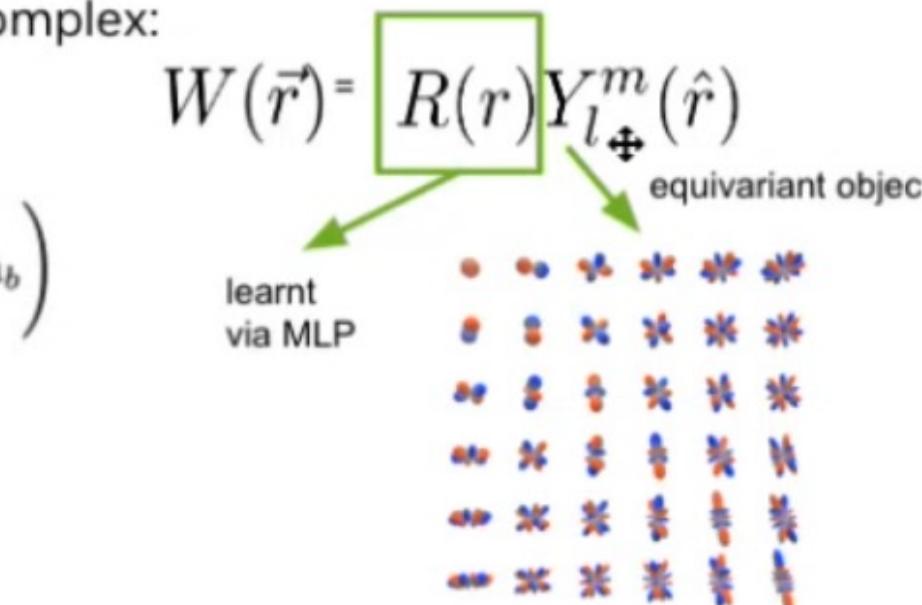
- Input are atomic coordinates (vectors)
- Key output: translation vector for reverse diffusion -> position update for molecule relative to protein
- Architecture: SE(3) equiv-GNN (trans, rot and no reflection)



Equivariant math operations in the model are complex:

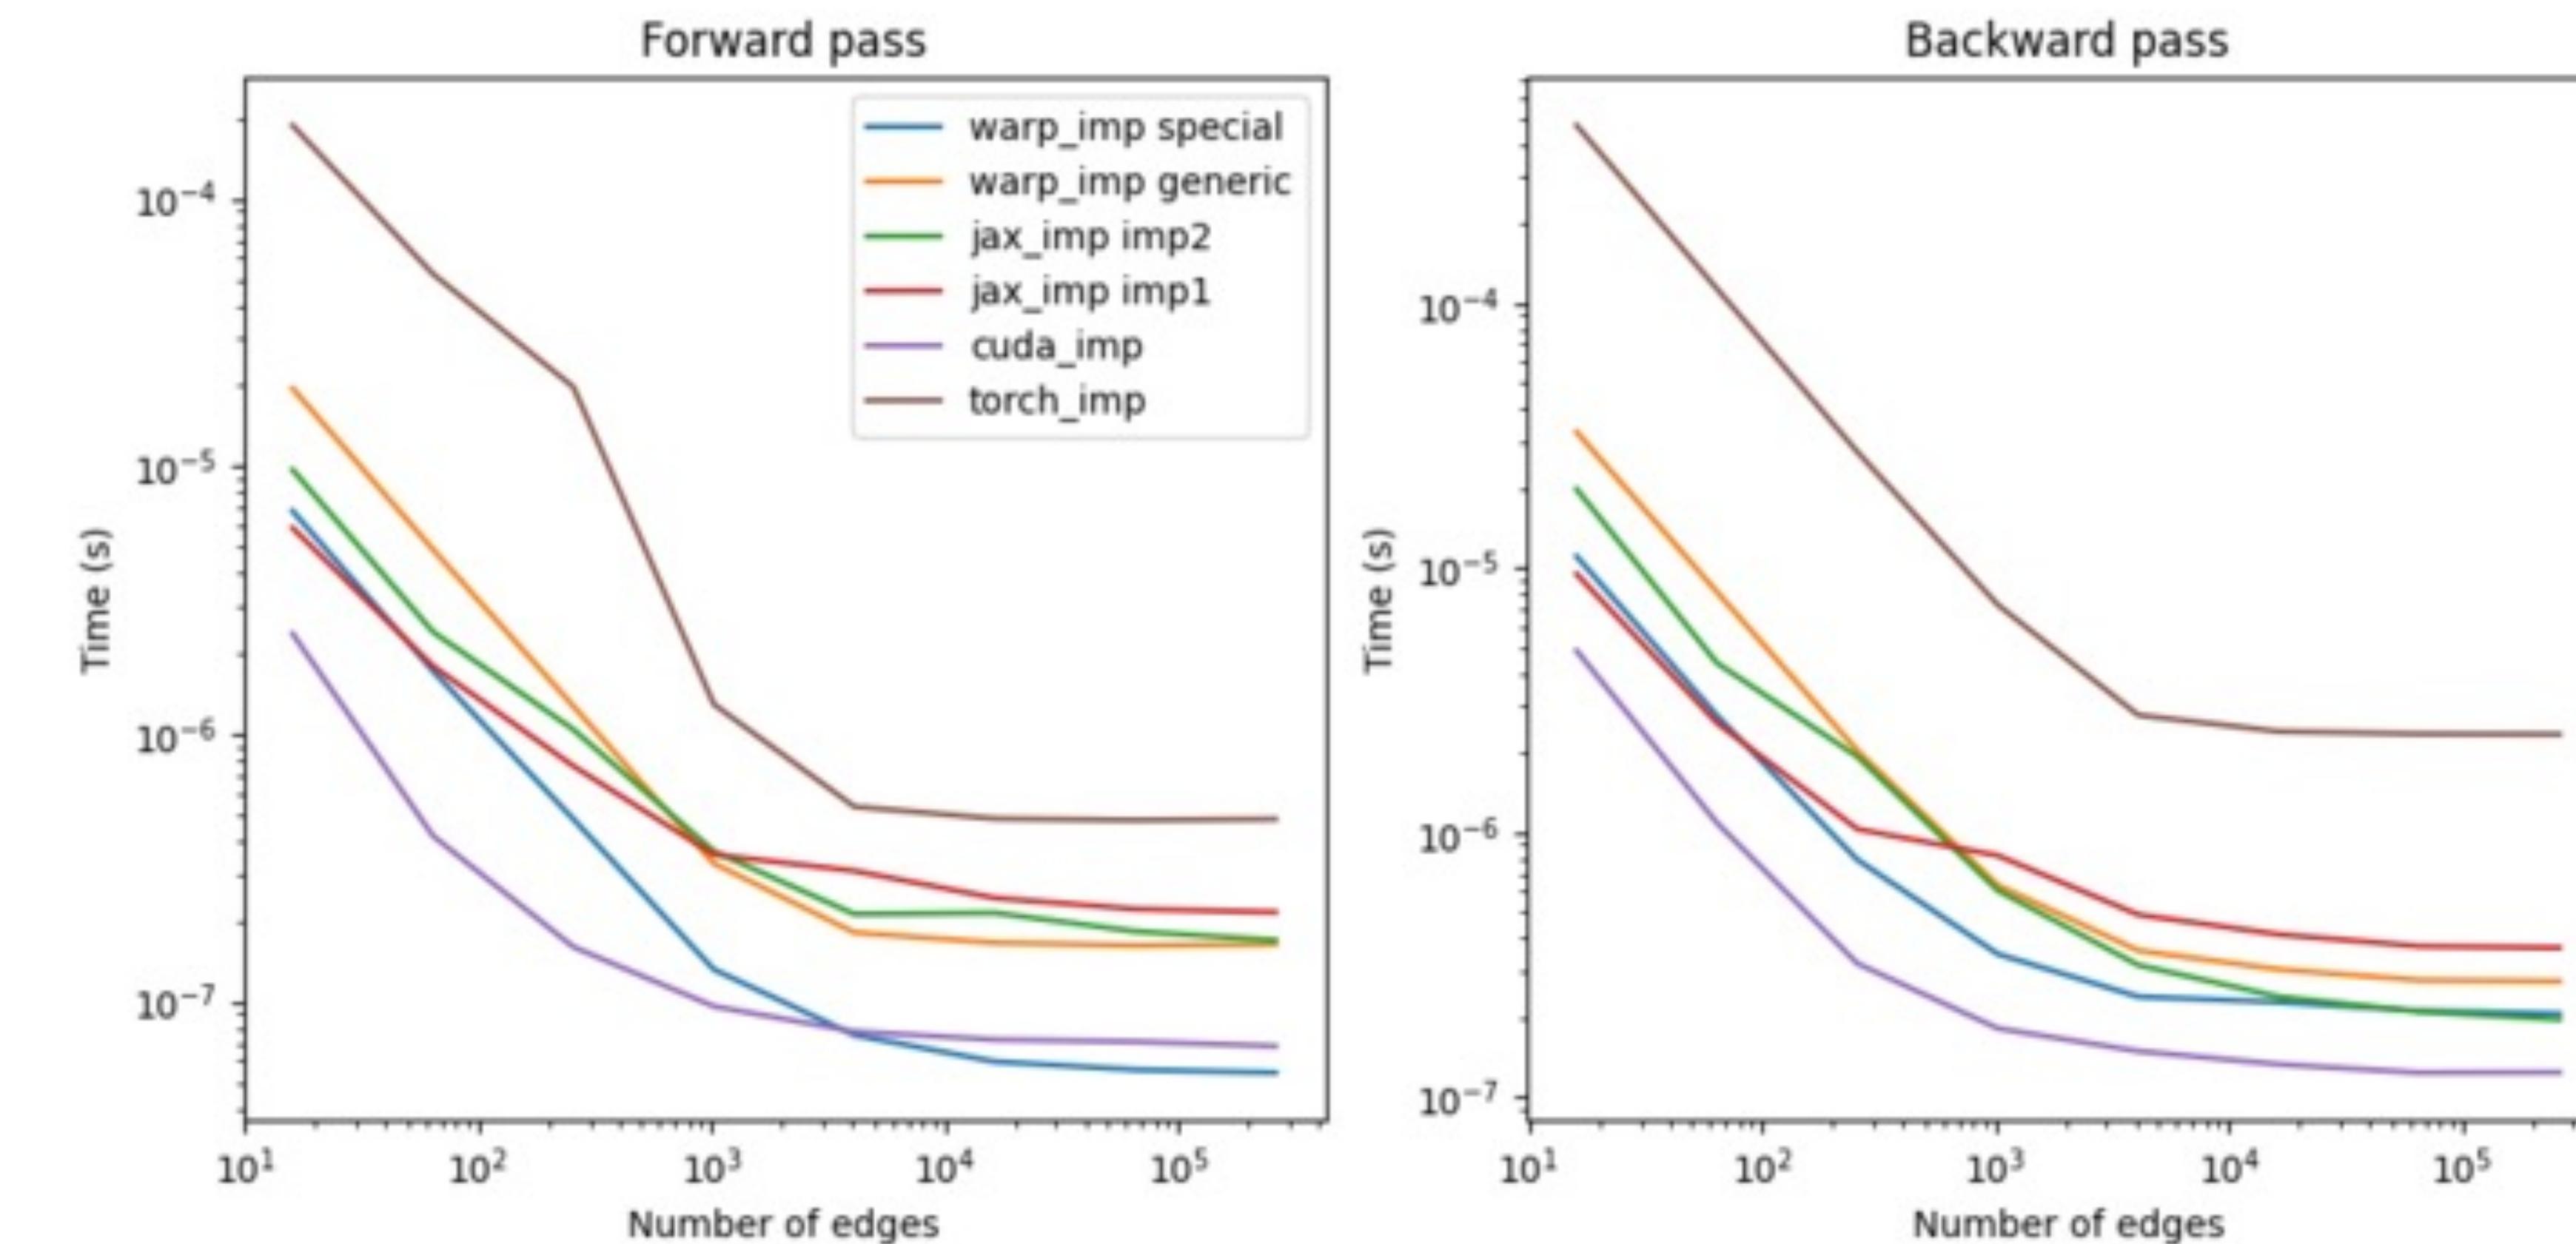
$$\mathbf{h}_a \leftarrow \mathbf{h}_a \bigoplus_{t \in \{\ell, r\}} \text{BN}^{(t_a, t)} \left( \frac{1}{|\mathcal{N}_a^{(t)}|} \sum_{b \in \mathcal{N}_a^{(t)}} Y(\hat{r}_{ab}) \otimes_{\psi_{ab}} \mathbf{h}_b \right)$$

with  $\psi_{ab} = \Psi^{(t_a, t)}(e_{ab}, \mathbf{h}_a^0, \mathbf{h}_b^0)$



Prot: 6agt  
Green: ground truth / Cyan: EquiBind

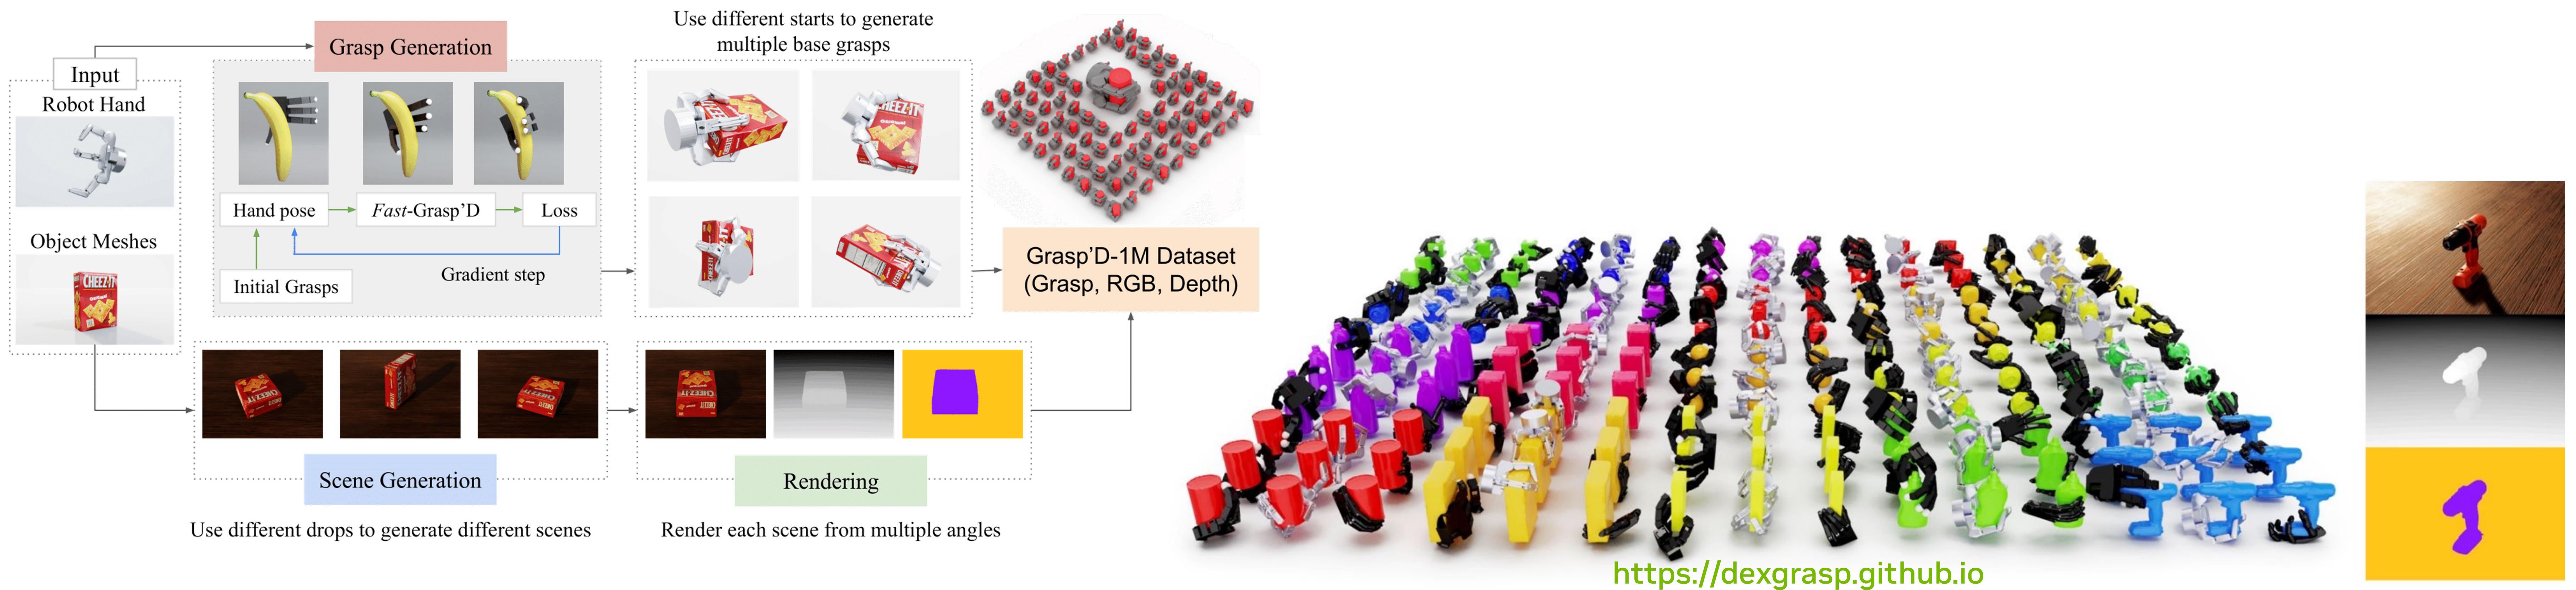
NVIDIA



# Fast-Grasp'D

## Dexterous Multi-finger Grasp Generation

- Goal: Generate realistic and stable grasps for robotic hand / object pairs
- Research from UofT + NVIDIA
- Differentiable SDF contact with rigid bodies
- 20x faster than previous SOTA
- Generated over 1M high quality grasps in 140 GPU hours over multiple hand types



# Aerial Vehicle Simulation

## Differentiable Dynamics simulation for Real-Time Control



# Robotic Perception

## ANYmal Quadruped Training

- Goal: Robot perception and processing for locomotion
- Extract the point cloud around 4000 robots using the built-in Warp hash grid
- ~12-16M points, 20ms for all the robots -> 27k FPS total
- Additional LIDAR simulation via. mesh ray-casts
- Deployed on Jetson AGX Orin hardware

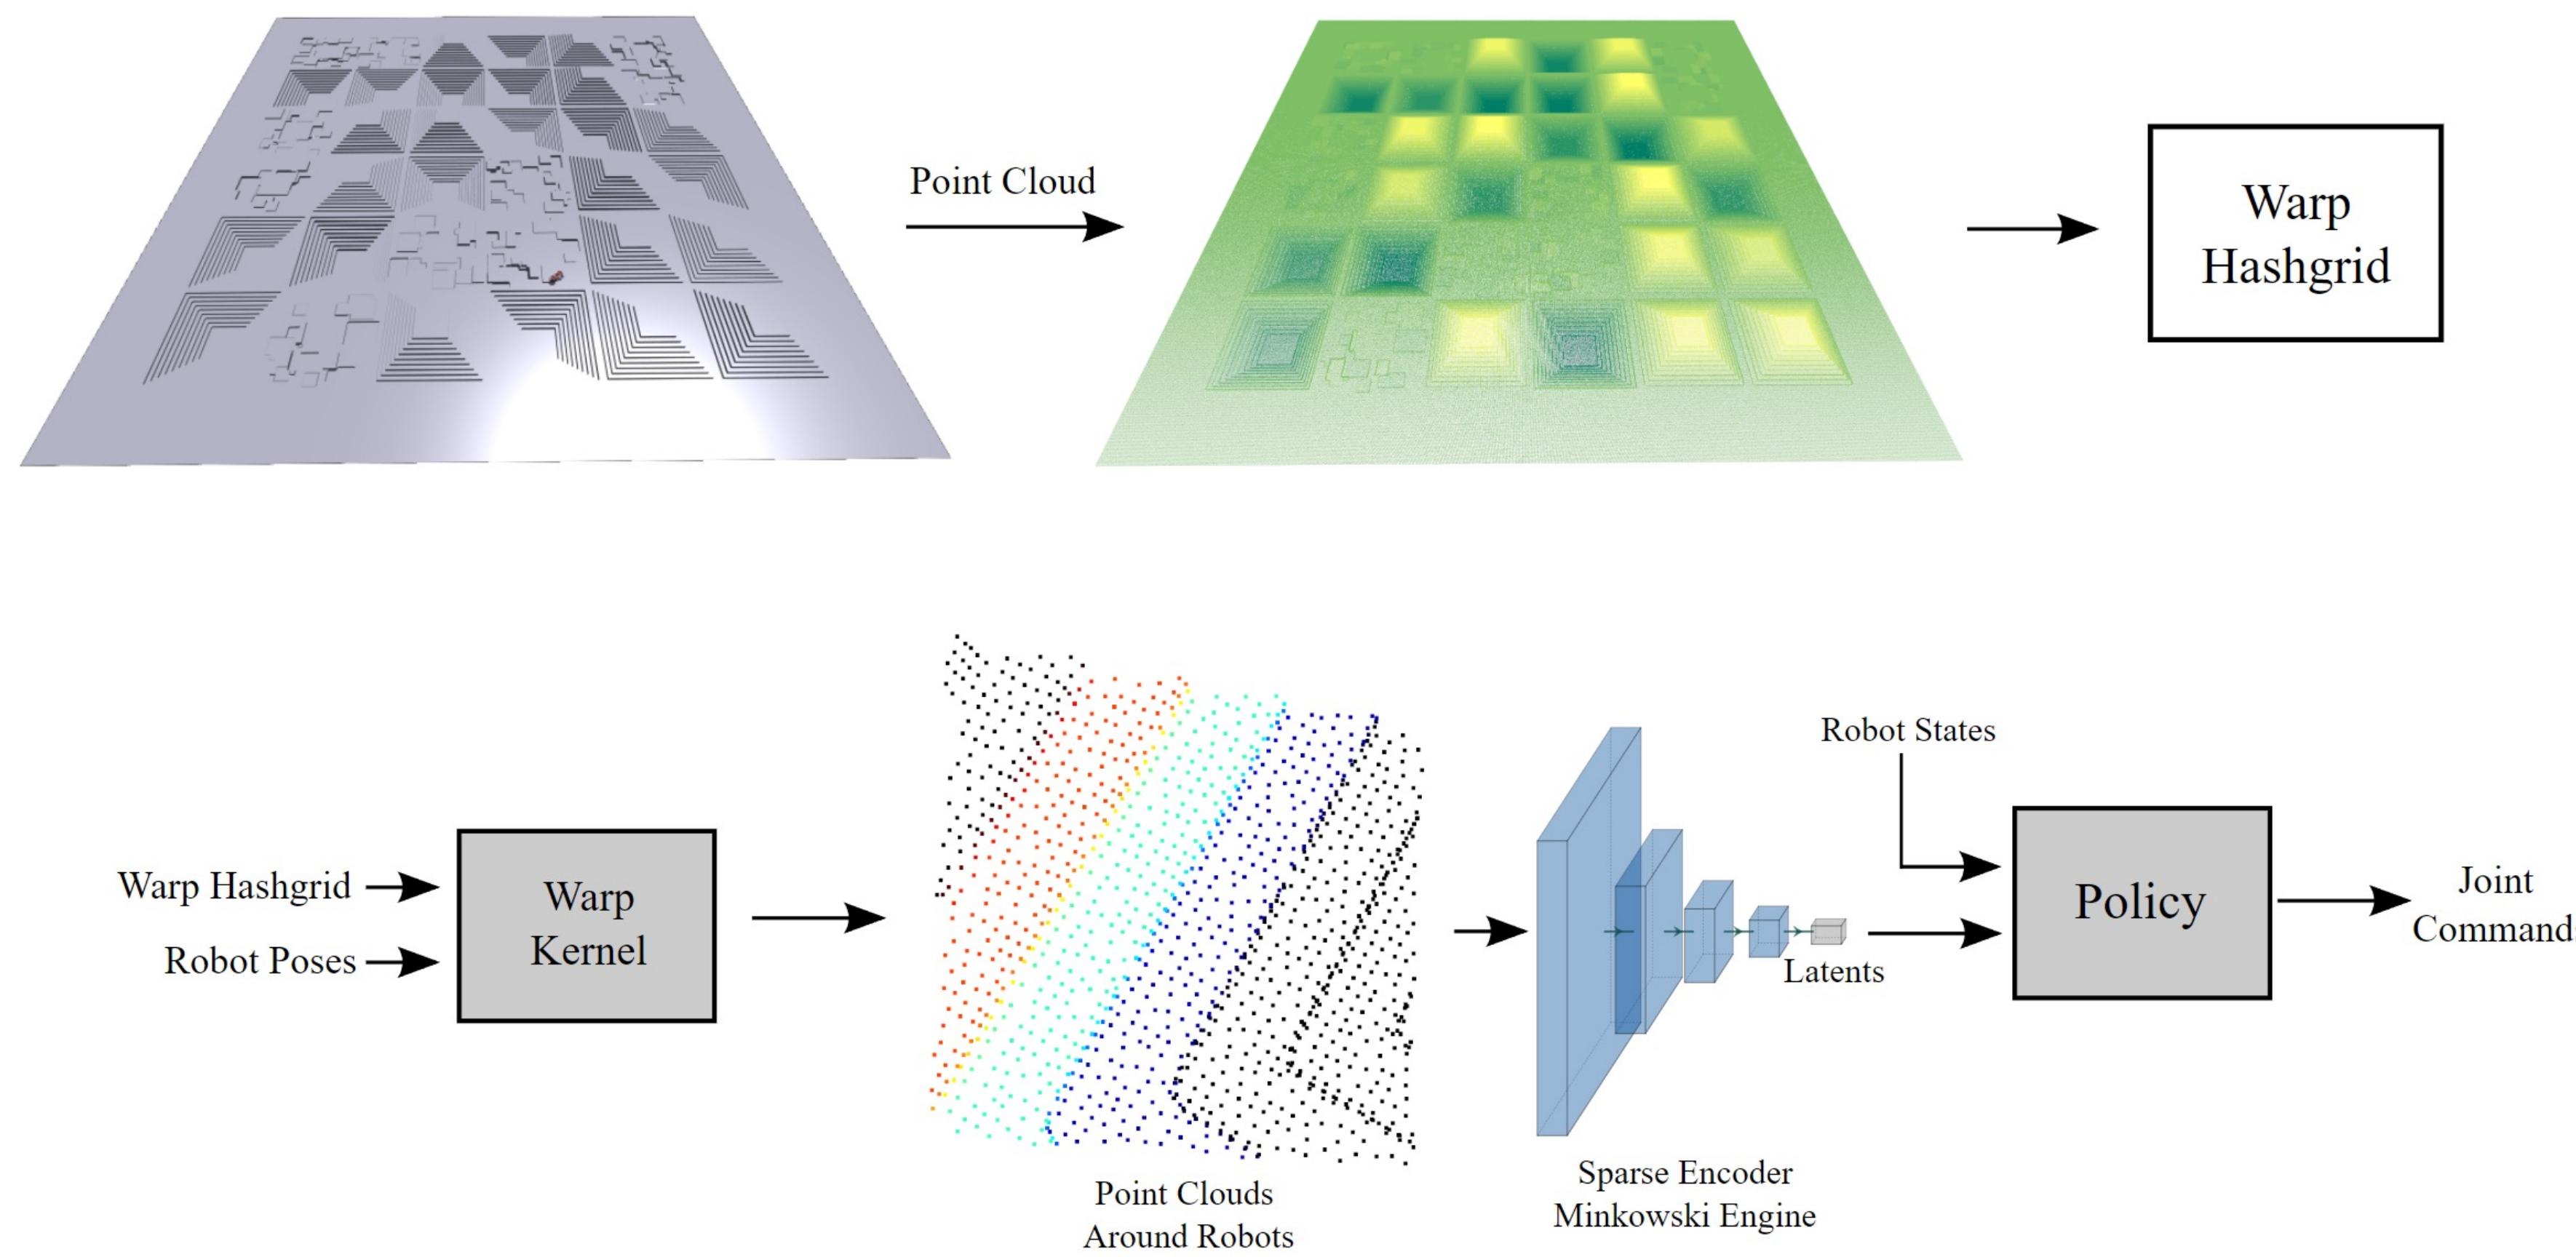
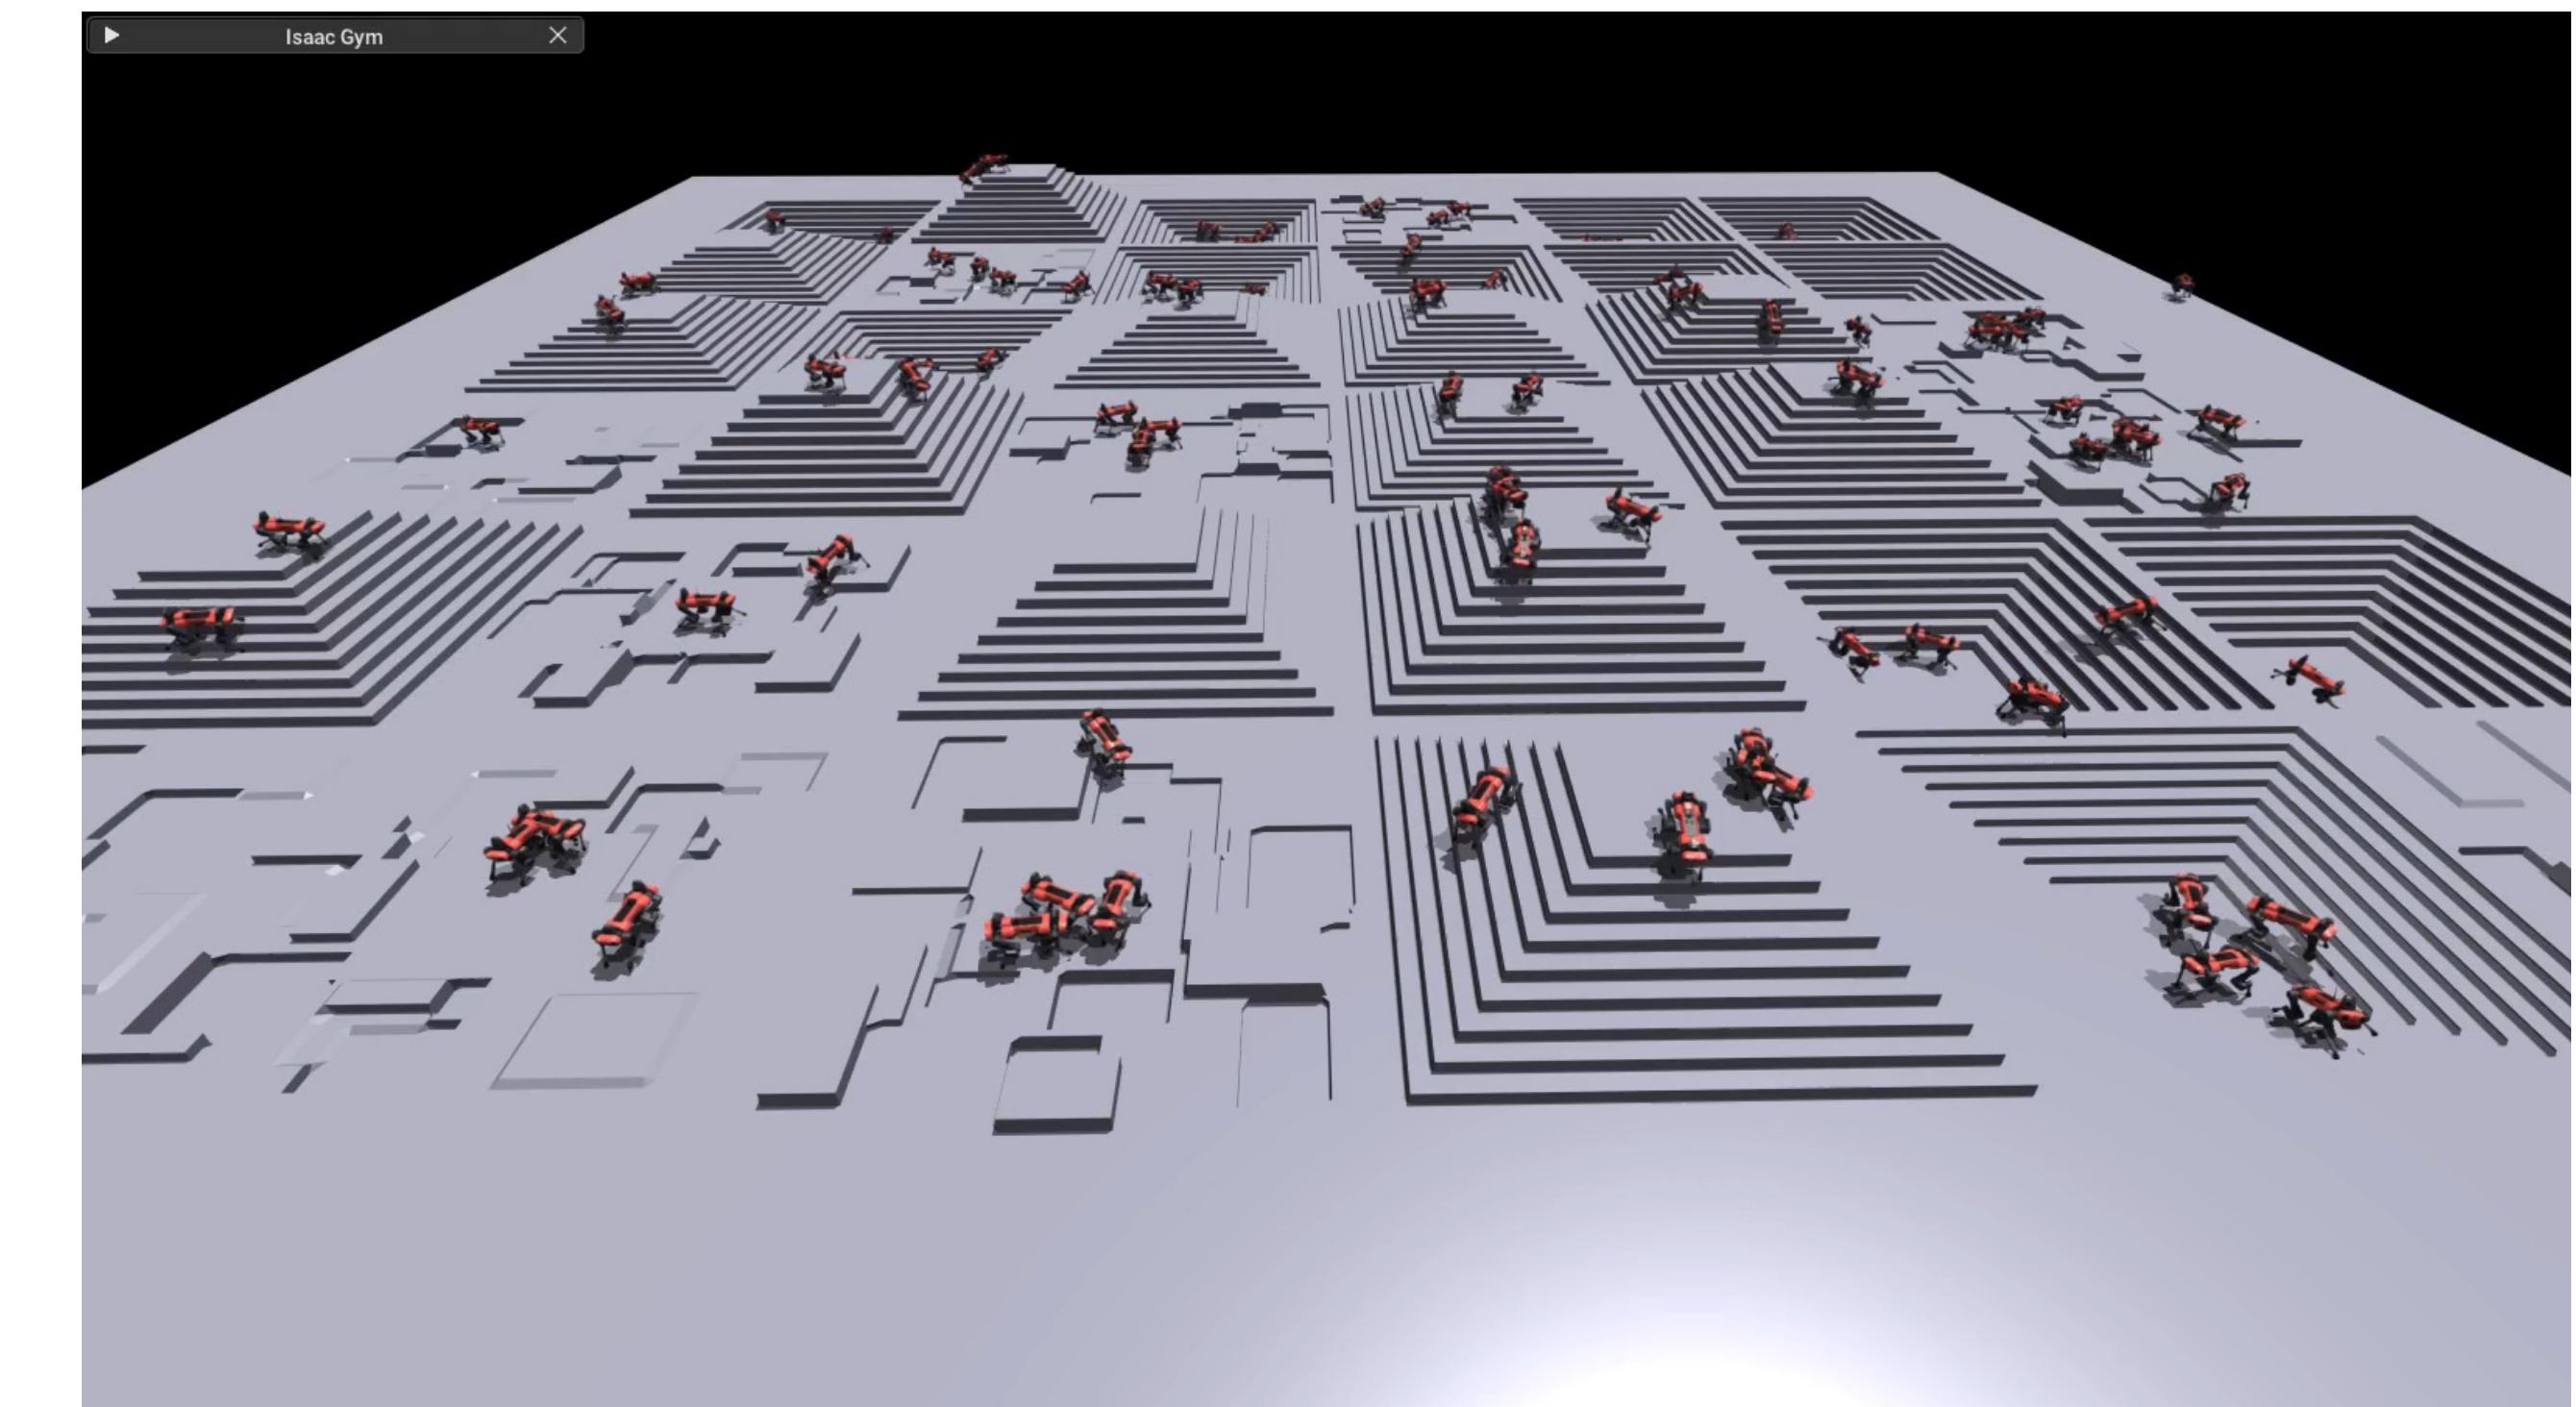


Image Courtesy Anybots ANYmal



# Roadmap

# Road Map

## Math Library Interop

- Integration with nvmath Python libraries
- Device-side versions of CUDA libraries:
  - cuBLASDx
  - cuFFTDx
- e.g. invoke FFT directly from inside Warp kernels
- Fully fused execution with user code
- See Math Libraries Python talk on Wednesday!

### Deep Dive into Math Libraries [S62162]

Arthy Sundaram, Senior Product Manager, NVIDIA

Harun Bayraktar, Director of Engineering, Math Libraries, NVIDIA

NVIDIA's GPU-accelerated Math Libraries, which are part of the CUDA Toolkit and the HPC SDK, are constantly expanding, providing industry-leading performance and coverage of common compute workflows across AI, ML, and HPC. We'll do a deep dive...

Add to Schedule ⓘ

Wednesday, Mar 20 | 9:00 AM - 9:50 AM PDT  
(Thursday, Mar 21 | 5:00 AM - 5:50 AM NZDT)

ⓘ Hilton Almaden 1 (L1)

```
import nvmath as nvm
import warp as wp

FFT_fwd = nvm.device.fft(type='c2c', size=size, precision=wp.float32, direction='forward', dsl='warp')
FFT_inv = nvm.device.fft(type='c2c', size=size, precision=wp.float32, direction='inverse', dsl='warp')

@wp.kernel
def compute(input: wp.array(dtype=float),
            output: wp.array(dtype=float),
            fwd: FFT_fwd,
            inv: FFT_inv):

    i = wp.tid()

    # epilogue
    thread_data = user_func(input[i])

    # execute forward FFT
    fwd(FFT_fwd, thread_data)

    # frequency space convolution
    thread_data /= scale

    # execute inverse FFT
    inv(FFT_inv, thread_data)

    # prologue
    output[i] = thread_data
```

Example: Fused Forward/Reverse FFT

# Road Map

## Warp Neural Networks

- Leverage TensorCore operations (MMA) inside Warp kernels, e.g.:
  - `wp.linear()`
  - `wp.map()`
- Ideal for building **fused** neural network inference in Warp kernels
- Keep all intermediate values local to kernel
- Network inference in-situ

```
@wp.kernel
def eval(layer0: Layer,
         layer1: Layer,
         loss: wp.array(dtype=float)):

    # setup feature vectors / input
    x = wp.vector(dtype=wp.float16, length=DIM_IN)
    x[0] = 0.0
    x[1] = 1.0
    x[2] = 2.0
    x[3] = 3.0

    # network eval
    z = wp.linear(layer0.weights, layer0.bias, DIM_HID, z)
    z = wp.map(relu, z)
    z = wp.linear(layer1.weights, layer1.bias, DIM_OUT, z)
    z = wp.map(relu, z)

    # compute loss
    l = wp.length(z)

    # write loss
    wp.atomic_add(loss, 0, l)
```

Example: Fused MLP Inference

# Road Map

## Warp Neural Networks

```

@wp.kernel
def eval_tetrahedra(args):
    tid = wp.tid()
    ...
    act = activation[tid]
    k_mu = materials[tid, 0]
    k_lambda = materials[tid, 1]
    k_damp = materials[tid, 2]
    Ds = wp.mat33(x10, x20, x30)
    Dm = pose[tid]
    inv_rest_volume = wp.determinant(Dm) * 6.0
    rest_volume = 1.0 / inv_rest_volume
    alpha = 1.0 + k_mu / k_lambda - k_mu / (4.0 * k_lambda)
    # scale stiffness coefficients to account for area
    k_mu = k_mu * rest_volume
    k_lambda = k_lambda * rest_volume
    k_damp = k_damp * rest_volume
    # F = Xs*Xm^-1
    F = Ds * Dm
    dFdt = wp.mat33(v10, v20, v30) * Dm
    col1 = wp.vec3(F[0, 0], F[1, 0], F[2, 0])
    col2 = wp.vec3(F[0, 1], F[1, 1], F[2, 1])
    col3 = wp.vec3(F[0, 2], F[1, 2], F[2, 2])
    # -----
    # Neo-Hookean (with rest stability [Smith et al 2018])
    Ic = dot(col1, col1) + dot(col2, col2) + dot(col3, col3)
    # deviatoric part
    P = F * k_mu * (1.0 - 1.0 / (Ic + 1.0)) + dFdt * k_damp
    H = P * wp.transpose(Dm)
    f1 = wp.vec3(H[0, 0], H[1, 0], H[2, 0])
    f2 = wp.vec3(H[0, 1], H[1, 1], H[2, 1])
    f3 = wp.vec3(H[0, 2], H[1, 2], H[2, 2])
    J = wp.determinant(F)
    s = inv_rest_volume / 6.0
    dJdx1 = wp.cross(x20, x30) * s
    dJdx2 = wp.cross(x30, x10) * s
    dJdx3 = wp.cross(x10, x20) * s
    f_volume = (J - alpha + act) * k_lambda
    f_damp = (wp.dot(dJdx1, v1) + wp.dot(dJdx2, v2) + wp.dot(dJdx3, v3)) * k_damp
    f_total = f_volume + f_damp
    f1 = f1 + dJdx1 * f_total
    f2 = f2 + dJdx2 * f_total
    f3 = f3 + dJdx3 * f_total
    f0 = (f1 + f2 + f3) * (0.0 - 1.0)
    # apply forces
    wp.atomic_sub(f, i, f0)
    wp.atomic_sub(f, j, f1)
    wp.atomic_sub(f, k, f2)
    wp.atomic_sub(f, l, f3)

```

Traditional Elastic Model

Neural Network

```

@wp.kernel
def eval_tetrahedra(args):
    ...
    # setup feature vectors
    x = wp.vector(dtype=wp.float16, Length=DIM_IN)
    # evaluate layers
    z = wp.linear(layer0.weights, layer0.bias, DIM_HID, x)
    z = wp.map(relu, z)
    z = wp.linear(layer1.weights, layer1.bias, DIM_OUT, z)
    z = wp.map(relu, z)
    # apply forces
    wp.atomic_sub(f, i, f0)
    wp.atomic_sub(f, j, f1)
    wp.atomic_sub(f, k, f2)
    wp.atomic_sub(f, l, f3)

```

Learned Constitutive Model

# Road Map

## Warp Kernels in JAX

- Wrap Warp kernels as JAX primitives with one line of code
- Invoke JAX primitive inside `@jax.jit` code
- Very convenient way to mix JAX and Warp
- Available now in `wp.jax_experimental`
- Sharding, vmap(), pmap() support coming soon
- See JAX session tomorrow!

### What's New in JAX [S62659]

Frederic Bastien, Principal Software Engineer, NVIDIA

Matthew Johnson, Research Scientist, Google

JAX is a framework for high-performance machine learning and numerical computing research that achieves performance through compilation, automatic differentiation, and provides composable transformations for automatic parallelization, making it...



Add to Schedule

Wednesday, Mar 20 | 8:00 AM - 8:50 AM PDT  
(Thursday, Mar 21 | 4:00 AM - 4:50 AM NZDT)

© SJCC LL20A (LL)

```
import warp as wp
import jax
import jax.numpy as jp

# import Warp->JAX adapter
from warp.jax_experimental import jax_kernel

@wp.kernel
def triple_kernel(input: wp.array(dtype=float), output: wp.array(dtype=float)):
    tid = wp.tid()
    output[tid] = 3.0 * input[tid]

# create a JAX primitive from a Warp kernel
jax_triple = jax_kernel(triple_kernel)

# use the Warp kernel in a JAX jit'd function
@jax.jit
def f():
    x = jp.arange(0, 64, dtype=jp.float32)
    return jax_triple(x)

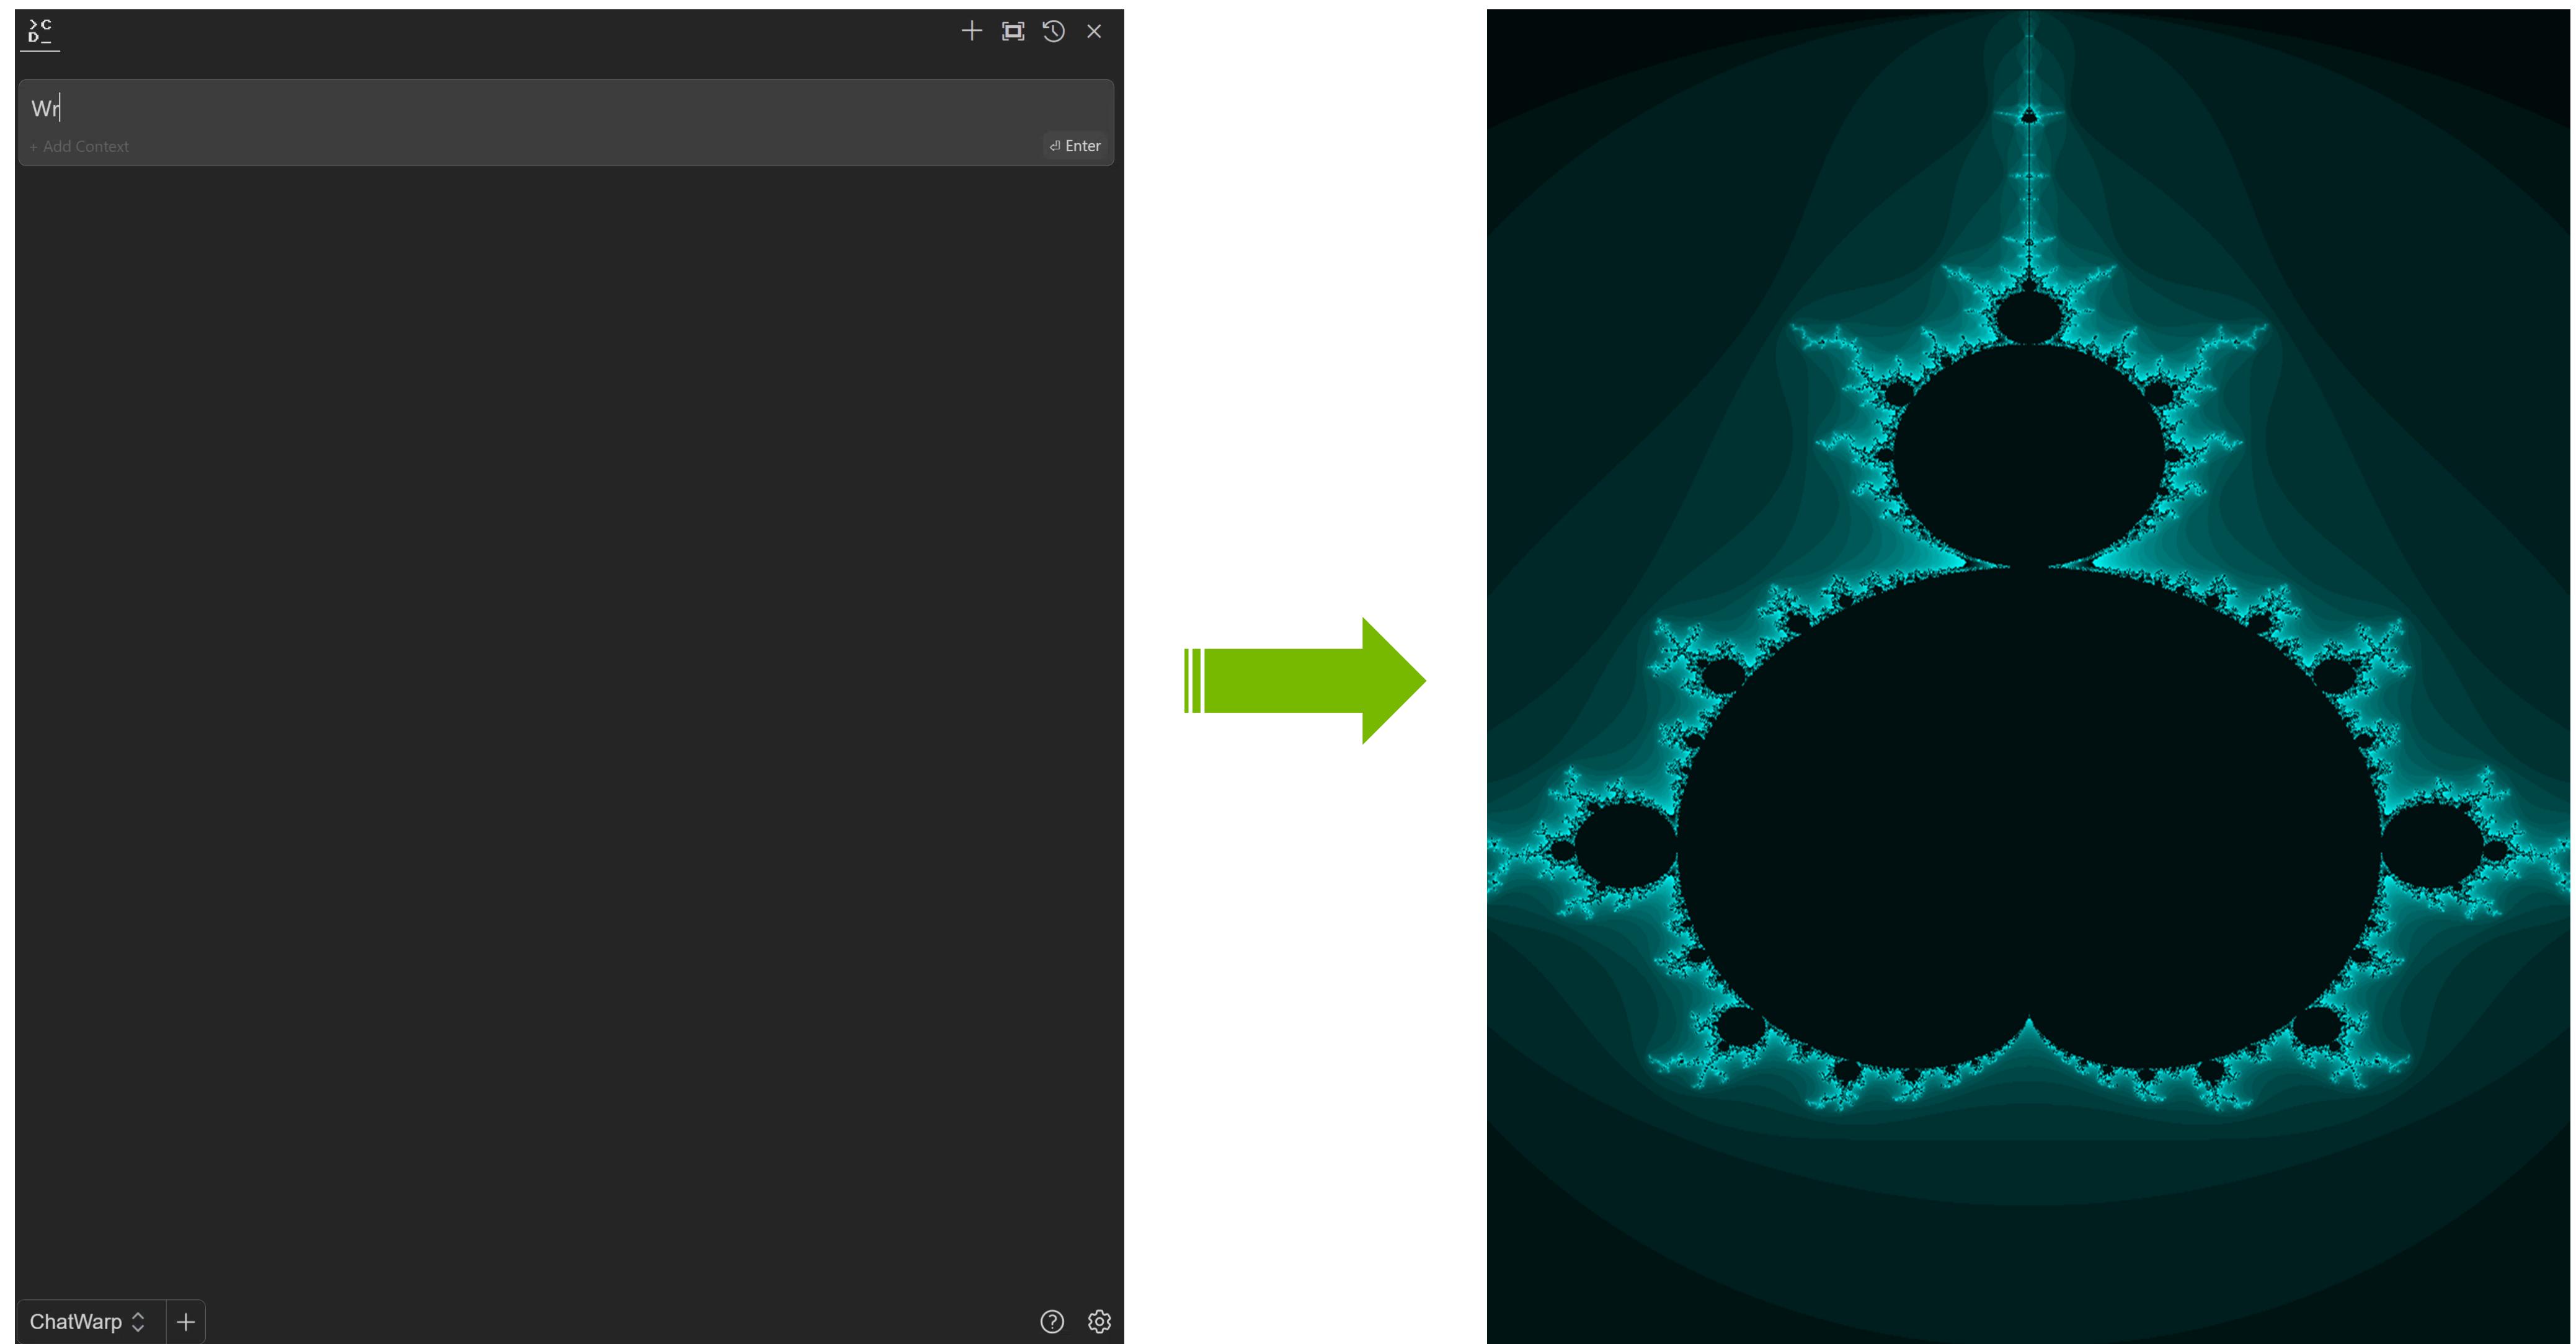
print(f())
```

Example: Warp kernel wrapped as a JAX primitive

# Road Map

## ChatWarp

- ChatWarp: a specialized LLM for generative Warp coding
- Combined RAG + Fine Tuning
- Prompt: Write a Warp kernel that a Mandelbrot texture
- Natural language -> loss function



Example: Warp Mandelbrot Generation using LLM

# Next Steps

## Installation:

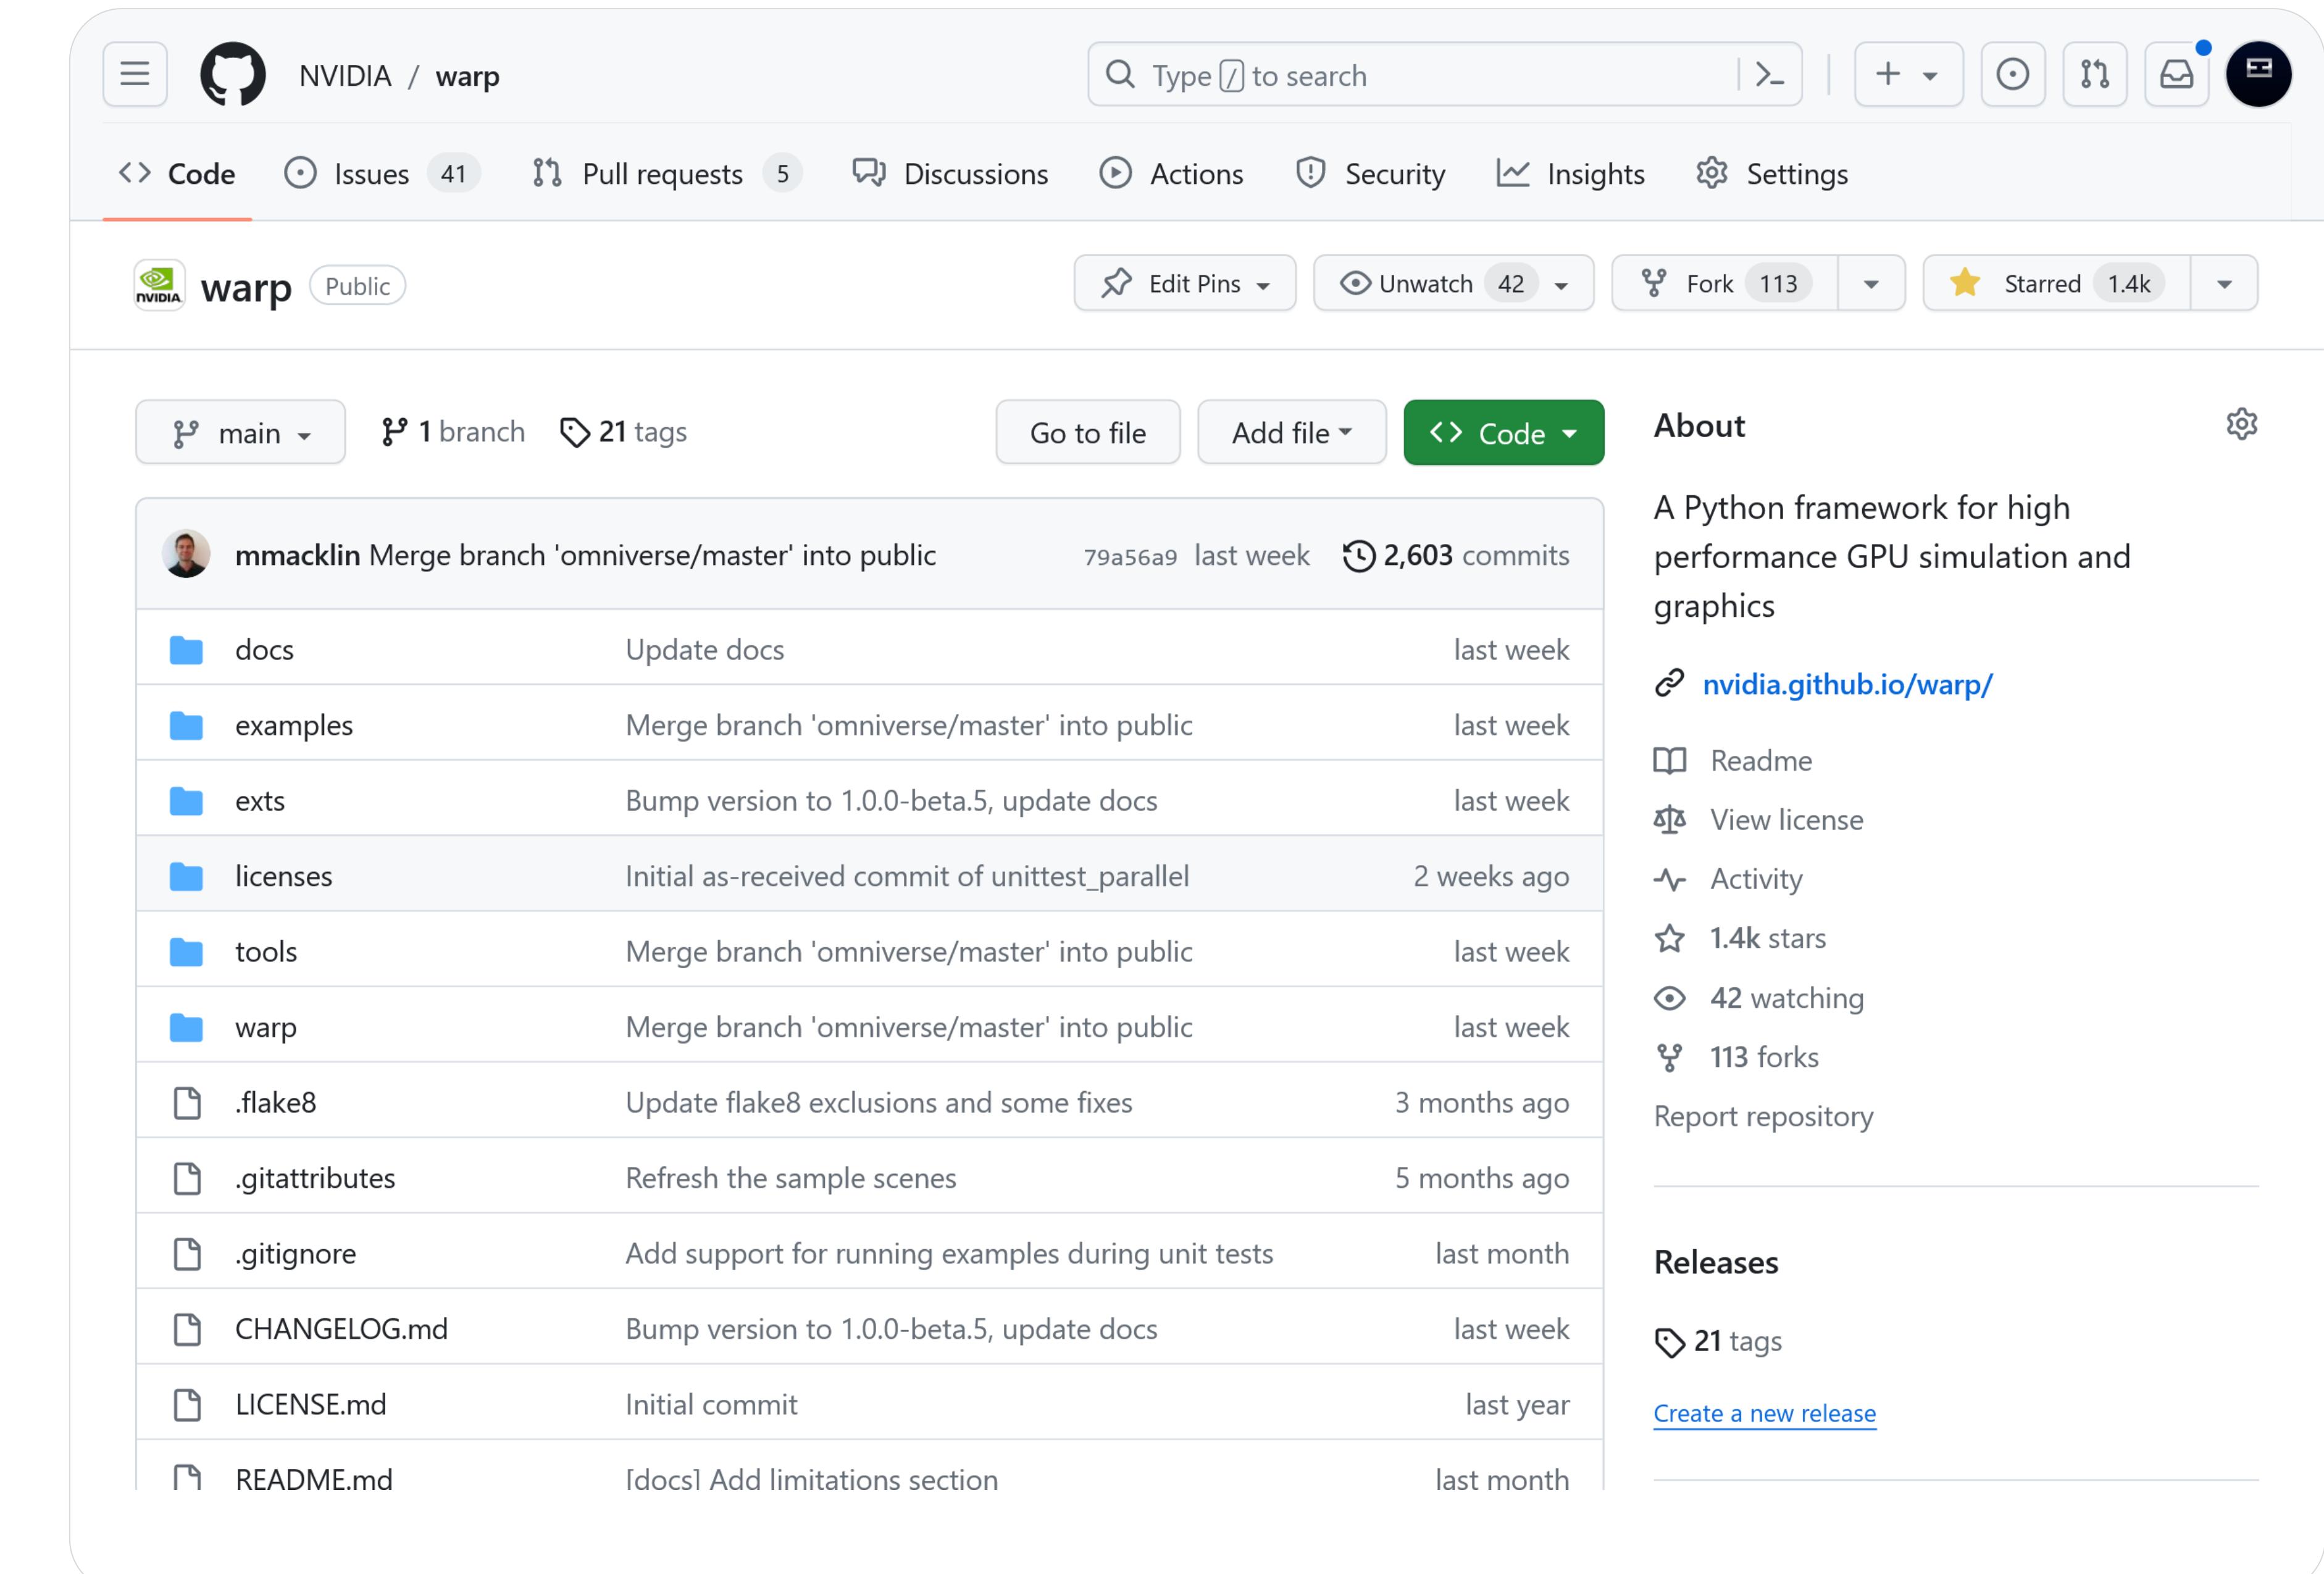
```
pip install warp-lang
```

## Platform support:

- Windows, Linux, macOS
- CUDA 11.5+, x86, aarch64, Jetson/Tegra

## Deployment:

- Released as a standalone, **open-source** library
- Binary packages with monthly release cadence on PyPi
- Diverse set of example scripts
- **omni.warp** available in Omniverse extensions registry
- Links:
  - Repo: <https://github.com/NVIDIA/warp>
  - Docs: <https://nvidia.github.io/warp/>





Questions?